



# COMPUTER ORGANIZATION

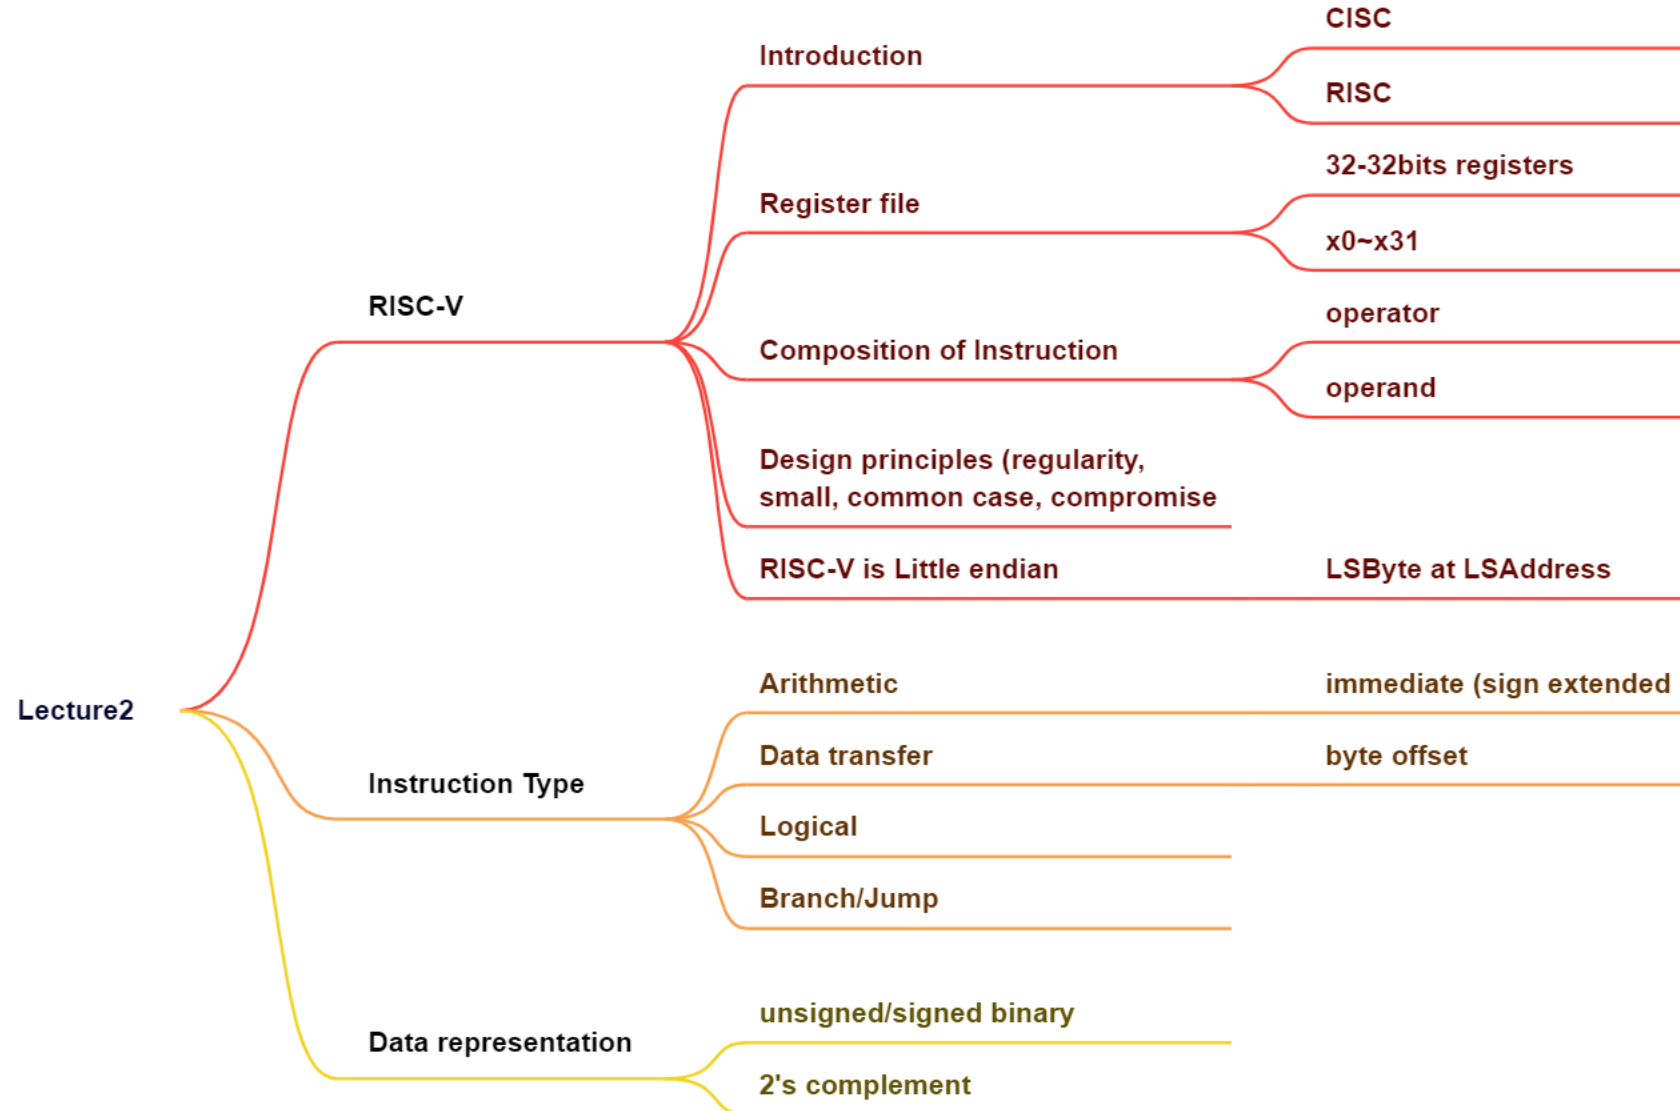
## Lecture 3 RISC-V Procedure

2025 Spring

This PowerPoint is for internal use only at Southern University of Science and Technology. Please do not repost it on other platforms without permission from the instructor.



# Recap



# Conditional Operations

- Branch to a labeled instruction if a condition is true
  - Otherwise, continue sequentially
- Conditional branch
  - *beq rs1, rs2, L1*  
if (rs1 == rs2) branch to instruction labeled L1;
  - *bne rs1, rs2, L1*
    - if (rs1 != rs2) branch to instruction labeled L1;
- Unconditional branch
  - *beq x0, x0, L1*
    - unconditional jump to instruction labeled L1

# Labels in Assembly

- We commonly see "labels" in the code
  - `foo: add x2, x1, x0`
- The assembler converts these into positions in the code
  - At what address in the code is that label ...
- Labels give control flow instructions, such as jumps and branches, a place to go ...
  - e.g. `bne x0, x2, foo`
- The assembler in outputting the code does the necessary calculation so the jump or branch will go to the right place

# Compiling If Statements

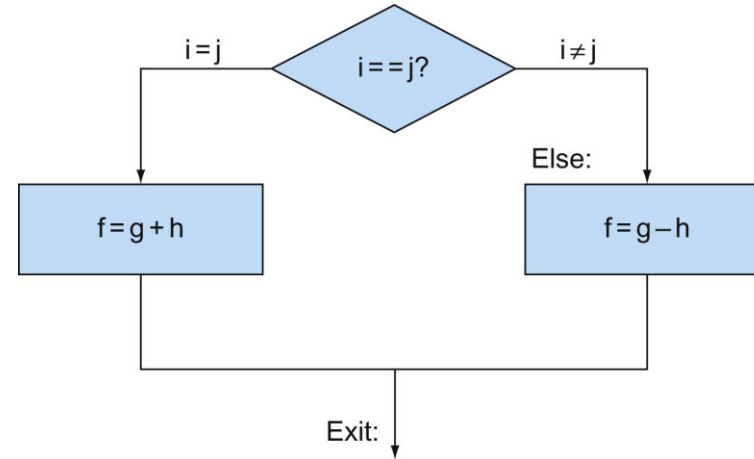
- C code

```
if (i==j) f = g+h;  
else f = g-h;
```

- i and j are in x22 and x23,
- f,g and h are in x19, x20 and x21

- Compiled RISC-V code:

```
        bne x22, x23, Else  # go to Else if i ≠ j  
        add x19, x20, x21   # f=g+h, skipped if i ≠ j  
        beq x0, x0, Exit    # unconditional go to Exit  
Else:   sub x19, x20, x21   # f=g-h, skipped if i = j  
Exit:
```



# Compiling Loop Statements

- C code:

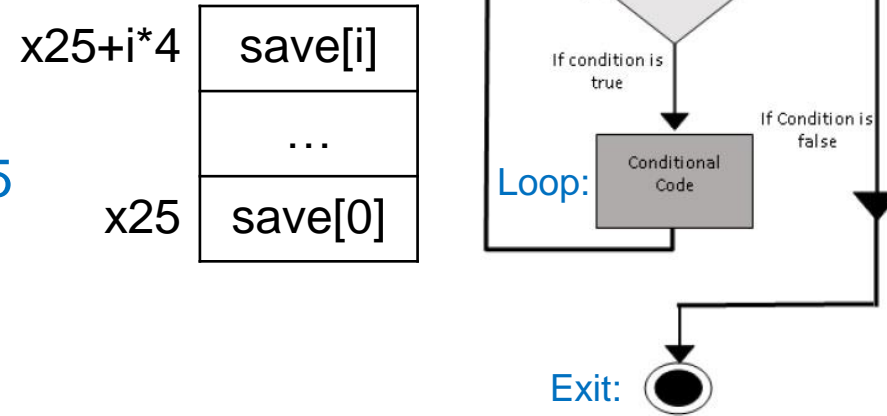
```
while (save[i] == k)
    i += 1;
```

- i in x22, k in x24, address of save in x25

- Compiled MIPS code:

```
Loop: slli    x10, x22, 2    # Temp reg x10 = i * 4
      add     x10, x10, x25  # x10 = address of save[i]
      lw      x9, 0(x10)    # Temp reg x9 = save[i]
      bne     x9, x24, Exit  # go to Exit if save[i]≠k
      addi    x22, x22, 1    # i = i + 1
      j       Loop         # pseudocode: jump to Loop

Exit:
```



# More Conditional Operations

- Signed comparison
  - `blt rs1, rs2, L1`
    - if ( $rs1 < rs2$ ) branch to instruction labeled L1
  - `bge rs1, rs2, L1`
    - if ( $rs1 \geq rs2$ ) branch to instruction labeled L1
  - Example, C to RISC-V
    - `if (a > b) a += 1;`
    - a in x22, b in x23
      - `bge x23, x22, Exit # signed comparison`
      - `addi x22, x22, 1`
  - Exit:
- Unsigned comparison
  - `bltu, bgeu`

# What if we need more instructions?

- RISC-V doesn't have "branch if greater than" or "branch if less than or equal"
- Instead you can reverse the arguments, as:
  - $A > B$  is equivalent to  $B < A$
  - $A \leq B$  is equivalent to  $B \geq A$
- The assembler defines **pseudo-instructions** for your convenience:  
`bgt x2, x3, foo (pseudo)` will become  
`blt x3, x2, foo (basic)`




# Pseudo-instructions

- For more pseudo-instructions, refer to RARS Help (see in lab).

Basic Instructions	Extended (pseudo) Instructions	Directives	Syscalls	Exco
lhu t1,10000000	Load Halfword Unsigned : Set t1 to zero-extended 16-bit value			
lhu t1,label	Load Halfword Unsigned : Set t1 to zero-extended 16-bit value			
li t1,-100	Load Immediate : Set t1 to 12-bit immediate (sign-extended)			
li t1,10000000	Load Immediate : Set t1 to 32-bit immediate			
lui t1,%hi(label)	Load Upper Address : Set t1 to upper 20-bit label's address			
lw t1,%lo(label)(t2)	Load from Address			
lw t1,(t2)	Load Word : Set t1 to contents of effective memory word address			
lw t1,-100	Load Word : Set t1 to contents of effective memory word address			
lw t1,10000000	Load Word : Set t1 to contents of effective memory word address			
lw t1,label	Load Word : Set t1 to contents of memory word at label's address			
mv t1,t2	MoVe : Set t1 to contents of t2			
neg t1,t2	NEGate : Set t1 to negation of t2			
nop	NO OPeration			
not t1,t2	Bitwise NOT (bit inversion)			

# Basic Blocks

- A basic block is a sequence of instructions with
  - No embedded branches (except at end)
  - No branch targets (except at beginning)



- A compiler identifies basic blocks for optimization
- An advanced processor can accelerate execution of basic blocks

# C to RISC-V Example

```
int A[20];  
int sum = 0;  
for (int i=0; i<20; i++)  
    sum += A[i];
```

Loop has 7 instructions

```
# Assume x8 holds pointer to A  
# Assign x10=sum, x11=i  
1  add x10, x0, x0          # sum=0  
2  add x11, x0, x0          # i=0  
3  addi x12,x0,20           # x12=20  
4  loop:  
   bge x11, x12, exit  
5  slli x13, x11, 2         # i * 4  
6  add x13, x13, x8         # A + i  
7  lw x13, 0(x13)          # *(A + i)  
8  add x10, x10, x13        # increment sum  
9  addi x11, x11, 1         # i++  
10 beq x0, x0, loop        # iterate  
11 exit:
```

# C to RISC-V Example Optimized

```
int A[20];  
int sum = 0;  
for (int i=0; i<20; i++)  
    sum += A[i];
```

Loop now has 6 instructions

```
# Assume x8 holds base address of A  
# Assign x10=sum, x11=i*4  
1  add x10, x0, x0  # sum=0  
2  add x11, x0, x0  # i=0  
3  addi x12,x0,80   # x12=20*4  
4  loop:  
   bge x11, x12, exit  
5  add x13, x11, x8  # A + i  
6  lw x13, 0(x13)    # *(A + i)  
7  add x10, x10, x13 # increment sum  
8  addi x11, x11, 4  # i++  
9  beq x0, x0, loop  # iterate  
10 exit:
```

# C to RISC-V Example Optimum

```
int A[20];
int sum = 0;
for (int i=0; i<20; i++)
    sum += A[i];
```

Loop now has 4 instructions

- Directly increment ptr into A array
- And only 1 branch/jump rather than two
  - Because first time through is always true so can move check to the end
  - The compiler will often do this automatically for optimization

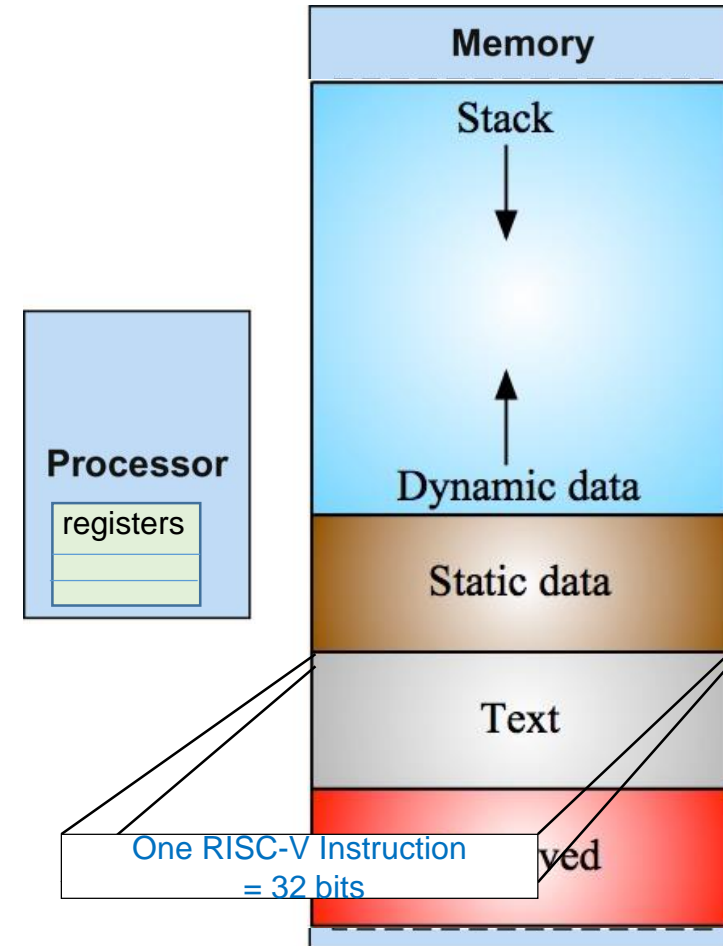
```
# Assume x8 holds base address of A
# Assign x10=sum
# Assume x11 holds ptr to next A
1 add x10, x0, x0          # sum=0
2 add x11, x0, x8          # Copy of A
3 addi x12, x8, 80         # x12=80 + A
4 loop:
  lw x13, 0(x11)
5 add x10, x10, x13
6 addi x11, x11, 4
7 blt x11, x12, loop
```

# Procedure Calling

- A procedure or function is one tool used by the programmers to structure programs
  - Benefit: easy to understand, reuse code
- We can think of a procedure like a spy
  - acquires resources → performs task → covers his tracks → returns back with desired result
- Six Steps of Calling a Function
  1. Put parameters **in a place** where the procedure can access them.
  2. **Transfer control** to the procedure.
  3. Acquire the **storage resources** needed for the procedure.
  4. **Perform** the desired task.
  5. Put the result value **in a place** where the calling program can access it.
  6. **Return control** to the **point of origin**, since a procedure can be called from several points in a program.

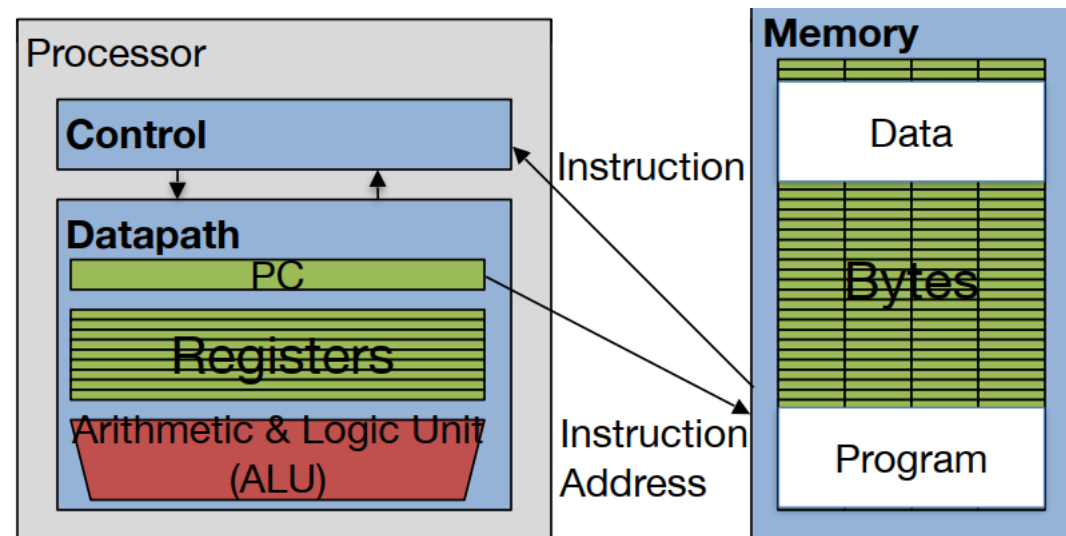
# Recall: How Program is Stored

- Instructions(programs) are represented in binary, just like data
- Programs are stored in *Text Segment*
- Constants and other static variables are stored in *Static data segment*
- Dynamic data: *Heap*
  - E.g., malloc in C, new in Java
- Automatic data: *Stack*



# Program Execution

- **PC (program counter)** is special internal register inside processor holding **byte** address of next instruction to be executed
- Instruction is fetched from memory, then control unit executes instruction using datapath and memory system, and updates program counter (default is add +4 bytes to PC, to move to next sequential instruction)





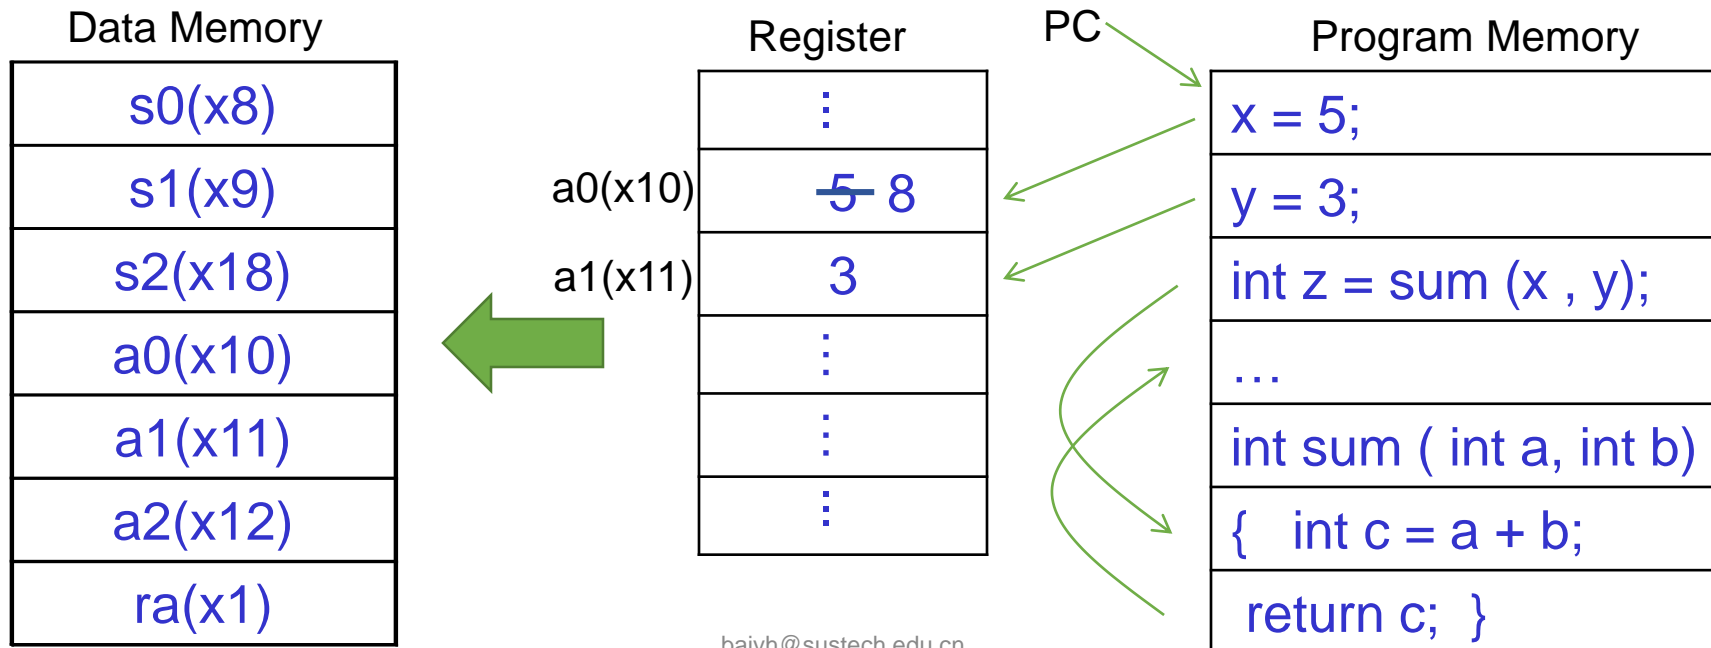
# Procedure Calling

Caller:

```
int x = 5;
int y = 3;
int z = sum (x , y);
x = x + 7;
...
```

callee:

```
int sum ( int a, int b)
{
    int c = a + b;
    return c;
}
```



# Procedure Calling: Question1

- Step 1, 5 and 6: Where should we put the arguments and return values?
  - Registers way faster than memory, so use them whenever possible
  - Symbolic register names
    - E.g., a0-a7 (x10-x17) for function arguments, a0-a1 for return values
    - E.g., ra (x1) for return address, used to save where a function is called from so we can get back
  - If need extra space, use memory (the stack!)

# RISC-V Registers and Convention

- ABI: Application Binary Interface, defines “Calling Convention” – How to call other functions
- So going forward, no more x3, x6... type nomenclature

Register	ABI name	Description	Saved by Callee?
x0	zero	the constant value 0	N/A
x1	ra	Return address	No
x2	sp	Stack pointer	Yes
x3	gp	Global pointer	N/A
x4	tp	Thread pointer	N/A
x5 – 7	t0 – 2	Temporaries	No
x8	s0 / fp	Saved register/Frame pointer	Yes
x9	s1	Saved register	Yes
x10 – 17	a0 – 7	Function arguments/Return values	No
x18 – 27	s2 – 11	Saved registers	Yes
x28 – 31	t3 – 6	Temporaries	No

# Procedure Calling: Question2

- Step 2 and 6: How do we Transfer Control?
  - Procedure call: **jump and link** (the real order is link and jump)  
`jal ra, target` #pseudo code: `jal target`
    - Address of following instruction (return address) put in ra
    - Jumps to target address
    - Used by Caller
  - Procedure return: **jump register**  
`jalr zero, 0(ra)` #pseudo code: `jr ra`
    - Similar to “Jump and Link” except in specification of target
    - Jump to ra and simultaneously saves the address of following instruction in zero register (value put into zero is meant to be throw away)
    - Used by Callee

# Transfer Control

Caller:

```
int x = 5; // x in s0
int y = 3; // y in s1
int z = sum (x , y);
...
```

callee: // a in a0, b in a1, c in a0

```
int sum ( int a, int b) {
    int c = a + b;
    return c;
}
```

address (shown in decimal)

```
1000 mv a0,s0      # parse argument x → a
1004 mv a1,s1      # parse argument y → b
1008 jal sum       # ra=1012, goto sum
1012 ...           # next instruction
```

```
...
2000 sum: add a0,a0,a1
```

```
2004 jr ra        # next instr at 1012...
```

Note: these are pseudo code

PC's decimal value: ...→1000→1004→1008→2000→2004→1012→...

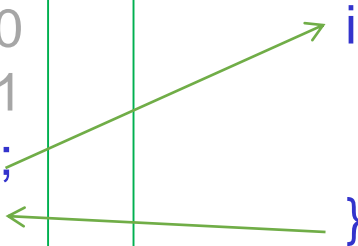
# Transfer Control

Caller:

```
int x = 5; // x in s0  
int y = 3; // y in s1  
int z = sum (x , y);  
...
```

callee: // a in a0, b in a1, c in a0

```
int sum ( int a, int b) {  
    int c = a + b;  
    return c;  
}
```



- Question: Why use **jr** here? Why not use **j**?
- Answer: **sum** might be called by many places, so we can't return to a fixed place. The calling proc to **sum** must be able to say "return here" somehow.

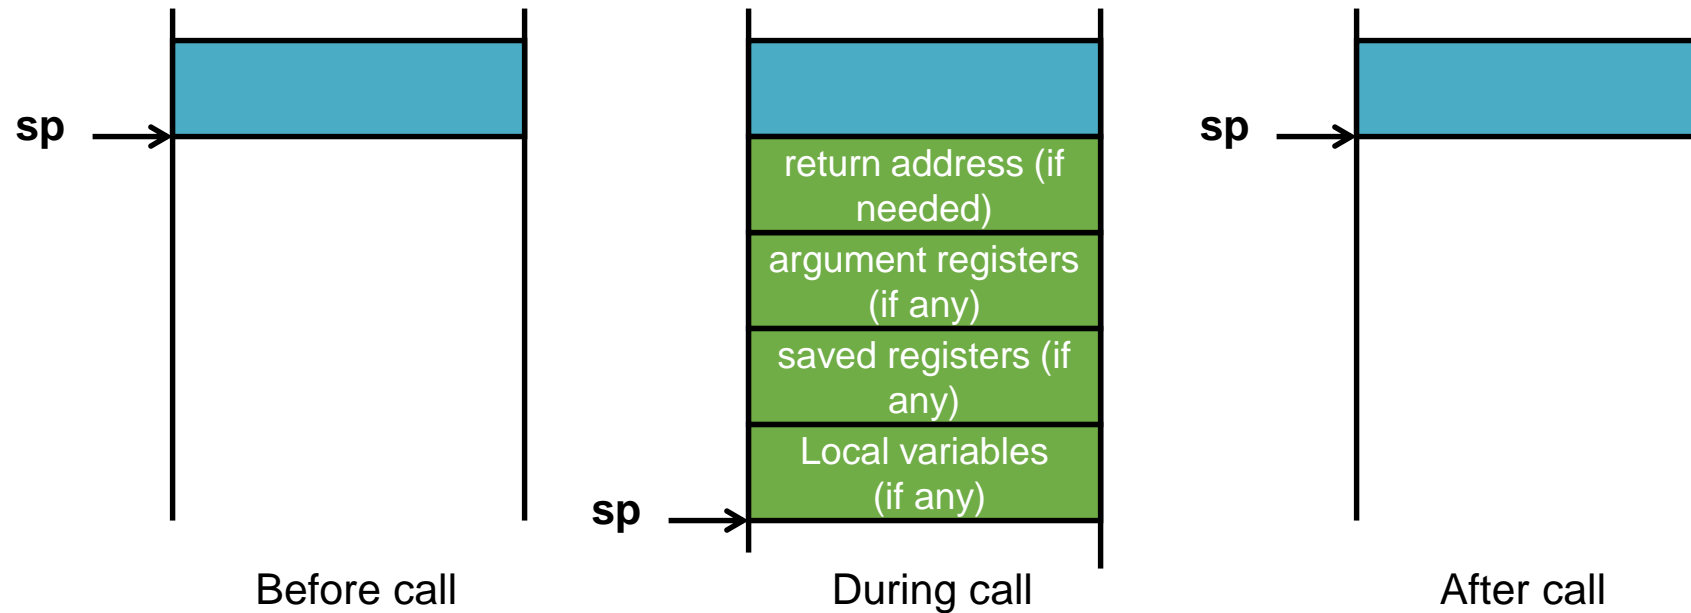
```
2000 sum: add a0, a0, a1
```

```
2004 jr ra # jump to the return address
```

# Procedure Calling: Question3

- Step 3: What's the Local storage for variables?
  - C has three storage classes:
    - Stack: automatic variables are local to function and discarded when function exits
    - Static: variables exist across exits from and entries to procedures
    - Heap: variables declared dynamically via malloc
  - Use stack for automatic (local) variables that don't fit in registers
    - Push: placing data onto stack
    - Pop: removing data from stack
  - **sp (x2) (stack pointer)** points to the last used place at the stack
    - Push decrements sp
    - Pop increments sp

# Stack Before, During, After Call





# Procedure Calling: Question 4

- Step 4: Function Calling Conventions?
  - It is effectively a contract between functions
  - By convention, registers are classified as one of ...
  - **caller-saved**
    - The function invoked (the callee) can do whatever it wants to them!
    - Means that the caller can not count on their contents not being destroyed
  - **callee-saved**
    - The function invoked must restore them before returning (if used)

# The Calling Convention

- Caller saved:
  - **a0–a7 (x10-x17)**: eight argument registers to pass parameters and two return values (**a0-a1**)
  - **t0-t6** Temporaries
  - **ra**: one return address register for return to the point of origin
- Callee saved:
  - **s0-s11** Saved registers: Preserved across function calls
- If both the Caller and Callee obey the procedure conventions, there are significant benefits
  - People who have never seen or even communicated with each other can write functions that work together
  - Recursion functions work correctly

# Leaf Procedure Example

- A “leaf” function - it calls nothing

```
int Leaf(int g, int h, int i, int j) {  
    int f;  
    f = (g + h) - (i + j);  
    return f;  
}
```

- Parameter variables g, h, i, and j in argument registers a0, a1, a2, and a3.
- Return variable f uses register a0
- Assume we compute f by using s0 and s1
- s0 and s1 are callee saved, so it's the responsibility of “Leaf” to save and restore

# Leaf Procedure Example

- RISC-V code:

a0, a1, a2, a3

```
int Leaf(int g, int h, int i, int j)
{
    int f;
    f = (g + h) - (i + j);
    return f;
}
```

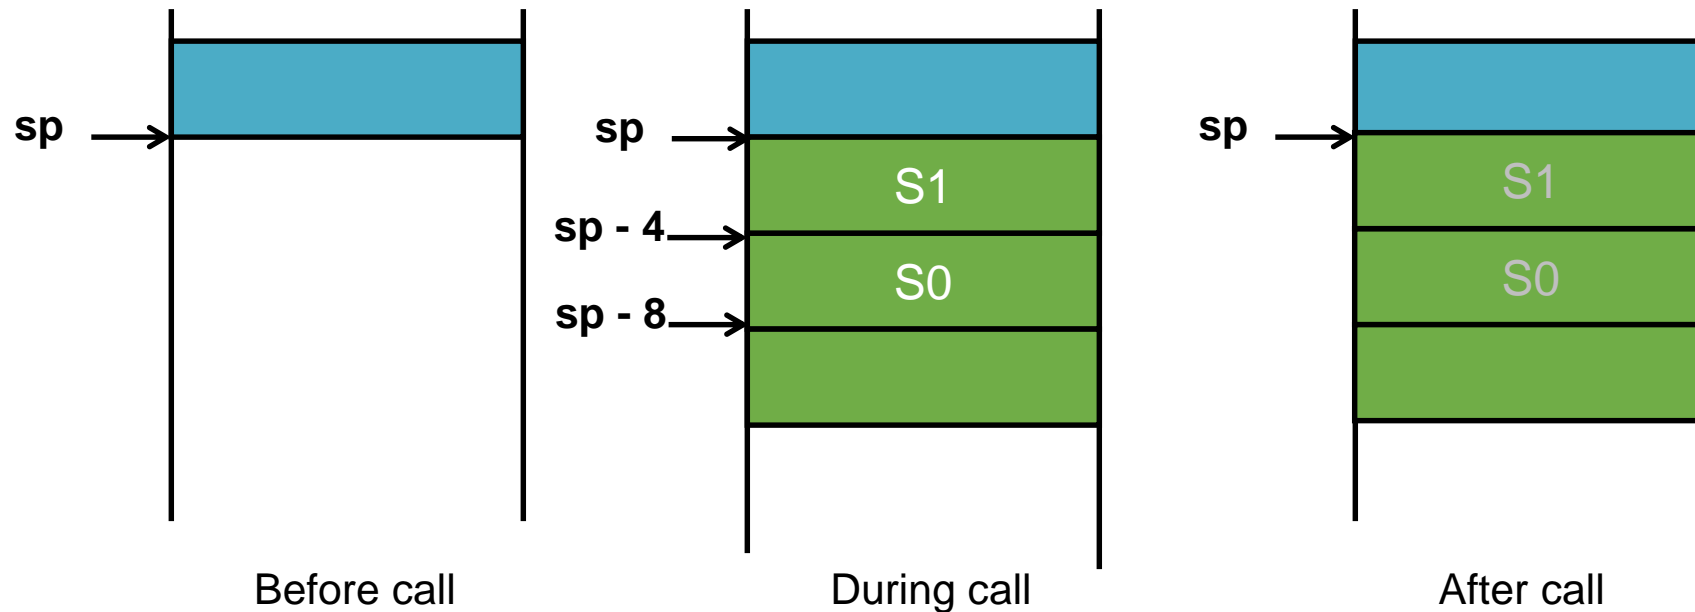
a0                      s0      s1

Leaf:

```
1  addi sp,sp,-8 # adjust stack for 2 items
2  sw s1, 4(sp)  # save s1 for use afterwards
3  sw s0, 0(sp)  # save s0 for use afterwards
4  add s0,a0,a1  # s0 = g + h
5  add s1,a2,a3  # s1 = i + j
6  sub a0,s0,s1  # return value (g + h) - (i + j)
7  lw s0, 0(sp)  # restore register s0 for caller
8  lw s1, 4(sp)  # restore register s1 for caller
9  addi sp,sp,8  # adjust stack to delete 2 items
10 jr ra         # jump back to calling routine
                        # pseudo of jalr zero, 0(ra)
```

# Leaf Procedure Example: Stack

- In the previous example we need to save old values of s0 and s1
- Push doesn't actually delete content from memory, it just increase the stack pointer



# Leaf Procedure Example: Observation

- This is a "leaf function": it calls no other function
  - We didn't need to save ra (because leaf didn't call any other function and therefore ra never changed)
  - Instead of s0 and s1 can just use temporary (caller-saved registers) only
- So we could have just as easily used t0 and t1 instead...

leaf:

```
add t0,a0,a1 # t0 = g + h
add t1,a2,a3 # t1 = i + j
sub a0,t0,t1 # return value (g+h)-(i+j)
jr ra
```

Caller saved:

a0–a7

t0-t6

ra

Callee saved:

s0-s11

# Non-Leaf Procedures

- Procedures that call other procedures (nested call)
  - Note: this could mean a function calling itself – recursion
- For nested call, caller needs to save on the stack:
  - Its return address
  - Any arguments and temporaries needed after the call
- Restore from the stack after the call

# Non-Leaf Procedure Example

- C code:

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y;  
}
```

- Function **sumSquare** is calling **mult**
- So there's a value in **ra** that **sumSquare** wants to jump back to, but this will be overwritten by the call to **mult**
- Need to save sumSquare return address before call to mult



# Non-Leaf Procedure Example

a0, a1

```
int sumSquare(int x, int y) {  
    return mult(x,x)+ y;  
}
```

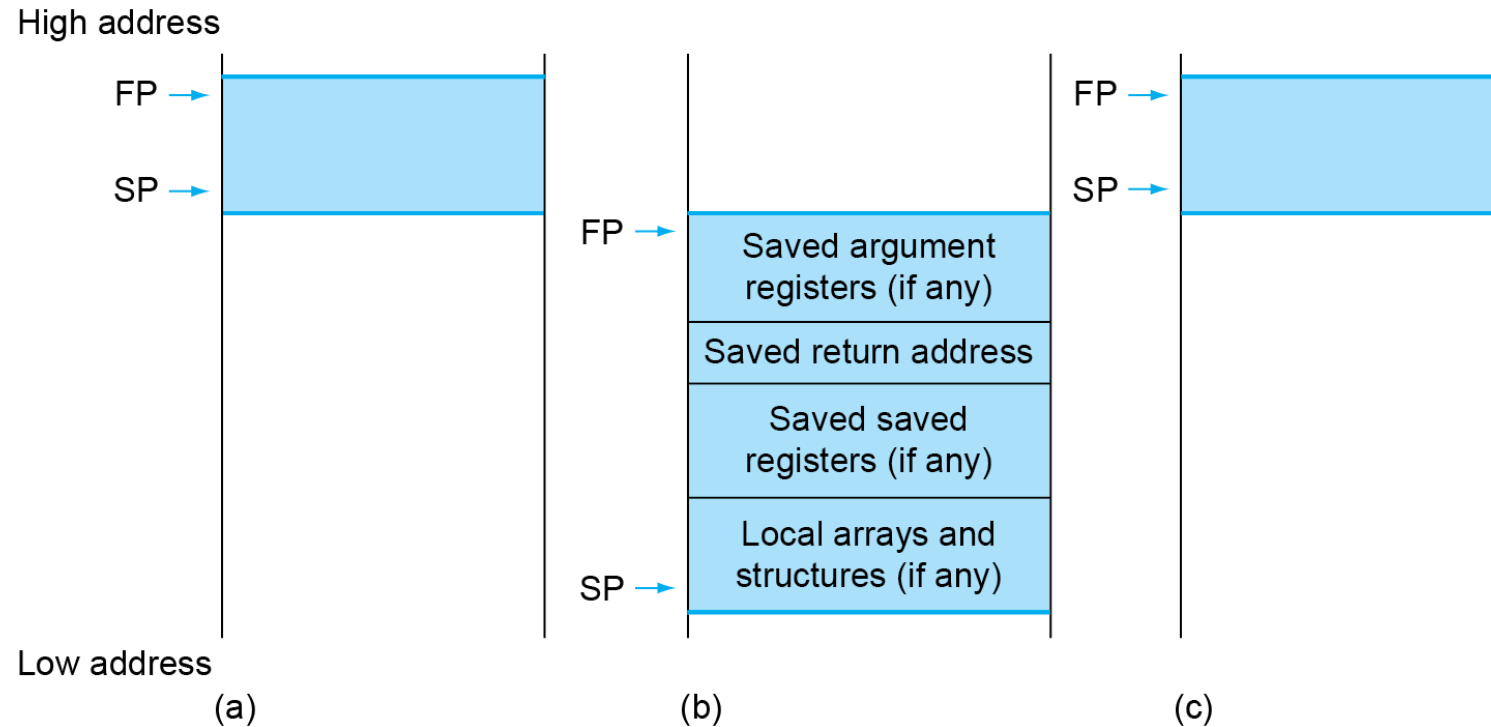
a0, a1

```
sumSquare:  
1  addi sp,sp,-8 # reserve space on stack  
2  sw ra, 4(sp) # save ret addr          Push  
3  sw a1, 0(sp) # save y  
4  mv a1,a0     # prepare 2nd argument for mult  
5  jal mult     # call mult  
6  lw a1, 0(sp) # restore y  
7  add a0,a0,a1 # mult()+y  
8  lw ra, 4(sp) # get ret addr  
9  addi sp,sp,8 # restore stack          Pop  
10 jr ra  
    mult:
```

# Register Conventions Summary

- Calle<sup>R</sup> must save arguments and temporary registers it is using onto the stack before making a procedure call
- Calle<sup>R</sup> can trust saved registers to maintain values
- Calle<sup>E</sup> must save any saved registers it intends to use by putting them on the stack before overwriting their values
- Notes:
  - Calle<sup>R</sup> and calle<sup>E</sup> only need to save the **appropriate using** (not all!)
  - Don't forget to restore the values later

# Local Data on the Stack



- Local data allocated by callee
  - e.g., C automatic variables
- Procedure frame (activation record)
  - Used by some compilers to manage stack storage

# Summary

- Functions called with `jal`, return with `jalr (jr)`
- The stack is your friend: Use it to save anything you need. Just leave it the way you found it!
- Instructions we know so far...
  - Arithmetic
  - Memory
  - Conditional Branches
  - Unconditional Branches (Jumps)
- Registers

# Non-Leaf Procedure Example 2

```
int fact (int n)
{
    if (n < 1) return (1);
    else return (n * fact(n-1));
}
```

## Notes:

The caller saves a0 and ra in its stack space.

Temps are never saved.

Compare  $n < 1$

Return 1

Fact(n-1)

Return  $n * \text{fact}(n-1)$   
baiyh@sustech.edu.cn

```
fact:
    addi    sp, sp, -8
    sw      ra, 4(sp)
    sw      a0, 0(sp)
    addi    t0, a0, -1
    bge     t0, zero, L1
    addi    a0, zero, 1
    addi    sp, sp, 8
    jr      ra

L1:
    addi    a0, a0, -1
    jal     fact
    addi    t1, a0, 0
    lw      a0, 0(sp)
    lw      ra, 4(sp)
    addi    sp, sp, 8
    mul     a0, a0, t1
    jr      ra
```

# Non-Leaf Procedure Example 2

- RISC-V code: suppose  $n=2$

```

fact:
1      addi    sp, sp, -8
2      sw      ra, 4(sp)
3      sw      a0, 0(sp)
4      addi    t0, a0, -1
5      bge     t0, zero, L1
6      addi    a0, zero, 1
7      addi    sp, sp, 8
8      jr      ra
9  L1: addi    a0, a0, -1
10     jal     fact
11 Y:  mv      t1, a0
12     lw      a0, 0(sp)
13     lw      ra, 4(sp)
14     addi    sp, sp, 8
15     mul     a0, a0, t1
16     jr      ra
    
```

X: return address of caller  
(not shown in the code)  
Y: the address of line 11

```

int fact (int n)
{ if (n < 1) return 1;
  else return n * fact(n - 1); }
    
```

