

CS 305: Computer Networks

Fall 2025

Lecture 5: Application Layer

Ming Tang

Department of Computer Science and Engineering
Southern University of Science and Technology (SUSTech)

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP


Video Streaming and CDNs: ~~context~~

- Video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
 - ~1B YouTube users, ~75M Netflix users
- Challenge: scale - how to reach ~1B users?
 - single mega-video server won't work (why?)
- Challenge: heterogeneity
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- **Solution:** distributed, application-level infrastructure



Multimedia: video

- Video: sequence of images displayed at constant rate
 - e.g., 24 images/sec
- Digital image: array of pixels
 - each pixel represented by bits




Multimedia: video

Coding (Compression): use redundancy *within* and *between* images to decrease # bits used to encode image

- spatial (within image)
- temporal (from one image to next)

Spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (purple) and number of repeated values (N)



frame i



frame $i+1$


temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i

Multimedia: video

Type	Video Bitrate, Standard Frame Rate (24, 25, 30)	Video Bitrate, High Frame Rate (48, 50, 60)
2160p (4k)	35-45 Mbps	53-68 Mbps
1440p (2k)	16 Mbps	24 Mbps
1080p	8 Mbps	12 Mbps
720p	5 Mbps	7.5 Mbps
480p	2.5 Mbps	4 Mbps
360p	1 Mbps	1.5 Mbps

- **CBR: (constant bit rate):** video encoding rate fixed
- **VBR: (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes

HTTP Streaming



All clients receive the same encoding of the video:


- Human users may have different requirements
- Clients may have different available bandwidth, which may be time-varying

How to deal with this?



Streaming multimedia: DASH

- **DASH: D**ynamic, **A**daptive **S**treaming over **H**TTP
- **Server:**
 - divides video file into **multiple chunks**
 - each chunk stored, encoded at different rates
 - **manifest file:** provides URLs for different chunks encoded at different rates
- **Client:**
 - periodically measures server-to-client bandwidth
 - consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time)

Streaming multimedia: DASH




Streaming multimedia: DASH



Media presentation description (MPD), also known as manifest

Streaming multimedia: DASH



Streaming multimedia: DASH

“intelligence” at client: client determines

- **when** to request chunk (so that buffer starvation, or overflow does not occur)
- **what encoding rate** to request (higher quality when more bandwidth available)
- **where** to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications


2.6 video streaming and **content distribution networks**

2.7 socket programming with UDP and TCP



Content distribution networks




来源: 小黑羊JoinWings
https://baijia.baidu.com/s?old_id=126615&wfr=pc&fr=app_list



Content distribution networks



Content distribution networks



Content distribution networks

- **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- **Option 1:** single, large “mega-server”
 - single point of failure
 - huge traffic
 - long path to distant clients
 - multiple copies of video sent over outgoing link

....quite simply: this solution **doesn't scale**


Content distribution networks

- **Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- **Option 2: Content distribution networks (CDN)** store/serve multiple copies of videos at multiple geographically distributed sites
 - **Enter deep:** push CDN servers deep into many access networks; inside ISPs
 - close to users
 - used by Akamai, 1700 locations
 - **Bring home:** smaller number (10's) of larger clusters in Internet Exchange Point (IXP); outside ISPs
 - used by Limelight

Content Distribution Networks (CDNs)

CDN: stores copies of content at CDN nodes

- e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
 - directed to nearby copy, retrieves content
 - may choose different copy if network path congested



Content Distribution Networks (CDNs)


Challenges: Coping with a congested Internet

- what content to place in CDN node?
 - Simple pull strategy: request, then store
- from which CDN node to retrieve content?
 - Cluster selection strategy
- the operation for retrieving content?

CDN Operation

Bob (client) requests video <http://netcinema.com/6Y7B23V>

- video stored in CDN at <http://KingCDN.com/NetC6y&B23V>



CDN: Cluster Selection Strategy

One simple strategy is to assign the client to the cluster that is

geographically closest:

- When a DNS request is received from a particular LDNS, the CDN chooses the geographically closest cluster
- may not be the closest cluster in terms of the length or number of hops
- ignore the variation in delay and available bandwidth over time

Periodic **real-time measurements** of delay and loss performance between their clusters and clients:

- a CDN can have each of its clusters periodically send probes to all of the LDNSs around the world.
- many LDNSs are configured to not respond to such probes.

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP