

# Principles of Database Systems (CS307)

## Lecture 7: Advanced SQL and Database Design using E-R Model

Zhong-Qiu Wang

Department of Computer Science and Engineering  
Southern University of Science and Technology

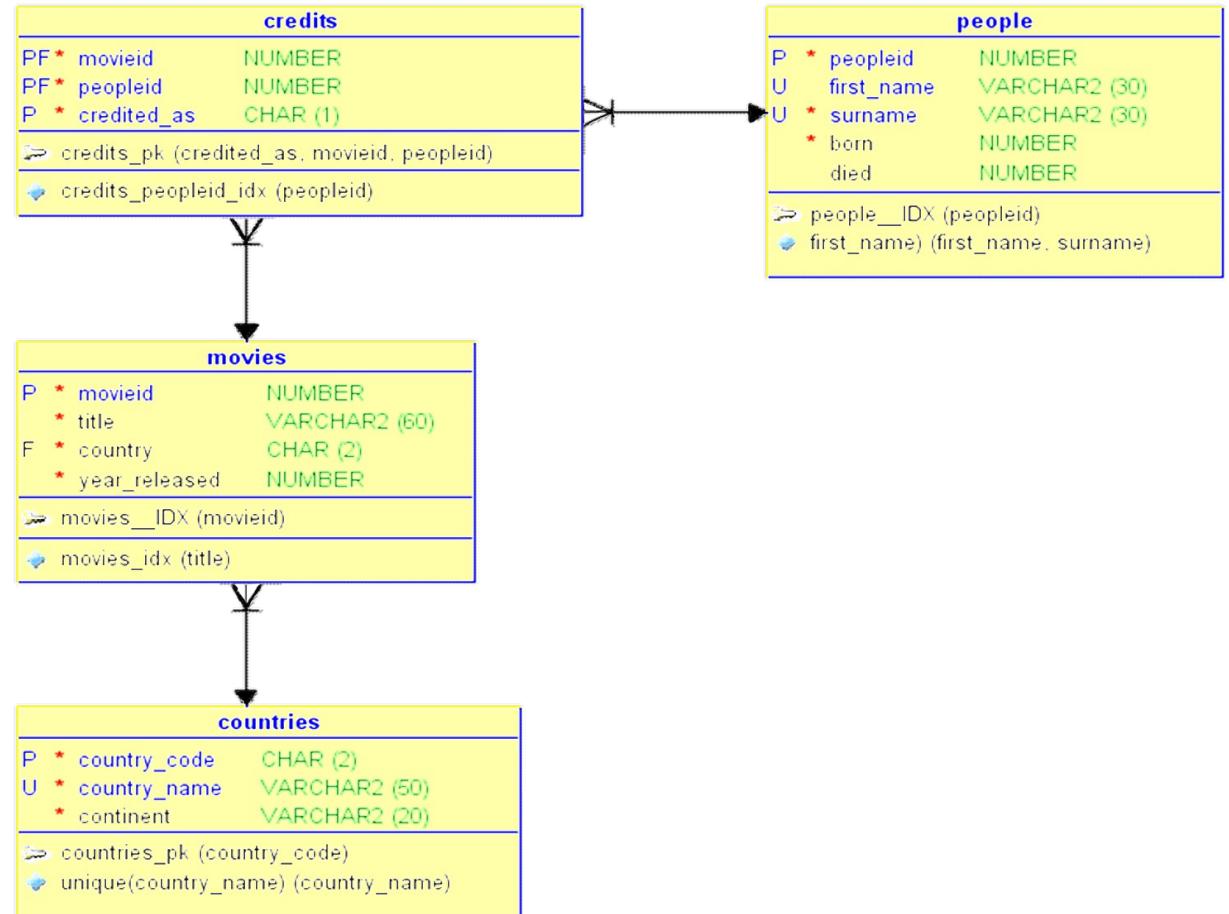
- Most contents are from slides made by Stéphane Faroult and the authors of Database System Concepts (7<sup>th</sup> Edition).
- Their original slides have been modified to adapt to the schedule of CS307 at SUSTech.
- The slides are largely based on the slides provided by Dr. Yuxin Ma

# Announcements

- Project I is out, due date: 23:59 on November 10th 2024, Beijing Time
  - Only accept project report in English
    - The class is designed as an English class
    - Can avoid potential unfairness when grading Chinese and English reports

# Entity and Relationship

- Starring -> Actor table
- Country -> Country and Region table
  - You can also link the movies with corresponding actors, countries/regions, etc.
- **Entity Relationship Diagram (E/R Diagram, ER Diagram, ERD)**
  - A way of representing entity tables and their relationships (relationship tables)
  - Connect tables via **foreign keys** and **relationship tables**



# Create Tables

- An SQL relation is defined using the create table command:

```
create table r (
    A1 D1, A2 D2, ..., An Dn,
    (integrity-constraint1),
    ...,
    (integrity-constraintk)
)
```

- **Relation schema**: r is the name of the relation
- **Attribute**: each  $A_i$  is an attribute name in the schema of relation r
- **Data type**:  $D_i$  is the data type of values in the domain of attribute  $A_i$
- **Constraints**: not null, unique, primary key, check function, referential integrity & foreign key

# Select

- `select * from [tablename]`
  - The select clause lists the attributes desired in the result of a query
  - To display the full content of a table, you can use `select *`
    - \* : all columns

```
select A1, A2, ..., An  
from r1, r2, ..., rm  
where P
```

# Some Functions

- When to use functions
  - An example of showing a result that isn't stored as such is **computing an age**
    - You should never store an age; it changes all the time!
    - If you want to display the age of people who are alive, you must compute their age by subtracting the year when they were born from the current year.
  - In the table people:
    - Alive – died is null
    - Age: <this year> - born



```
select peopleid, surname,  
       date_part('year', now()) - born as age  
from people  
where died is null;
```

|   |            |           |      |        |   |
|---|------------|-----------|------|--------|---|
| 7 | 7 Caroline | Aaron     | 1952 | <null> | F |
| 8 | 8 Quinton  | Aaron     | 1984 | <null> | M |
| 9 | 9 Dodo     | Abashidze | 1924 | 1990   | M |

# Aggregate Functions

`count(*)/count(col), min(col), max(col), stddev(col), avg(col)`

- These aggregate functions are supported in almost every DBMS
  - Most DBMS implement other more advanced functions
  - Some functions work with any datatype, others only work with numerical columns
- It is strongly recommended to refer to user manual for details
  - For example, SQLite doesn't have `stddev()`

# Retrieving Data from Multiple Tables

- This is done with this type of query. We retrieve, and display as a single set, pieces of data coming from two different tables.

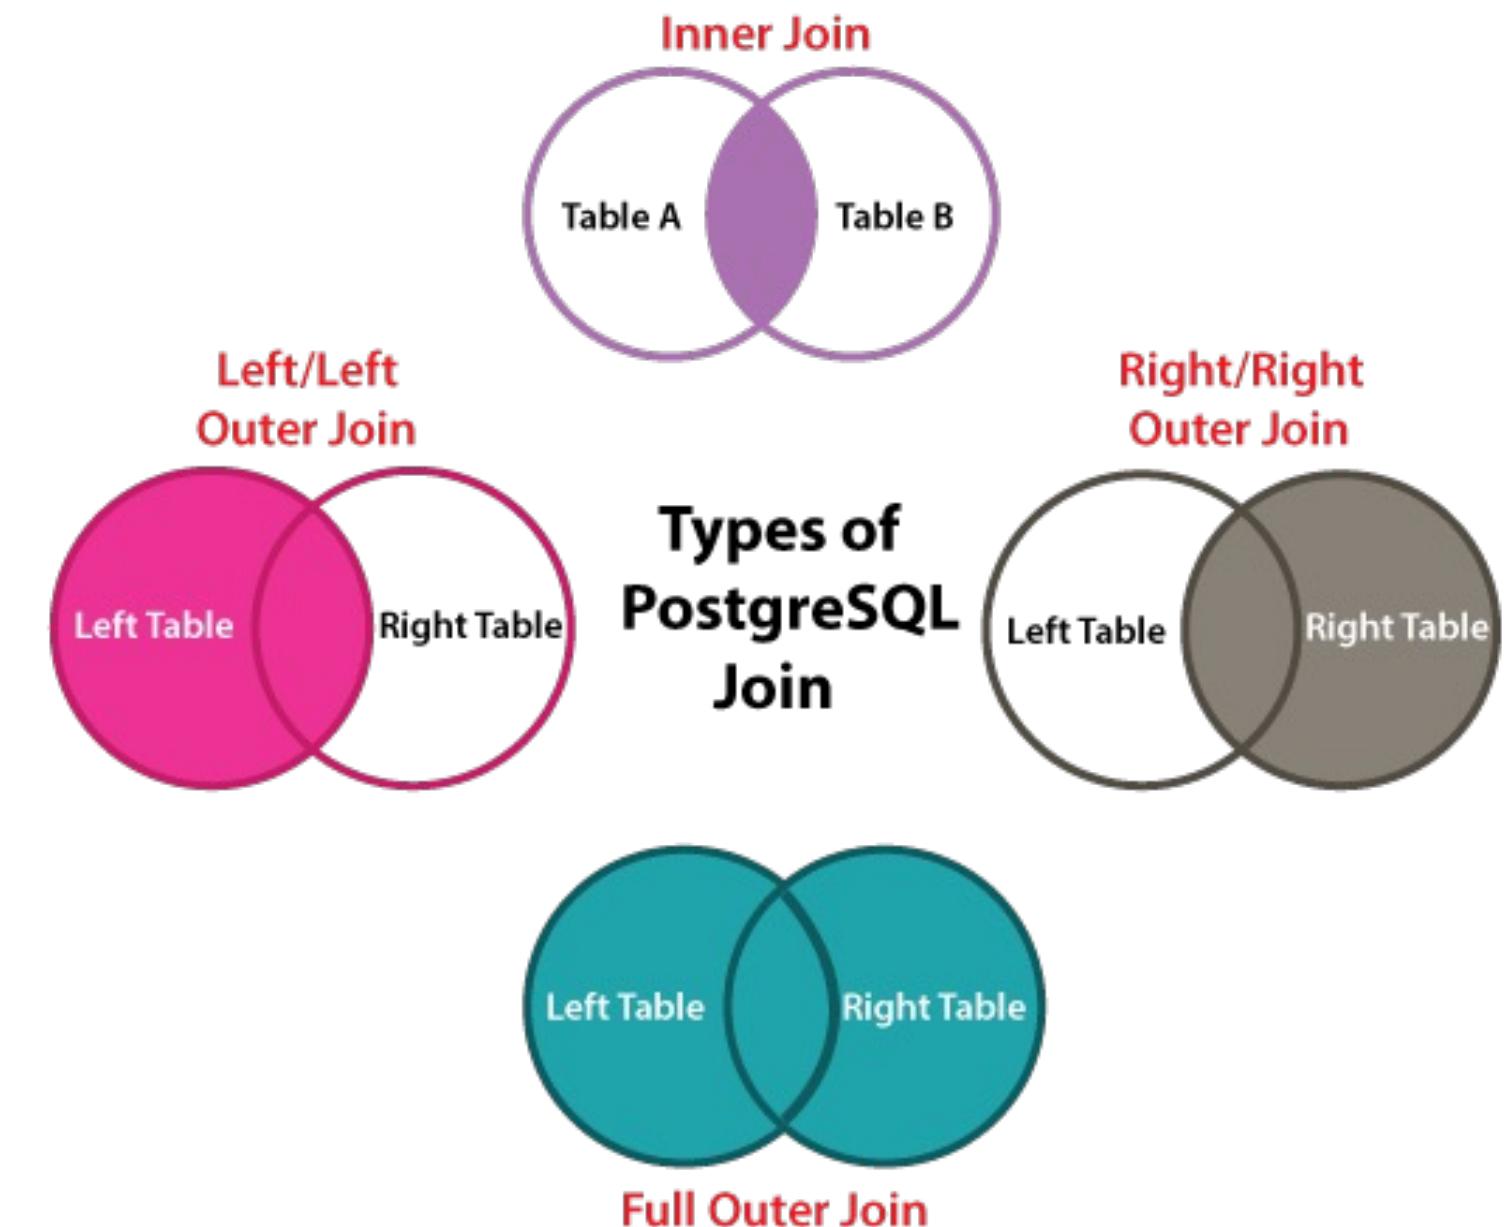


```
select title,  
       country_name,  
       year_released  
  from movies  
 join countries  
  on country_code = country;
```

| title                                 | country_name  | year_released |
|---------------------------------------|---------------|---------------|
| 12 stulyev                            | Russia        | 1971          |
| Al-mummia                             | Egypt         | 1969          |
| Ali Zaoua, prince de la rue           | Morocco       | 2000          |
| Apariencias                           | Argentina     | 2000          |
| Ardh Satya                            | India         | 1983          |
| Armaan                                | India         | 2003          |
| Armaan                                | Pakistan      | 1966          |
| Babettes gæstebud                     | Denmark       | 1987          |
| Banshun                               | Japan         | 1949          |
| Bidaya wa Nihaya                      | Egypt         | 1960          |
| Variety                               | United States | 2008          |
| Bon Cop, Bad Cop                      | Canada        | 2006          |
| Brilliantovaja ruka                   | Russia        | 1969          |
| C'est arrivé près de chez vous        | Belgium       | 1992          |
| Carlota Joaquina - Princesa do Brasil | Brazil        | 1995          |
| Cicak-man                             | Malaysia      | 2006          |
| Da Nao Tian Gong                      | China         | 1965          |
| Das indische Grabmal                  | Germany       | 1959          |
| Das Leben der Anderen                 | Germany       | 2006          |
| Den store gavtyv                      | Denmark       | 1956          |

# Inner and Outer Joins

- So far, we only join the rows with matching values on the corresponding columns
- There are more things we can do with join



# Subquery

- We have used subqueries after the keyword **from** before
  - ... in order to build queries upon a query result
- And, we can add subqueries after **select** and **where** as well

```
select A1, A2, ..., An  
  from r1, r2, ..., rm  
    where P
```

# Advanced Ordering

- Multiple columns
  - The result set will be `order by col1` first
  - Rows with the same value on `col1` will be ordered by `col2`
- Ascending or descending order
  - Add `desc` or `asc` after the column
    - `asc` is the default option and thus always omitted



`order by col1, col2, ...`



`-- Order col1 descendingly  
order by col1 desc`

`-- Order based on col1 first, then col2.  
-- col1 will be in the descending order, col2 ascending.  
order by col1 desc, col2 asc, ...`

# Window Function

- Syntax:

```
<function> over (partition by <col_p> order by <col_o1, col_o2, ...>)
```

- **<function>**
  - Ranking window functions
  - Aggregation functions
- **partition by**
  - Specify the column for grouping
- **order by**
  - Specify the column(s) for ordering in each group

# Ranking Window Function

- Example
  - How can we rank the movies in each country separately based on the released year?
    - “order by” for subgroups



```
select country,
       title,
       year_released,
       rank() over (
           partition by country order by year_released
       ) oldest_movie_per_country
from movies;
```

Partitioned by country

- i.e., a country in a group

| country | title      | year_released | oldest_movie_per_country |
|---------|------------|---------------|--------------------------|
| ar      | some title | 1948          | 1                        |
|         | some title | 1959          | 2                        |
|         | some title | 1980          | 3                        |
| cn      | some title | 1987          | 1                        |
|         | some title | 2002          | 2                        |
| uk      | some title | 1985          | 1                        |
|         | some title | 1992          | 2                        |
|         | some title | 2010          | 3                        |

rank()

- A function to say that “I want to order the rows in each partition”
- No parameters in the parentheses

order by year\_released

- In each group (partition), the rows will be ordered by the column “year\_released”

An order value is computed for each row in a partition.

- Only inside the partition, not across the entire result set

# Update

- Make changes to the existing rows in a table
- **update** is the command that changes column values
  - You can even set a non-mandatory column to NULL
  - The change is applied to all rows selected by the **where**



```
update table_name  
set column_name = new_value,  
    other_col = other_new_val,  
    ...  
where ...
```

# Delete

- As the name shows, **delete** removes rows from tables



- If you omit the WHERE clause, then (as with UPDATE) the statement **affects all rows** and you **end up with an empty table!**
- Well,
  - Many database products provide **a roll-back mechanism** when deleting rows
  - Transactions can also protect you (to some extent)

# Constraints

- One important point with constraints (esp., foreign keys) is that **they guarantee that data remains consistent**
  - They not only work with **insert**, but with **update** and **delete**
  - Example: Try to delete some rows in the country table
    - Foreign-key constraints are especially useful in controlling delete operations**



```
delete from countries where country_code = 'us';
```

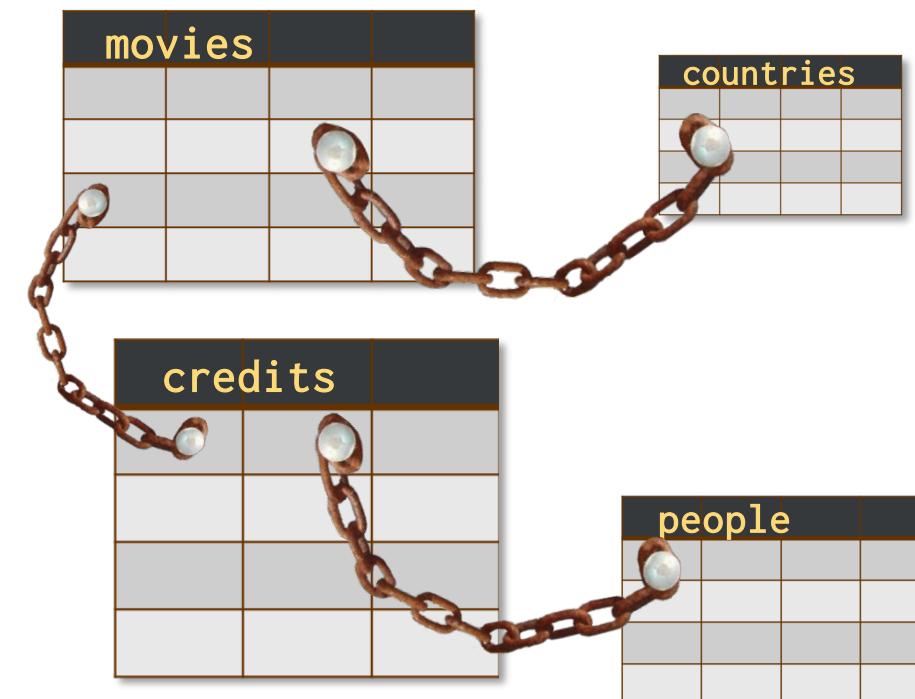
[23503] ERROR: update or delete on table "countries" violates foreign key constraint "movies\_country\_fkey" on table "movies"  
Detail: Key (country\_code)=(us) is still referenced from table "movies".

| movieid | title                 | country | year_released |
|---------|-----------------------|---------|---------------|
| 1       | Casab                 | us      | 1942          |
| 2       | Goodfellas            | us      | 1990          |
| 3       | Bronenosets Potyomkin | ru      | 1925          |

| country_code | country_name  | continent |
|--------------|---------------|-----------|
| ru           | Russia        | Europe    |
| us           | United States | America   |

# Constraints

- This is why constraints are so important:
  - They ensure that whatever happens, you'll **always be able to make sense of ALL pieces of data in your database**
  - For any changes made to the tables, all declared constraints need to be satisfied



# Self-defined Function

- Sometimes the built-in functions cannot fulfill our requirements
  - And the power of declarative language (SQL) is not strong enough
- Most DBMS implement a **built-in, SQL-based programming language**
  - A **procedural extension** to SQL

# Procedural vs. Declarative

- Two different programming paradigms
  - Imperative programming (命令式编程)
    - Describe the algorithms step-by-step (i.e., how to do)
    - Procedural (过程式) : C (and many other legacy languages)
    - Object-oriented: Java
  - Declarative programming (声明式编程)
    - Describe the result without specifying the detailed steps (i.e., what to do)
    - (Pure) declarative: SQL, Regular Expressions, Markup (HTML, XML), CSS
    - Functional: Scheme, Haskell, Scala, Erlang
    - Logic programming: Prolog

# Function in (Postgre)SQL

- How do we rewrite the name conversion expression into a function?

```
create function full_name(p_fname varchar, p_sname varchar)
returns varchar
as $$ 
begin
    return case
        when p_fname is null then ''
        else p_fname || ' '
    end || case position('('' in p_sname)
        when 0 then p_sname
        else trim(')' from substr(p_sname, position('('' in p_sname) + 1))
        || ''
        || trim(substr(p_sname, 1, position('('' in p_sname) - 1))
    end;
end;
$$ language plpgsql;
```

## Language Type

PostgreSQL supports 4 procedural languages: PL/pgSQL, PL/Tcl, PL/Perl, and PL/Python

- Tcl, Perl, and Python are famous scripting languages in case you don't know

# Function in (Postgre)SQL

- How do we rewrite the name conversion expression into a function?

```
create function full_name(p_
returns varchar
as $$
begin
    return case
        when p_fname is null
        else p_fname || ' '
    end || case position('('' in p_sname)
        when 0 then p_sname
        else trim(')' from substr(p_sname, position('('' in p_sname) + 1))
    || ''
    || trim(substr(p_sname, 1, position('('' in p_sname) - 1))
end;
end;
$$ language plpgsql
```

```
create function append_test(p_code varchar)
returns varchar
as $$
    if p_code == 'cn':
        return 'China'
    else:
        return 'not China'
$$ language plpython3u;
```

**Yes, we can even use Python to write functions**



## Language Type

PostgreSQL supports 4 procedural languages:

PL/pgSQL, PL/Tcl, PL/Perl, and PL/Python

- Tcl, Perl, and Python are famous scripting languages in case you don't know

# Functions and Procedures in (Postgre)SQL

- It follows the general definition of functions and procedures
  - **Function**: return a value
  - **Procedure**: return **NO** value
- However,
  - For some historical reasons, PostgreSQL actually has no implementation specifically for procedures
    - It shares the same mechanism with functions
    - Treats procedures as **void functions**
  - \* But for some other database systems, there are separate implementations for functions and procedures

# When to Use Procedures

- For business logics
  - One requirement may need a series of SQL querys and statements
  - Transactions may be used
  - Example: Insert a new movie into the databases
    - movies table
      - Basic information for the movie
    - countries table
      - Transformation between country names and codes
    - people table
      - new actors / directors
    - credits table
      - new credit information
  - Problem: Update all the tables? Input validation? Code reuse? Security?

# When to Use Procedures

- To add a movie:
  - We may have a series queries to execute when inserting only one movie
  - How about one call for all the processes?
    - Benefit 1: Network overhead
      - When running multiple queries, you are going to waste time chatting over the network with the remote server
    - Benefit 2: Security
      - Prevent users from modifying data otherwise than by calling carefully written and well tested procedures
      - Ensure that users can only modify data via carefully written and well tested procedures

# The Procedure



```
create function movie_registration
    (p_title      varchar,
     p_country_name varchar,
     p_year       int,
     p_director_fn  varchar,
     p_director_sn  varchar,
     p_actor1_fn   varchar,
     p_actor1_sn   varchar,
     p_actor2_fn   varchar,
     p_actor2_sn   varchar)
returns void
as $$

declare
    n_rowcount int;
    n_movieid int;
    n_people int;
begin
    insert into movies(title, country, year_released)
        select p_title, country_code, p_year
        from countries
        where country_name = p_country_name;
    get diagnostics n_rowcount = row_count;

    if n_rowcount = 0
    then
        raise exception 'country not found in table COUNTRIES';
    end if;
```

```
n_movieid := lastval();
select count(surname)
into n_people
from (select p_director_sn as surname
      union all
      select p_actor1_sn as surname
      union all
      select p_actor2_sn as surname) specified_people
where surname is not null;

insert into credits(movieid, peopleid, credited_as)
select n_movieid, people.peopleid, provided.credited_as
from (select coalesce(p_director_fn, '*') as first_name,
            p_director_sn as surname,
            'D' as credited_as
         union all
         select coalesce(p_actor1_fn, '*') as first_name,
                p_actor1_sn as surname,
                'A' as credited_as
         union all
         select coalesce(p_actor2_fn, '*') as first_name,
                p_actor2_sn as surname,
                'A' as credited_as) provided
inner join people
on people.surname = provided.surname
and coalesce(people.first_name, '*') = provided.first_name
where provided.surname is not null;

get diagnostics n_rowcount = row_count;
if n_rowcount != n_people
then
    raise exception 'Some people couldn''t be found';
end if;
end;
$$ language plpgsql;
```

Trigger (触发器)

# Trigger (触发器) - Actions When Changing Tables

A **trigger** is a specification that the database should automatically execute a particular function whenever a certain type of operation is performed.

-- Chapter 39, PostgreSQL Documentation

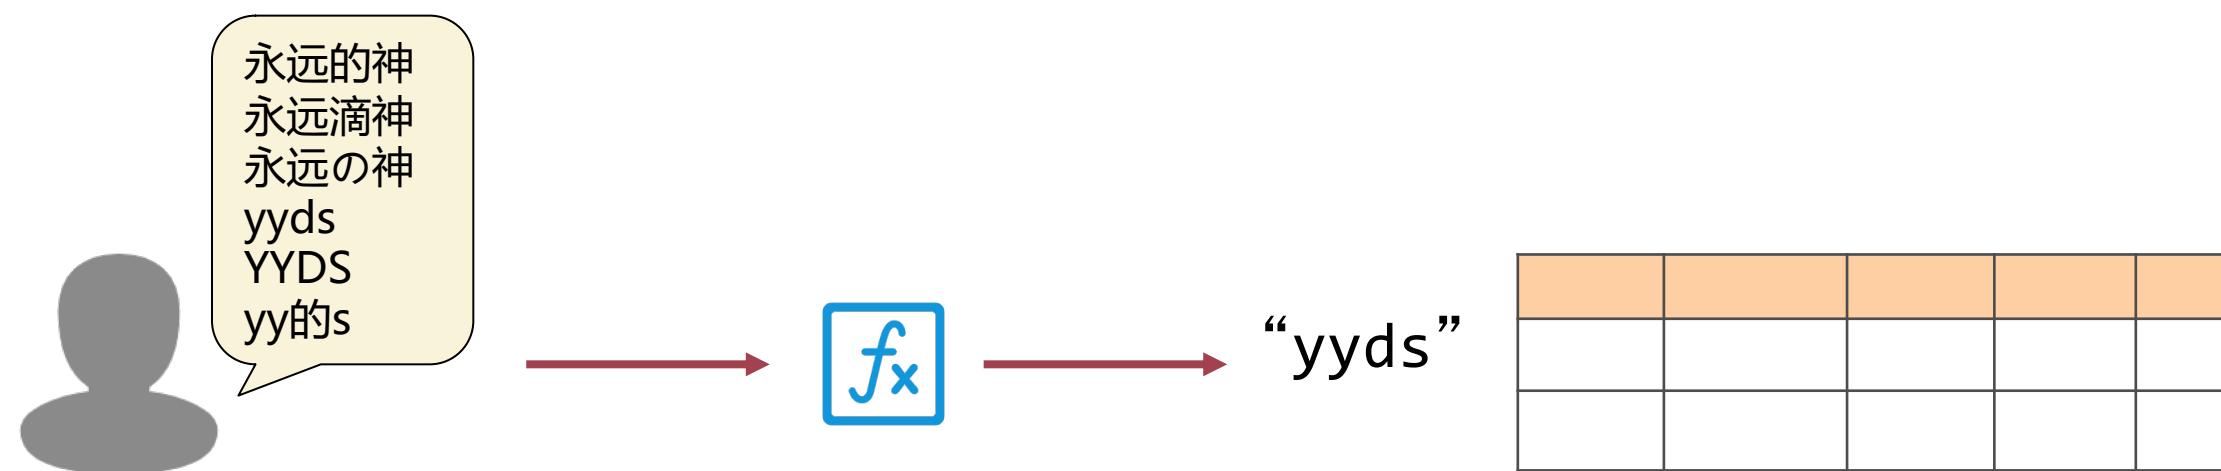
A **trigger** is a statement that the system executes automatically as a side effect of a modification to the database.

-- Chapter 5.3, Database System Concepts, 7th

- We can attach “actions” to a table
  - They will be **executed automatically whenever the data in the table changes**
- Purpose of using triggers
  - Validating data
  - Checking complex rules
  - Managing data redundancy

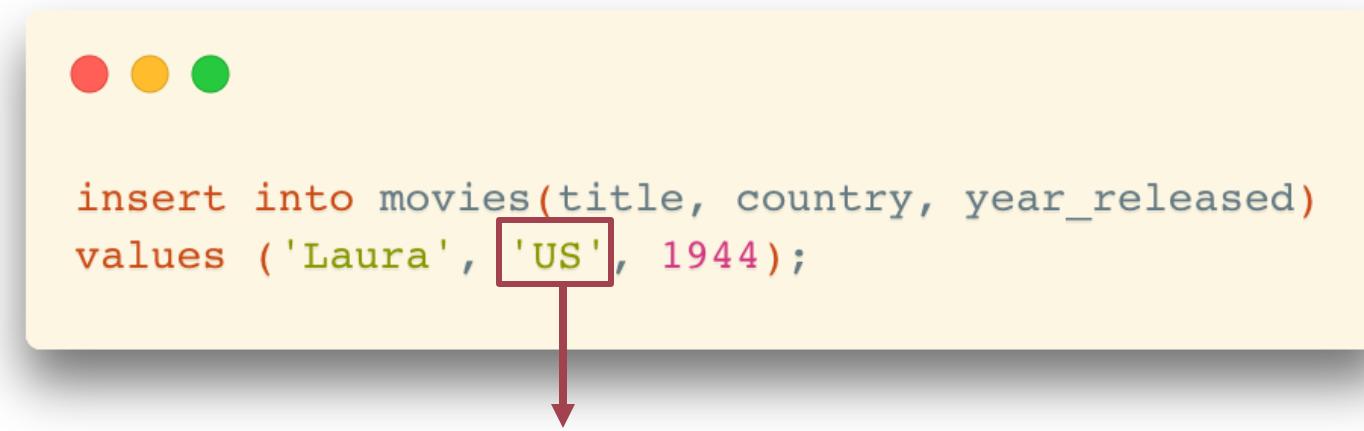
# Purpose of Using Triggers

- Validating data
  - Some data are badly formatted in programs before sending to the database
  - We need to validate such data before inserting them into the database
- “On-the-fly” modification
  - Change the input directly when the input arrives



# Purpose of Using Triggers

- Validating data
  - Example: insert a row in the movies table
    - In the JDBC program, an `insert` request is written like the following:



```
insert into movies(title, country, year_released)
values ('Laura', 'US', 1944);
```

Need to update it to 'us'  
before inserting

- Although,
  - Such validation or transformation should be better handled by the application programs

# Purpose of Using Triggers

- Checking complex rules
  - Sometimes, the business rules are so complicated
    - They **CANNOT** be checked via declarative integrity constraints

# Purpose of Using Triggers

- Managing data redundancy
  - Some data redundancy issues cannot be avoided by simply adding constraints
  - For example: inserted the same movie but in different languages



```
-- US
insert into movies(title, country, year_released)
values ('The Matrix', 'us', 1999);

-- China (Mainland)
insert into movies(title, country, year_released)
values ('黑客帝国', 'us', 1999);

-- Hongkong
insert into movies(title, country, year_released)
values ('22世紀殺人網絡', 'us', 1999);
```

It satisfies the unique constraint on (title, country, year\_released)  
• ... but they represent the same movie

# Trigger Activation

- Two key points:
  - When to fire a trigger?
  - What (command) fires a trigger?

# Trigger Activation

- When to fire a trigger?
  - In general: “During the change of data”
    - ... but we need a detailed discussion
  - Note: “During the change” means select queries won’t fire a trigger.

# Trigger Activation: When

- Example: Insert a set of rows with “insert into select”
  - One statement, multiple rows



```
insert into movies(title, country, year_released)
select titre, 'fr', annee
from films_francais;
```

- Option 1: Fire a trigger only once for the statement
  - Before the first row is inserted, or after the last row is inserted
- Option 2: Fire a trigger for each row
  - Before or after the row is inserted

# Trigger Activation: When

- Different options between DBMS products

PostgreSQL



ORACLE®



MySQL®



Microsoft®  
SQL Server®

- Before statement
  - Before each row
  - After each row
- After statement

- Before statement
  - Before each row
  - After each row
- After statement

- Before statement
  - Before each row
  - After each row
- After statement

# Trigger Activation: What

- What (command) fires a trigger?
  - insert
  - update
  - delete



# Before and After Triggers

- Differences between before and after triggers
  - “Before” and “after” the operation is done (insert, update, delete)
  - If we want to update the incoming values in an insert statement (e.g., “US” → “us”), “before trigger” should be used since the incoming values have not been written to the table yet

```
insert into movies(title, country, year_released)
values ('Laura', 'US', 1944);
```

Need to update it to 'us'  
before inserting

# Before and After Triggers

- Typical usage scenarios for trigger settings
  - Modify input on the fly
    - before insert / update
    - for each row
  - Check complex rules
    - before insert / update / delete
    - for each row
  - Manage data redundancy
    - after insert / update / delete
    - for each row

# Example of Triggers

- An Example
  - For the `people_1` table, when updating a person, count the number of movies and save the result in the `num_movies` column

```
● ● ●  
-- auto-generated definition  
create table people_1  
(  
    peopleid integer,  
    first_name varchar(30),  
    surname varchar(30),  
    born integer,  
    died integer,  
    gender bpchar,  
    num_movies integer  
);
```

|    | peopleid | first_name | surname  | born | died   | gender | num_movies |
|----|----------|------------|----------|------|--------|--------|------------|
| 1  | 13       | Hiam       | Abbass   | 1960 | <null> | F      | <null>     |
| 2  | 559      | Aleksandr  | Askoldov | 1932 | <null> | M      | <null>     |
| 3  | 572      | John       | Astin    | 1930 | <null> | M      | <null>     |
| 4  | 585      | Essence    | Atkins   | 1972 | <null> | F      | <null>     |
| 5  | 598      | Antonella  | Attili   | 1963 | <null> | F      | <null>     |
| 6  | 611      | Stéphane   | Audran   | 1932 | <null> | F      | <null>     |
| 7  | 624      | William    | Austin   | 1884 | 1975   | M      | <null>     |
| 8  | 637      | Tex        | Avery    | 1908 | 1980   | M      | <null>     |
| 9  | 650      | Dan        | Aykroyd  | 1952 | <null> | M      | <null>     |
| 10 | 520      | Zackary    | Arthur   | 2006 | <null> | M      | <null>     |
| 11 | 533      | Oscar      | Asche    | 1871 | 1936   | M      | <null>     |
| 12 | 546      | Elizabeth  | Ashley   | 1939 | <null> | F      | <null>     |

# Example of Triggers

- Create a trigger



```
create trigger test_trigger  
before update  
on people_1  
for each row  
execute procedure fill_in_num_movies();
```

# Example of Triggers

- Create a trigger

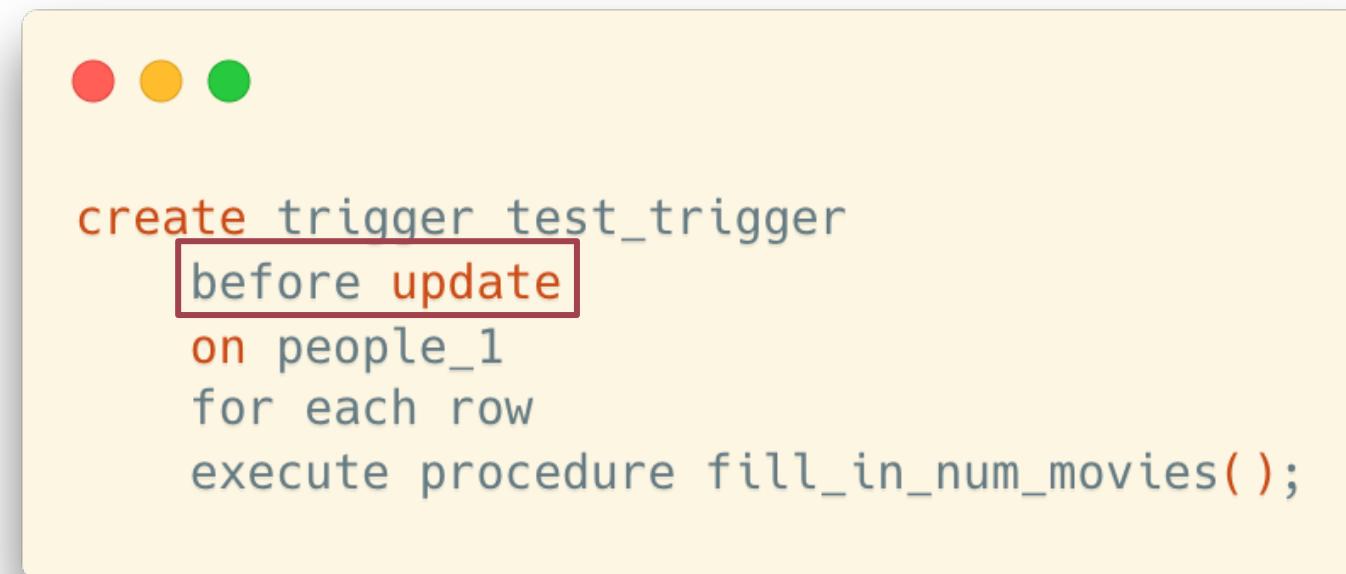


Name of the trigger

```
create trigger test_trigger
before update
on people_1
for each row
execute procedure fill_in_num_movies();
```

# Example of Triggers

- Create a trigger



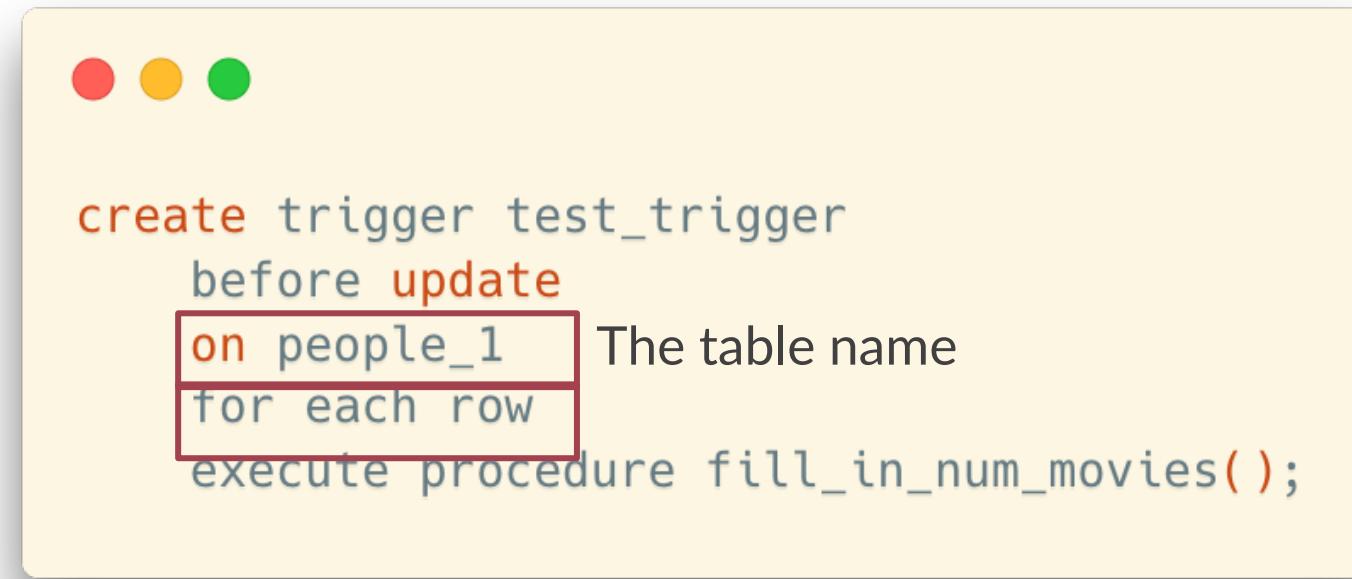
```
create trigger test_trigger
before update
on people_1
for each row
execute procedure fill_in_num_movies();
```

{ BEFORE | AFTER | INSTEAD OF } { event [ OR ... ] }

- Specify when the trigger will be executed
  - before | after
  - ... and on what operations the trigger will be executed
    - insert [or update [or delete]]

# Example of Triggers

- Create a trigger



```
create trigger test_trigger
before update
on people_1 The table name
for each row
execute procedure fill_in_num_movies();
```

“for each row”

or

“for each statement” (default)

# Example of Triggers

- Create a trigger

```
create trigger test_trigger
  before update
  on people_1
  for each row
    execute procedure fill_in_num_movies();
```

The actual procedure for the trigger

# Example of Triggers

- Create a trigger
  - Besides, a corresponding procedure should be created as well



```
create or replace function fill_in_num_movies()
    returns trigger
as
$$
begin
    select count(distinct c.movieid)
    into new.num_movies
    from credits c
    where c.peopleid = new.peopleid;
    return new;
end;
$$ language plpgsql;
```

# Example of Triggers

- Create a trigger
  - Besides, a corresponding procedure should be created as well

```
create or replace function fill_in_num_movies()
    returns trigger "trigger" is the return type
as
$$
begin
    select count(distinct c.movieid)
    into new.num_movies
    from credits c
    where c.peopleid = new.peopleid;
    return new;
end;
$$ language plpgsql;
```

# Example of Triggers

- Create a trigger
  - Besides, a corresponding procedure should be created as well

```
create or replace function fill_in_num_movies()
  returns trigger
as
$$
begin
  select count(distinct c.movieid)
  into new.num_movies
  from credits c
  where c.peopleid = new.peopleid;
  return new;
end;
$$ language plpgsql;
```

“new” and “old” are two internal variables that represents the row before and after the changes

# Example of Triggers

- Create a trigger
  - Besides, a corresponding procedure should be created as well

```
create or replace function fill_in_num_movies()
  returns trigger
as
$$
begin
  select count(distinct c.movieid)
  into new.num_movies
  from credits c
  where c.peopleid = new.peopleid;
  return new;
end;
$$ language plpgsql;
```

Remember to return the result which will be used in the **update** statement

# Example of Triggers

- Create a trigger
  - Besides, a corresponding procedure should be created as well
    - Remember to [create the procedure before creating the trigger](#)
- Run test updates



```
-- create the procedure fill_in_num_movies() first  
-- then, create the trigger  
-- finally, we can run some test update statements  
update people_1 set num_movies = 0 where people_1.peopleid <= 100;
```

# Example: Auditing

- Another example of trigger is **keeping an audit trail**
  - **Not really care about people who read data**
    - Remember that select cannot fire a trigger, although with the big products you can trace all queries
  - ... but it may be useful for checking people who modify data that they aren't supposed to modify

# Example: Auditing

- Trace the insertions and updates to employees in a company

```
create table company(
    id int primary key      not null,
    name          text      not null,
    age           int       not null,
    address        char(50),
    salary         real
);
```

```
create table audit(
    emp_id int not null,
    change_type char(1) not null,
    change_date text not null
);
```

# Example: Auditing

- Trace the insertions and updates to employees in a company



```
create trigger audit_trigger
    after insert or update
    on company
    for each row
execute procedure auditlogfunc();
```

```
create or replace function auditlogfunc() returns trigger as
$example_table$
begin
    insert into audit(emp_id, change_type, change_date)
    values (new.id,
            case
                when tg_op = 'UPDATE' then 'U'
                when tg_op = 'INSERT' then 'I'
                else 'X'
            end,
            current_timestamp);
    return new;
end ;
$example_table$ language plpgsql;
```

# Example: Auditing

- Trace the insertions and updates to employees in a company



```
insert into company (id, name, age, address, salary)
values (2, 'Mike', 35, 'Arizona', 30000.00);
```

company

|  | id | name   | age | address | salary |
|--|----|--------|-----|---------|--------|
|  | 1  | 2 Mike | 35  | Arizona | 30000  |

audit

|  | emp_id | change_type | change_date                   |
|--|--------|-------------|-------------------------------|
|  | 1      | 2 I         | 2022-04-25 18:37:35.515151+00 |

**View**

# Recall: Function

- Used for returning numbers, strings, dates, etc.
  - Or, return a single value



```
select full_name(first_name, surname)  
from people  
where surname like '%(von)';
```

... returns a string of the full name

# View

- Reuse the relational operations, like how we reuse some code in a program
  - ... which **returns relations (tables)** instead of simple values



```
create view viewname as  
select ...  
from ...  
where ...
```

# View

- Example: Create a view
  - ... where the result of the inner select query looks like:



```
create view vmovies as
select m.movieid,
       m.title,
       m.year_released,
       c.country_name
  from movies m join countries c
    on c.country_code = m.country;
```

|    | movieid | title                           | year_released | country_name  |
|----|---------|---------------------------------|---------------|---------------|
| 1  | 1       | 12 stulyev                      | 1971          | Russia        |
| 2  | 2       | Al-mummia                       | 1969          | Egypt         |
| 3  | 3       | Ali Zaoua, prince de la rue     | 2000          | Morocco       |
| 4  | 4       | Apariencias                     | 2000          | Argentina     |
| 5  | 5       | Ardh Satya                      | 1983          | India         |
| 6  | 6       | Armaan                          | 2003          | India         |
| 7  | 7       | Armaan                          | 1966          | Pakistan      |
| 8  | 8       | Babette's gæstebud              | 1987          | Denmark       |
| 9  | 9       | Banshun                         | 1949          | Japan         |
| 10 | 10      | Bidaya wa Nihaya                | 1960          | Egypt         |
| 11 | 11      | Variety                         | 2008          | United States |
| 12 | 12      | Bon Cop, Bad Cop                | 2006          | Canada        |
| 13 | 13      | Brilliantovaja ruka             | 1969          | Russia        |
| 14 | 14      | C'est arrivé près de chez vous  | 1992          | Belgium       |
| 15 | 15      | Carlota Joaquina - Princesa d.. | 1995          | Brazil        |
| 16 | 16      | Cicak-man                       | 2006          | Malaysia      |
| 17 | 17      | Da Nao Tian Gong                | 1965          | China         |
| 18 | 18      | Das indische Grabmal            | 1959          | Germany       |
| 19 | 19      | Das Leben der Anderen           | 2006          | Germany       |
| 20 | 20      | Den store gavtyv                | 1956          | Denmark       |

# View vs. Table

- In practice, there isn't much to a view
  - It's basically a named query

```
create view vmovies(v_id, v_title, v_year, v_country) as
select m.movieid,
       m.title,
       m.year_released,
       c.country_name
  from movies m
    join countries c
      on c.country_code = m.country;
```

# View vs. Table

- In practice, there isn't much to a view
  - It's basically a named query

The columns can be renamed

- Otherwise, the original names in the tables will be used

```
create view vmovies(v_id, v_title, v_year, v_country) as
select m.movieid,
       m.title,
       m.year_released,
       c.country_name
  from movies m
    join countries c
      on c.country_code = m.country;
```

# View vs. Table

- In practice, there isn't much to a view
  - It's basically a named query



The columns can be renamed

- Otherwise, the original names in the tables will be used

```
create view vmovies(v_id, v_title, v_year, v_country) as  
select m.movieid,  
       m.title,  
       m.year_released,  
       c.country_name  
  from movies m  
       join countries c  
         on c.country_code = m.country;
```

Drop the existing view before creating a new one with the same name:



```
drop view vmovies;
```

# View vs. Table

- Once the view is created,
  - ... we can **query the view exactly as if it were a table**

```
create view vmovies as  
select m.movieid,  
       m.title,  
       m.year_released,  
       c.country_name  
  from movies m join countries c  
    on c.country_code = m.country;
```

Use the view name in the **from** clause

- ... as if it were a “table”

```
select *  
  from vmovies  
 where country_name = 'Italy';
```

# View vs. Table

- A view can be as a complicated query as you want
  - It will usually return something that isn't as normalized as your tables, but easier to understand
- View is like a function of a select query
  - The view relation is **not** precomputed and stored
  - The DBMS stores the **query expression** associated with the view relation
  - The view is computed by executing the query, whenever the virtual relation is used
    - Virtual relation is created whenever needed, on demand

# Content in Views is Dynamic

- When the rows change in the tables where a view relies on, the result of a view will change as well
  - Think it like a query result, or a table variable
  - \* Or, think it like a relational function which returns a relation
  - A function of a select query
    - The view is computed by executing the query, whenever the virtual relation is used

# Content in Views is Dynamic

- In a view, the **columns** are the existing ones in tables when the view was created
  - Columns added later to tables **won't be added** in the view, even if the view was created with “select \*”
    - And, “select \*” is a bad practice in view. **We will not know what exact columns are selected** once the columns in the tables are changed.

```
● ● ●

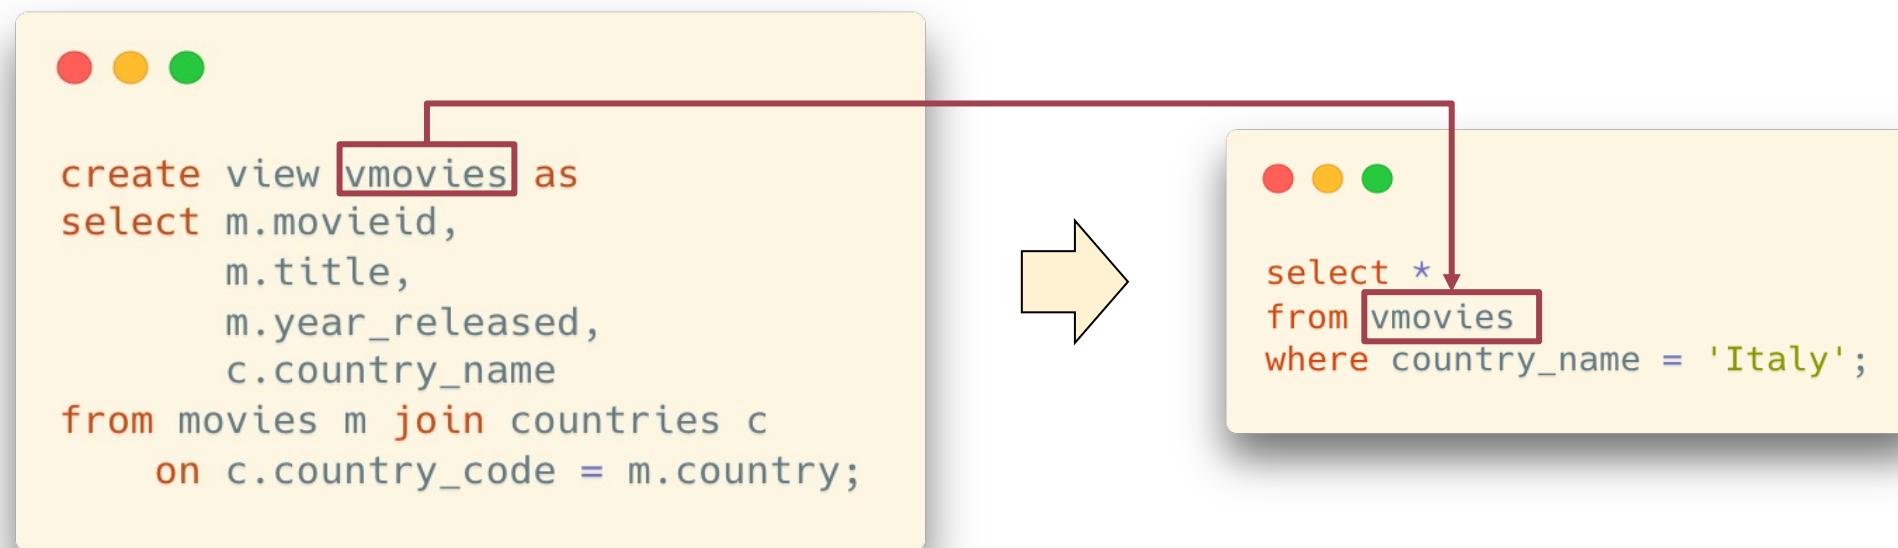
-- Create the view first
create view vmovies as
select * -- IT IS A BAD PRACTICE OF USING * HERE
from movies m join countries c
  on c.country_code = m.country;

-- Alter the table then
alter table movies add column test_col int not null default 0;

-- There won't be a column "test_col" in the query result
select * from vmovies;
```

# View vs. Table (Cont.)

- View is like Table
  - Sometimes it is used as a table
  - Usage Scenario 1: Simplify Complex Queries
    - Simplify a complicated query result into a single named query
    - e.g., Many business reports are based on the same set of joins, with just variations on the columns that you aggregate or order by



# View vs. Table (Cont.)

- View is like Table
  - Sometimes it is used as a table
  - Usage Scenario 2: An alternative way to implement E-R models
    - Sometimes, we may not be able to use tables to model entities and relationships
      - Access control
      - Dirty and messy original data
    - Since views look like tables, we can use views to represent entities or relationships
      - ... based on some existing objects (like tables)

# Drawback of Views

- View is like Table
  - But it is still not a table
- In some cases, it may cause performance issues or unnecessary operations
  - \* ... since the users of the views don't know the details of the views
- A function of a select query
  - The view is computed by executing the query, whenever the virtual relation is used

# Drawback of Views

- View looks like Table, tastes like Table
  - Example: A refined way to display the credit information with a view

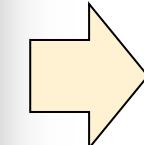
```
create view vmovie_credits
as
select m.title,
       m.year_released release,
       case c.credited_as
           when 'A' then 'Actor'
           when 'D' then 'Director'
           else '?'
           end duty,
       full_name(p.first_name, p.surname) name
from movies m
     inner join credits c
             on c.movieid = m.movieid
     inner join people p
             on p.peopleid = c.peopleid;
```

| #  | title                                    | release | duty     | name               |
|----|--|---------|----------|--------------------|
| 1  | "Erogotoshitachi" yori Jinruigaku nyūmon | 1966    | Director | Shohei Imamura     |
| 2  | "Erogotoshitachi" yori Jinruigaku nyūmon | 1966    | Actor    | Shoichi Ozawa      |
| 3  | "Erogotoshitachi" yori Jinruigaku nyūmon | 1966    | Actor    | Sumiko Sakamoto    |
| 4  | '76                                      | 2016    | Actor    | Ramsey Nouah       |
| 5  | '76                                      | 2016    | Actor    | Ibinabo Fiberesima |
| 6  | (T)Raumschiff Surprise                   | 2004    | Actor    | Michael Herbig     |
| 7  | (T)Raumschiff Surprise                   | 2004    | Director | Michael Herbig     |
| 8  | (T)Raumschiff Surprise                   | 2004    | Actor    | Til Schweiger      |
| 9  | (T)Raumschiff Surprise                   | 2004    | Actor    | Anja Kling         |
| 10 | (T)Raumschiff Surprise                   | 2004    | Actor    | Rick Kavanian      |
| 11 | (T)Raumschiff Surprise                   | 2004    | Actor    | Christian Tramitz  |
| 12 | (T)Raumschiff Surprise                   | 2004    | Actor    | Sky du Mont        |
| 13 | 002 operazione luna                      | 1965    | Actor    | Linda Sini         |
| 14 | 002 operazione luna                      | 1965    | Director | Lucio Fulci        |

# Drawback of Views

- View looks like Table, tastes like Table
  - Example: A refined way to display the credit information with a view
    - However, if we only want to get all distinct movie titles, it will return the correct result, but there are a lot of useless works in the internal query of the view

```
● ● ●  
select distinct title  
from vmovie_credits
```



```
● ● ●  
  
create view vmovie_credits  
as  
select m.title,  
       m.year_released release,  
       case c.credited_as  
         when 'A' then 'Actor'  
         when 'D' then 'Director'  
         else '?'  
       end duty,  
       full_name(p.first_name, p.surname) name  
     from movies m  
       inner join credits c  
         on c.movieid = m.movieid  
       inner join people p  
         on p.peopleid = c.peopleid;
```

Do we really need the joins?

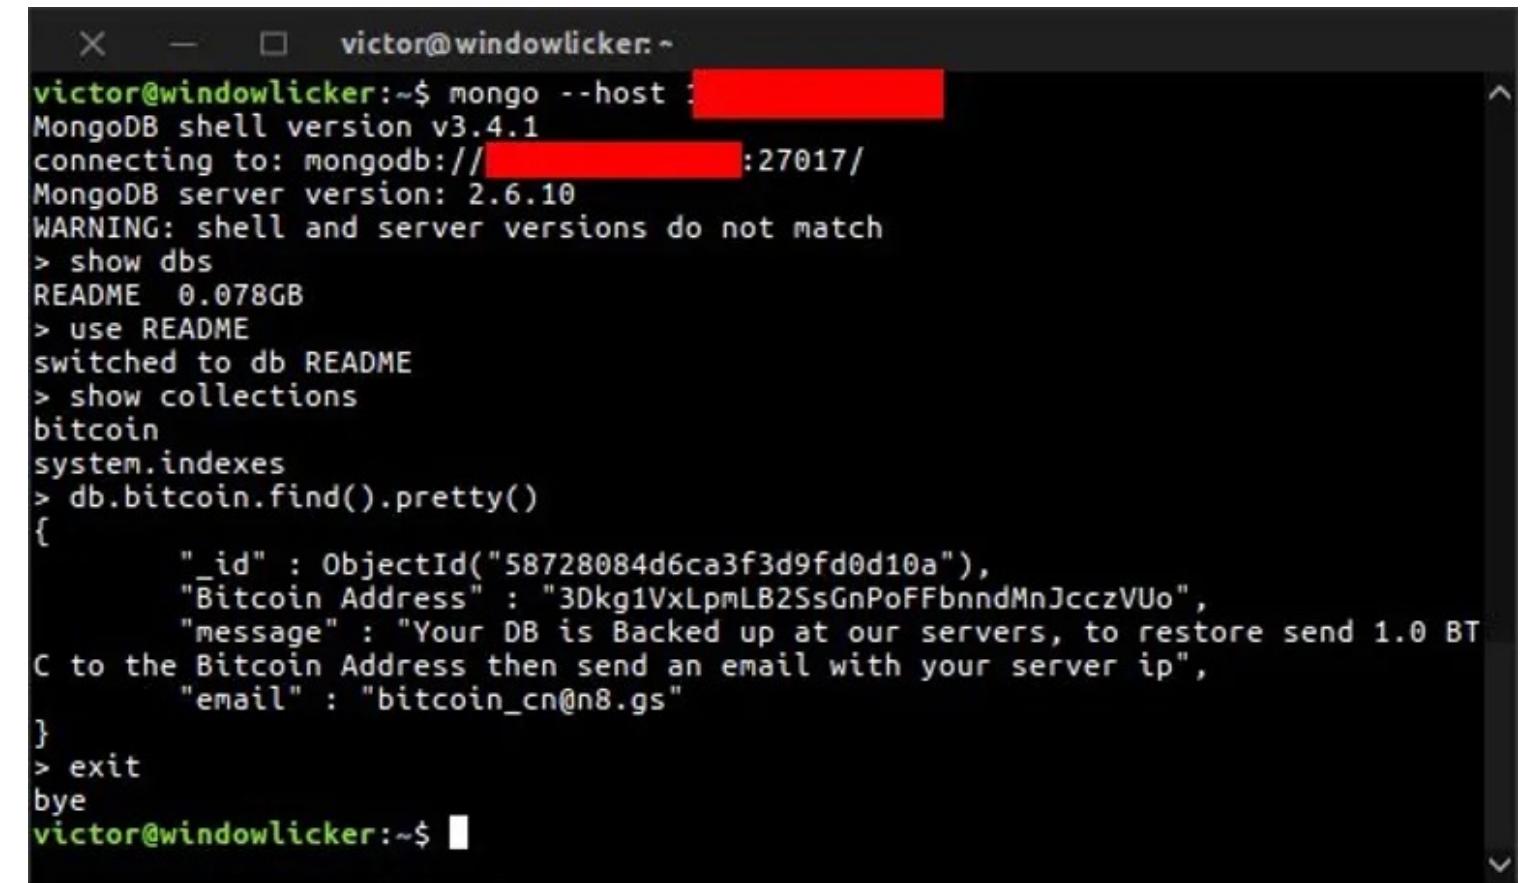
# View vs. Materialized View

- Normal views dynamically retrieve data from the original tables
  - However, the performance won't be good
- **Materialized View:** A cached result of a view
  - When the corresponding tables are not changing rapidly, we can **cache the view results** (i.e., materialize the results)
    - Better performance than dynamic retrieval
  - **The results are not updated each time the view is used**
    - Manual update, or use the help of triggers

# Access Control

# Authentication

- To access a database, you must be authenticated
  - Entering a username and a password.
- If you don't set a proper password, you will put your database system in danger



```
victor@windowlicker:~$ mongo --host :  
MongoDB shell version v3.4.1  
connecting to: mongodb://:27017/  
MongoDB server version: 2.6.10  
WARNING: shell and server versions do not match  
> show dbs  
README 0.078GB  
> use README  
switched to db README  
> show collections  
bitcoin  
system.indexes  
> db.bitcoin.find().pretty()  
{  
    "_id" : ObjectId("58728084d6ca3f3d9fd0d10a"),  
    "Bitcoin Address" : "3Dkg1VxLpmLB2SsGnP0FFbnndMnJcczVUo",  
    "message" : "Your DB is Backed up at our servers, to restore send 1.0 BT  
C to the Bitcoin Address then send an email with your server ip",  
    "email" : "bitcoin_cn@n8.gs"  
}  
> exit  
bye  
victor@windowlicker:~$
```

<https://www.zdnet.com/article/mongodb-ransacked-now-27000-databases-hit-in-mass-ransom-attacks/>

# Privileges

- Besides, DBMS usually provide another layer of access control on objects in the system
  - Operations (select, update, insert, delete, etc.)
  - Objects (table, database, views, trigger, etc.)

# Grant and Revoke Access

- grant and revoke
  - Can be used on the table or the database level



```
-- grant <right> to <account>
grant select on movies to test_user;

-- revoke <right> from <account>
revoke select on movies from test_user;
```

The possible privileges are:

SELECT  
INSERT  
UPDATE  
DELETE  
TRUNCATE  
REFERENCES  
TRIGGER  
CREATE  
CONNECT  
TEMPORARY  
EXECUTE  
USAGE

# How Can Views Help

- User access control
  - By creating a view that only returns rows associated with the user name, we can control the privileges for a specific user in this table



```
create view my_stuff  
as  
select * from stuff  
where username = user
```

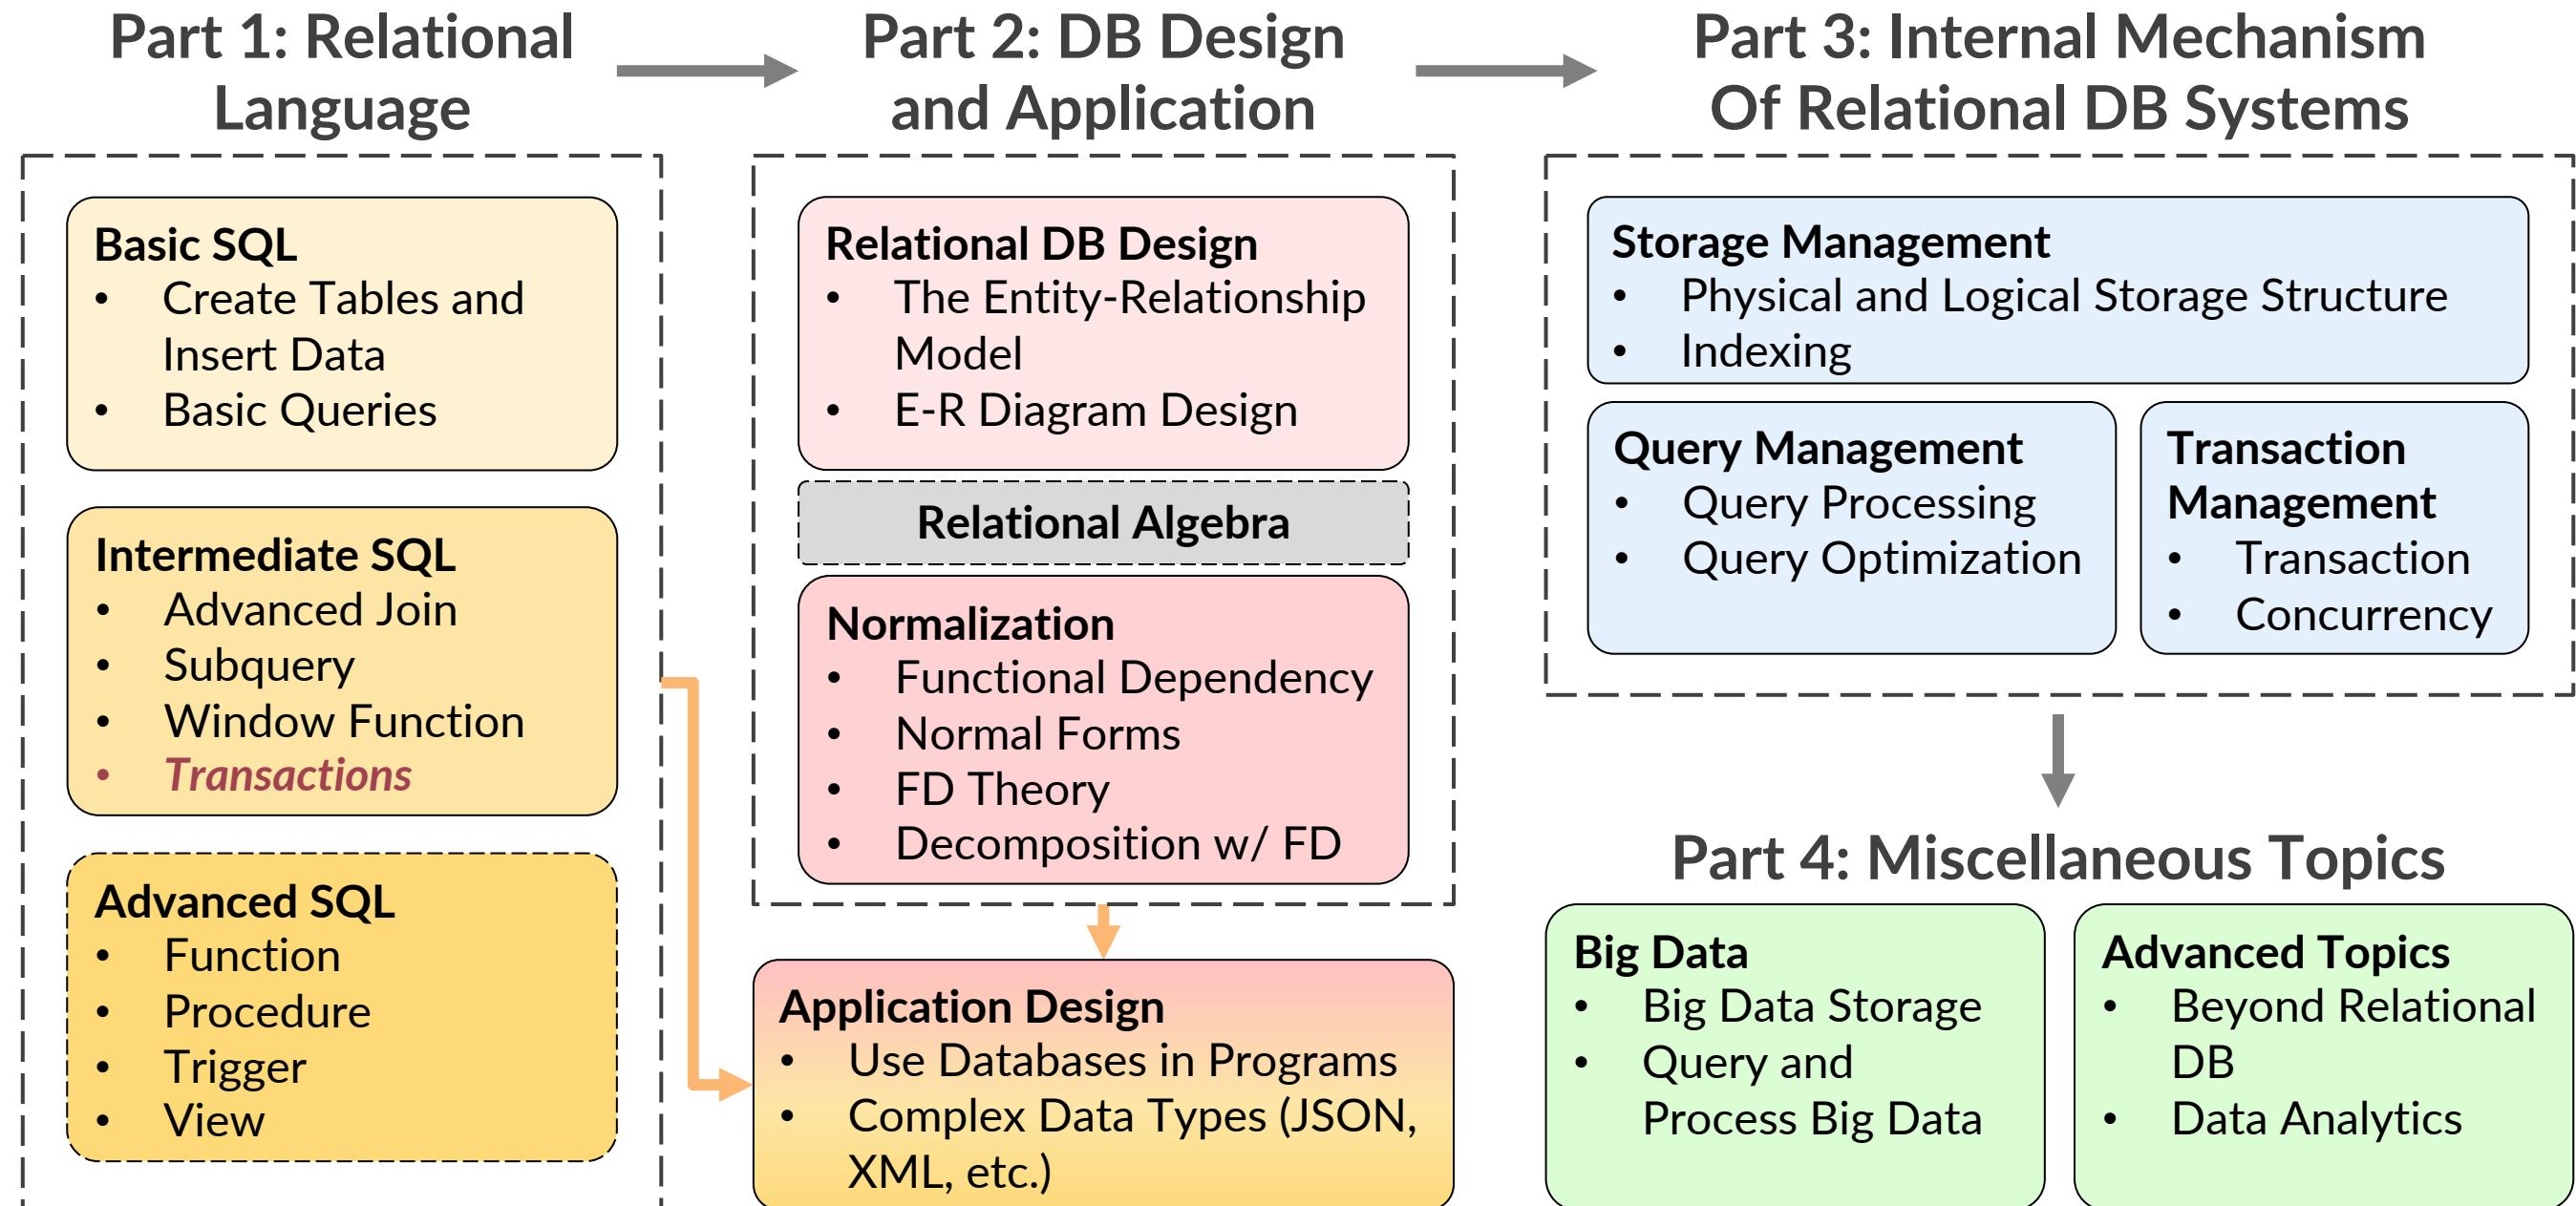


-- E.g., restrict the users to only access  
-- the movies assigned to their user names

```
create view user_view as  
select movieid,  
       title,  
       country,  
       year_released,  
       runtime  
     from movies  
   where user_name = user;
```

user: an internal variable that returns the current user name

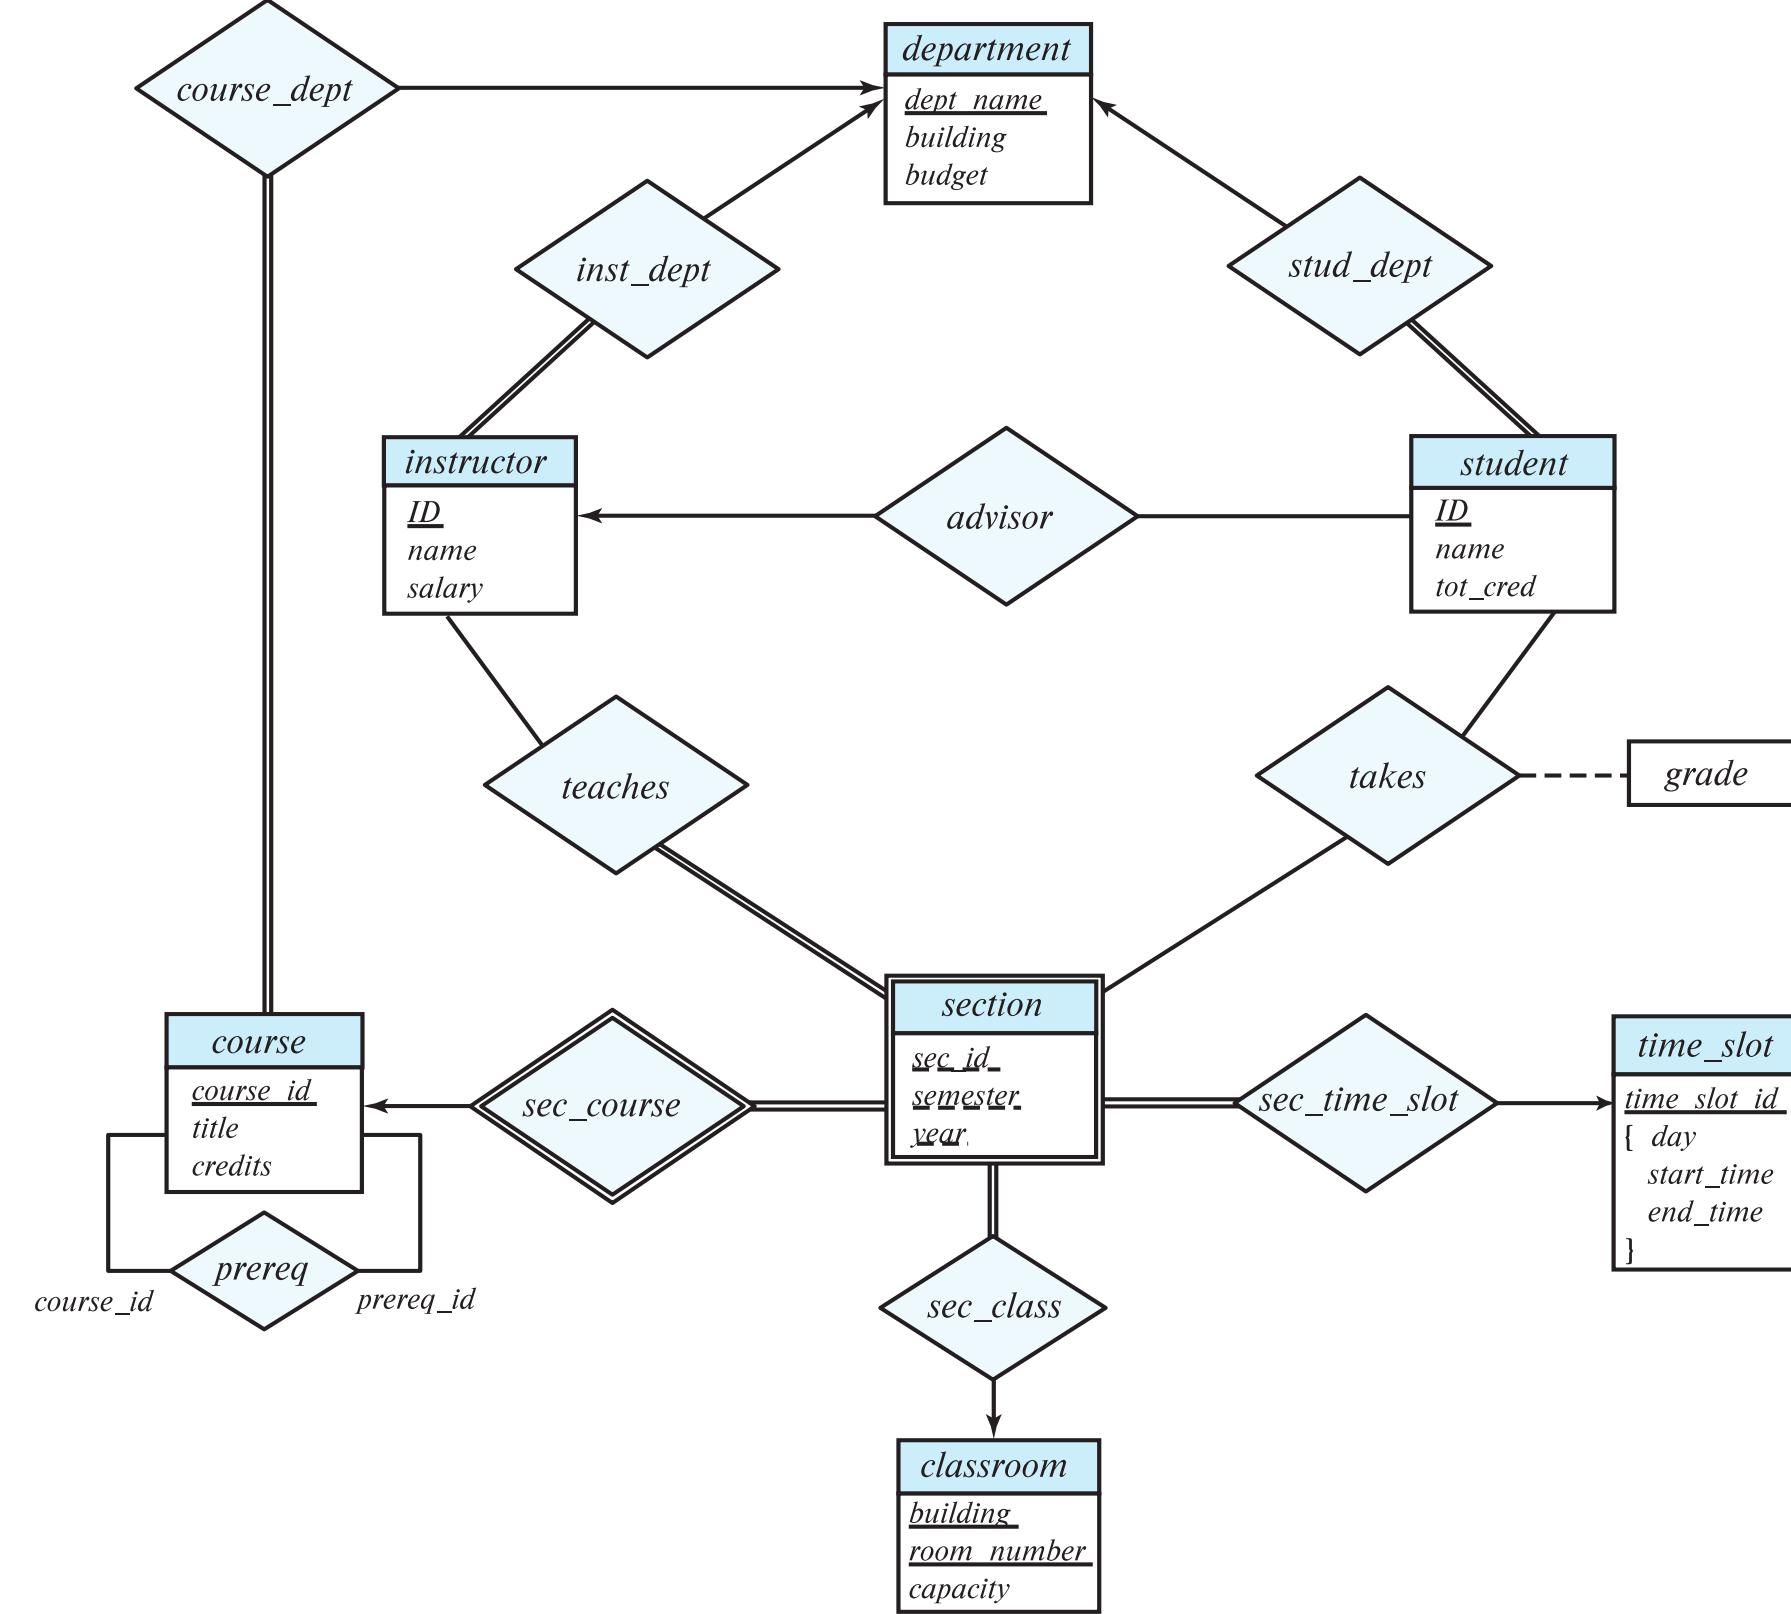
# Outline



# **Entity-Relationship Model (E-R Model)**

# **Entity-Relationship Diagram (E-R Diagram)**

# The New Running Example



# Design Phases

- Initial phase: characterize fully the data needs of the prospective database users
  - Interact extensively with domain experts
  - The outcome is a specification of user requirements

# Design Phases

- Second phase: choosing a data model
  - E.g., relational model, entity-relationship model, semi-structured data model, and object-oriented data model
  - Applying the concepts of the chosen data model
  - Translating these requirements into a **conceptual schema of the database**
    - Detailed overview of the enterprise
  - A fully developed conceptual schema indicates the functional requirements of the enterprise
    - Describes operations (e.g., update, retrieval, delete) that will be performed on the data
- Entity-relationship model is typically used
  - Outcome is an E-R diagram that provides a graphic representation of the schema
    - Entities that are represented in the database
    - Attributes of the entities
    - Relationships among entities
    - Constraints on the entities and relationship

# Design Phases

- Final Phase: Moving from an abstract data model to the implementation of the database
  - Logical Design – Deciding on the database schema
    - Database design requires that we find a “good” collection of relation schemas
    - Business decision
      - What attributes should we record in the database?
    - Computer Science decision
      - What relation schemas should we have and how should the attributes be distributed among the various relation schemas?
  - Physical Design – Deciding on the physical layout of the database
    - E.g., the form of file organization and choice of index structures
  - Physical schema is easy to change after application is built
    - But logical schema is not, because changes may affect a number of queries and updates scattered across application code

# Design Alternatives

- In designing a database schema, we must ensure that **we avoid two major pitfalls**
  - **Redundancy:** a bad design may result in repeated information
    - E.g., store course identifier and title of a course for each course offering
      - Only store course identifier is sufficient
    - Redundant representation of information may **lead to data inconsistency among the various copies of information**
      - E.g., update is not performed on all the copies
  - **Incompleteness:** a bad design may make certain aspects of the enterprise difficult or impossible to model
    - E.g., only have entity for course offering, but without entity for courses
      - Impossible to model new courses that are not offered yet

# Design Alternatives

- Avoiding bad designs is not enough
  - There may be many good designs from which we must choose
- For example, a customer who buys a product
  - The sale is a relationship between the customer and the product?
  - The sale is a relationship among the customer, the product, and the sale itself?
    - i.e., the sale can be considered as an entity
- Database design can be difficult
  - When #entities and #relationships are large

# Design Approaches

- Entity-Relationship Model (covered in this chapter)
  - Specifies an enterprise schema representing overall logical structure of a database
  - Models an enterprise as a collection of entities and relationships
    - **Entity**: a “thing” or “object” in the enterprise that is distinguishable from other objects
      - Described by a set of attributes
    - **Relationship**: an association among several entities
  - Represented diagrammatically by an **entity-relationship diagram (E-R diagram)**
    - Express overall logical structure of a database graphically
- Normalization Theory (coming in the next few weeks)
  - Formalize what designs are bad, and test for them

# Entity and Entity Sets

- An **entity** is **an object** that exists and is distinguishable from other objects
  - Concrete entity: specific person, company, plant, book
  - Abstract entity: flight reservation, course, course offering
  - An entity is represented by a set of attributes; i.e., descriptive properties possessed by all members of an entity set
  - Example:

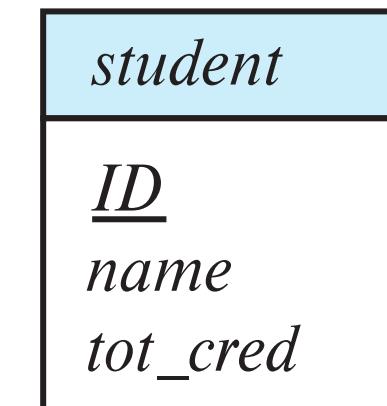
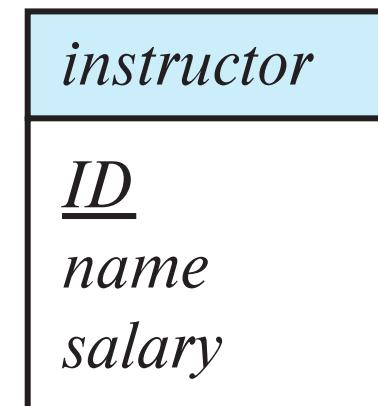
```
instructor = (ID, name, salary)
course = (course_id, title, credits)
```
- A subset of the attributes form a **primary key** of the entity set; i.e., uniquely identifying each member of the set
  - Government-issued ID number as the primary key
    - May have privacy and security issues
  - Enterprise-issued ID number

# Entity and Entity Sets

- An **entity set** is a set of entities of the same type that share the same properties
  - Example: set of all persons, companies, trees, holidays
  - Entity sets may not be disjoint (e.g., person vs. instructor and student)

# Representing Entity sets in ER Diagram

- Entity sets can be represented graphically as follows:
  - Rectangles represent entity sets.
  - Attributes listed inside entity rectangle
  - Underline indicates primary key attributes



# Relationship Sets

- A relationship is an association among several entities  
44553 (Peltier)      advisor                          22222 (Einstein)  
student entity      relationship set                      instructor entity
- A relationship set is a set of relationships of the same type (e.g., advising)
  - A mathematical relation among  $n \geq 2$  (possibly non-distinct) entities, each taken from entity sets

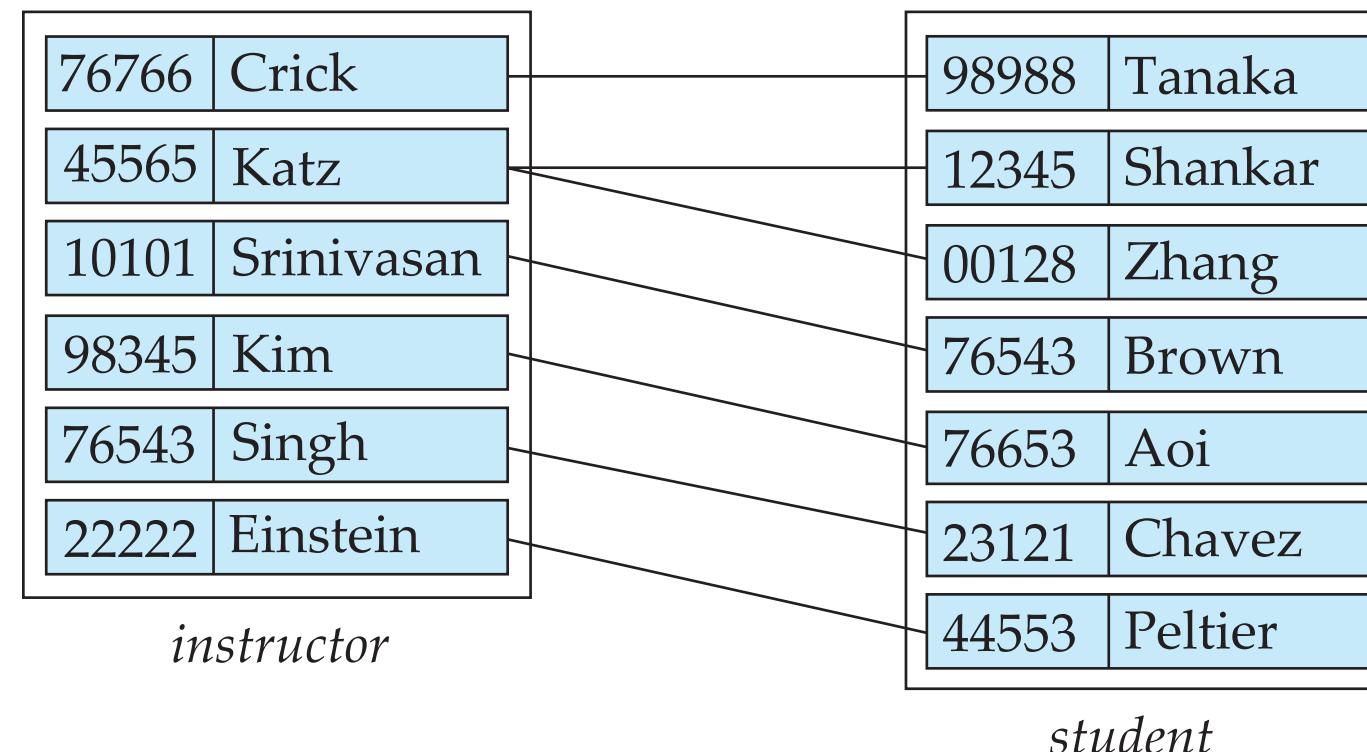
$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

where  $(e_1, e_2, \dots, e_n)$  is a relationship

- Example:  $(44553, 22222) \in \text{advisor}$
- The association between entity sets is referred to as participation
  - Entity sets  $E_1, E_2, \dots, E_n$  participate in relationship set R

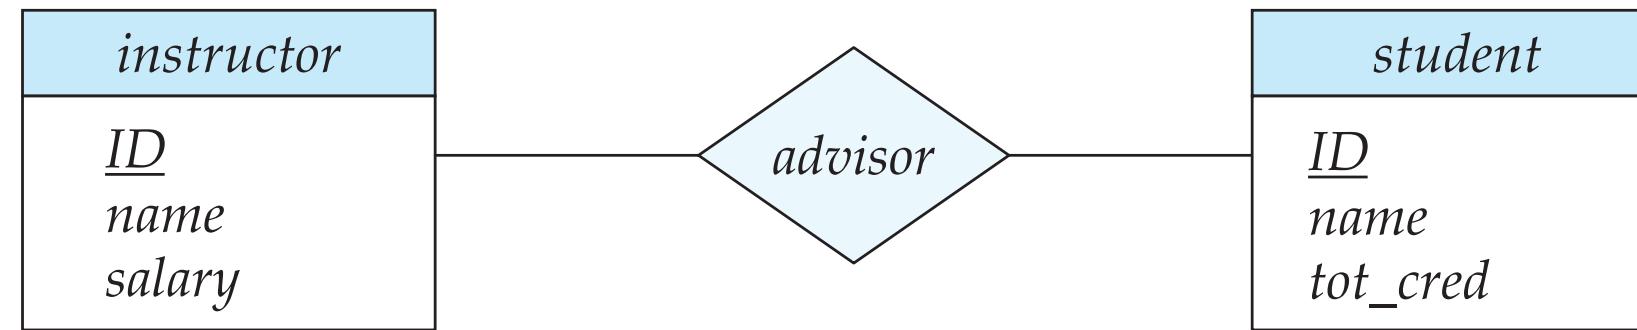
# Relationship Sets

- Example: we define the relationship set **advisor** to denote the associations between students and the instructors who act as their advisors.
  - Pictorially, we draw a line between related entities



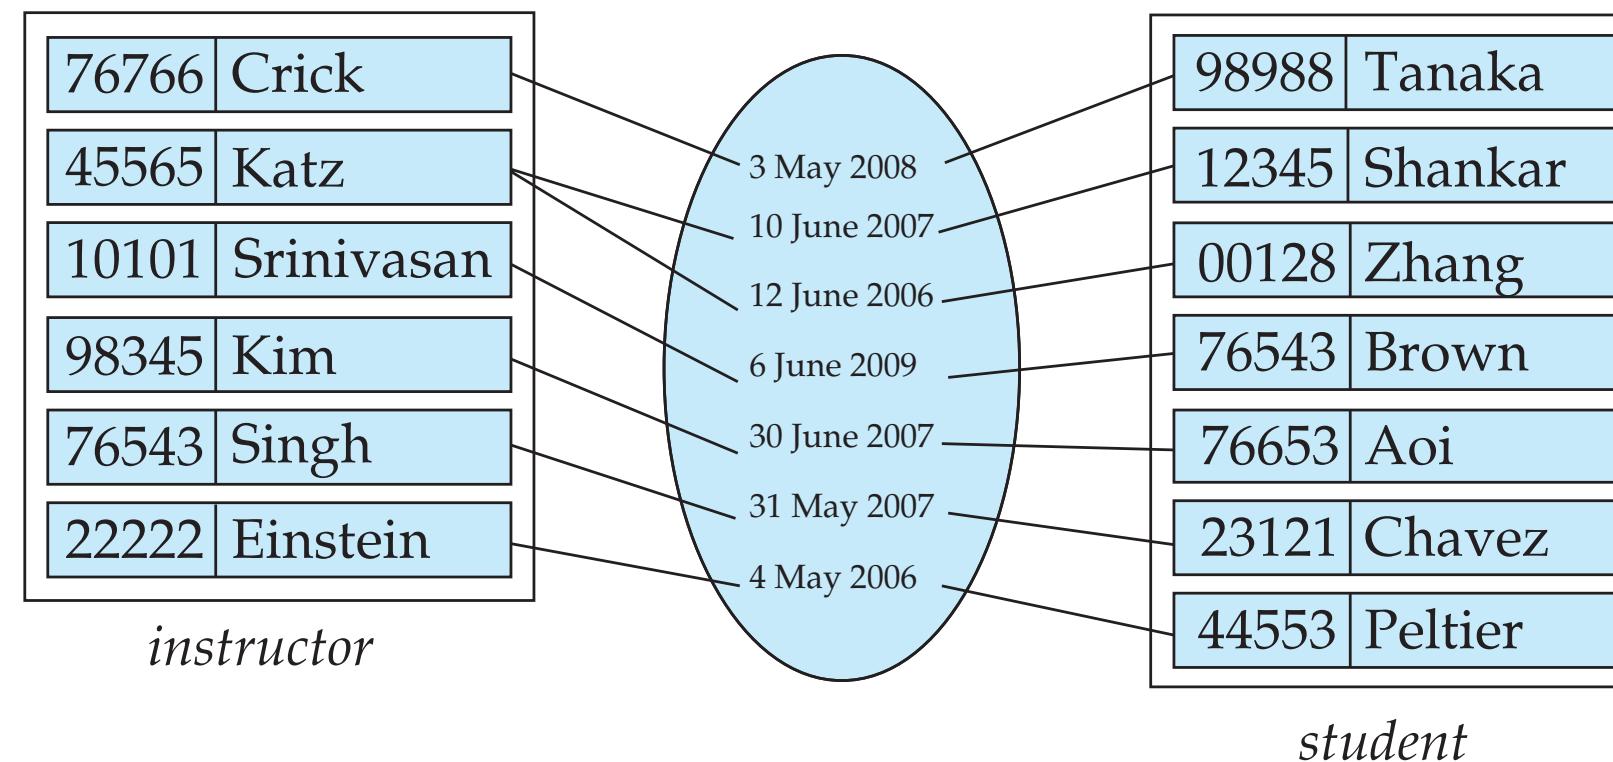
# Representing Relationship Sets via E-R Diagrams

- Use diamonds to represent relationship sets



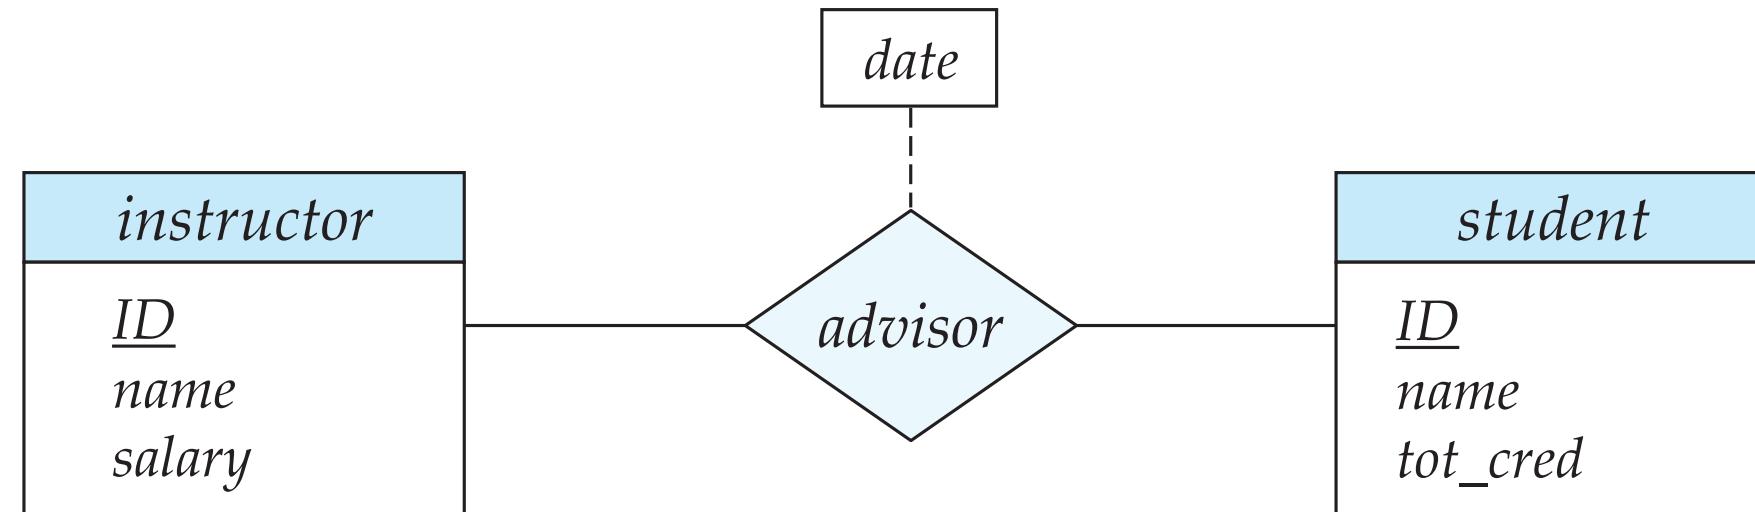
# Relationship Sets (Cont.)

- A **descriptive attribute** can be associated with a relationship set.
  - E.g., the advisor relationship set between entity sets **instructor** and **student** may have the attribute **date**, which tracks when the student started being associated with the advisor



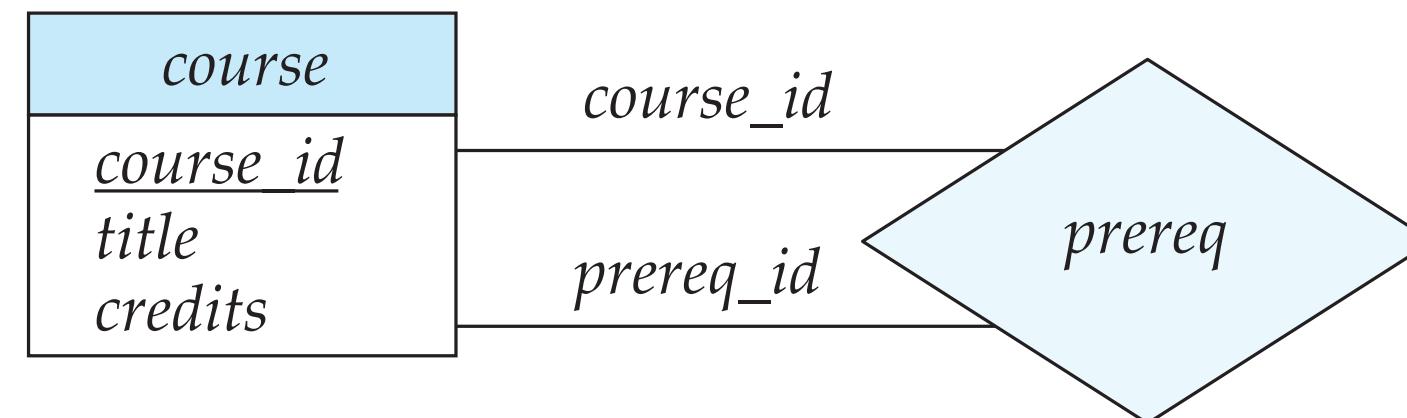
# Relationship Sets with Attributes

- Represented by using an **undivided rectangle**
- Linked to the diamond representing the relationship set via a **dashed line**



# Roles

- Entity sets of a relationship need not be distinct (i.e., can be the same entity set)
  - We can create **self-pointing relationships** for an entity set
  - Each occurrence of an entity set plays a “role” in the relationship
  - Example: A relationship set to represent the prerequisites of a course
    - E.g., Data Structure depends on Introduction to Programming
    - The labels “course\_id” and “prereq\_id” are called roles



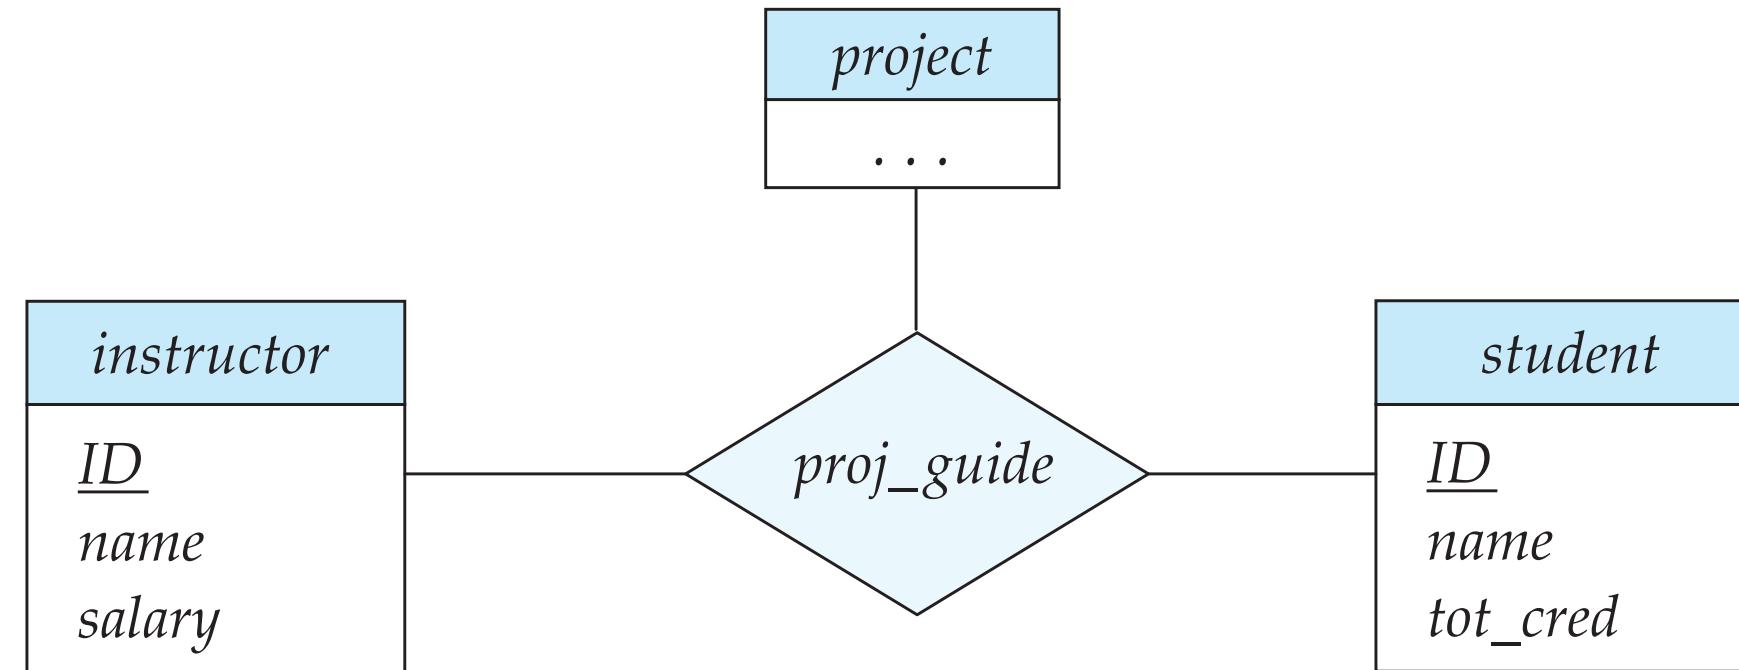
- If the entity sets participating in a relationship are **distinct**,
  - Their roles are implicit and usually are not specified

# Degree of a Relationship Set

- Denoted as the number of entity sets participating in a relationship set
- Binary relationship
  - Involve **two** entity sets (or degree two)
  - Most relationship sets in a database system are binary
- Relationships between more than two entity sets are rare
  - E.g., students work on research projects under the guidance of an instructor
    - relationship proj\_guide is a **ternary relationship** among instructor, student, and project
      - A particular student is guided by a particular instructor on a particular project

# Non-binary Relationship Sets

- Although most relationship sets are binary,
  - There are occasions when it is more convenient to represent relationships as non-binary
- E-R Diagram with a Ternary Relationship

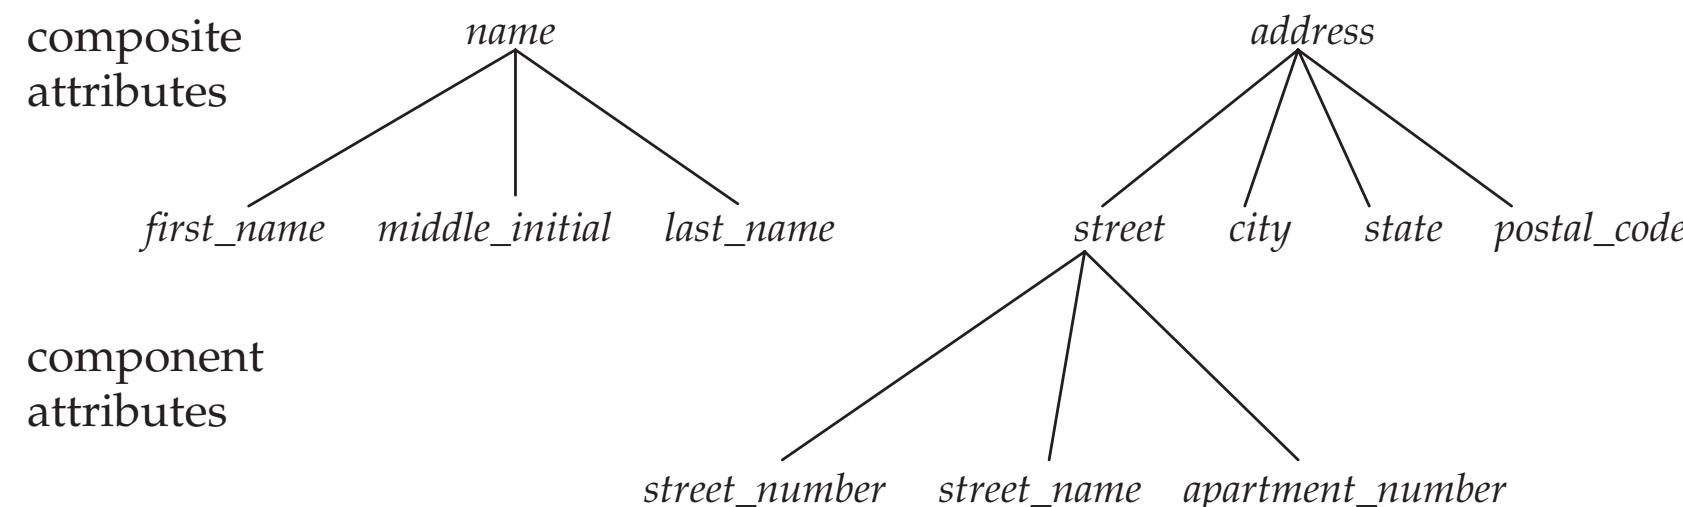


# Complex Attributes

- Attribute types
  - Simple (i.e., not divided into subparts) and composite (i.e., divided into subparts) attributes

# Composite Attributes

- Composite attributes allow us to divide attributes into subparts
  - Sometimes we may only use part of the attributes
  - In this case, composite attribute is a good design choice
- Allow us to group together related attributes, making the modeling cleaner
- A composite attribute may appear as a hierarchy



| <i>instructor</i>       |
|-------------------------|
| <i>ID</i>               |
| <i>name</i>             |
| <i>first_name</i>       |
| <i>middle_initial</i>   |
| <i>last_name</i>        |
| <i>address</i>          |
| <i>street</i>           |
| <i>street_number</i>    |
| <i>street_name</i>      |
| <i>apt_number</i>       |
| <i>city</i>             |
| <i>state</i>            |
| <i>zip</i>              |
| { <i>phone_number</i> } |
| <i>date_of_birth</i>    |
| <i>age ()</i>           |

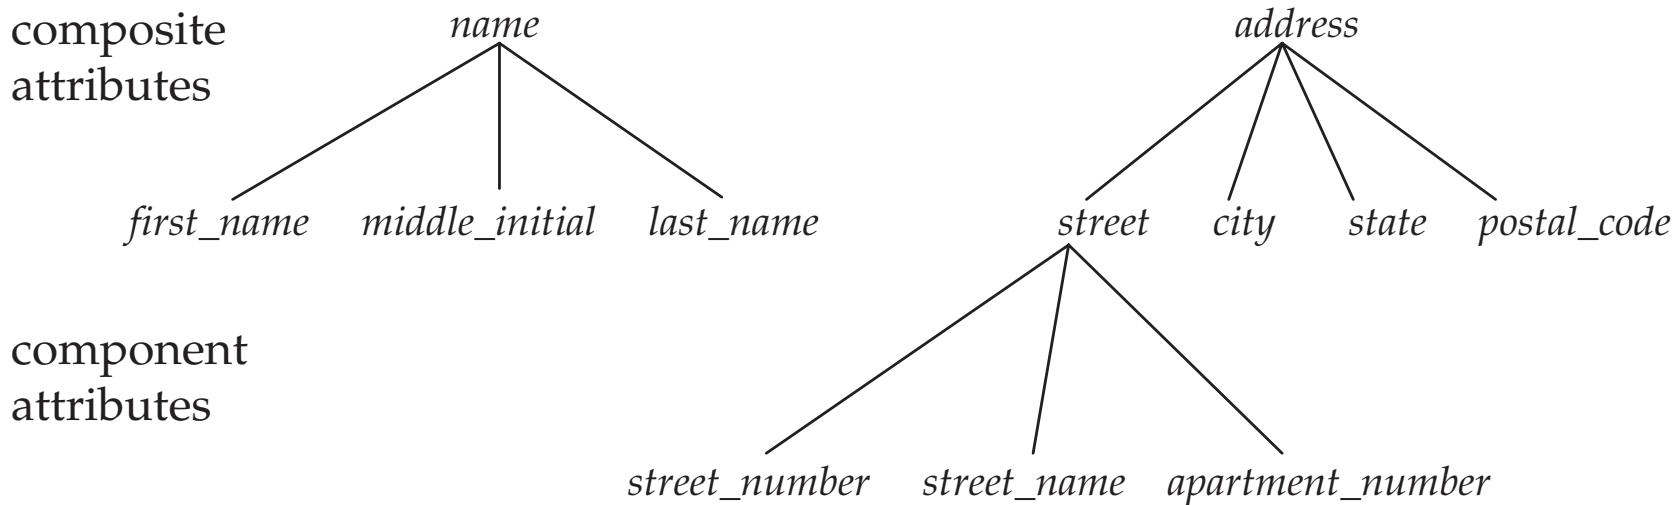
E-R  
notation

# Complex Attributes

- Attribute types
  - Simple (i.e., not divided into subparts) and composite (i.e., divided into subparts) attributes
  - Single-valued and multivalued attributes
    - Single-valued attribute: e.g., for a student, only one *student\_id*
    - Multivalued attribute
      - *phone\_numbers*: a person can have 0, 1, or multiple phone numbers at the same time
      - *grades*: each student may have multiples grades in a course
      - *dependent\_name*: an instructor may be hired by 0, 1, or multiple departments
  - Derived attributes
    - Can be computed from other attributes
    - Example: *age* computed based on *date\_of\_birth*, #students advised by an instructor
- Domain: The set of permitted values for each attribute
  - *course\_id* might be the set of all text strings of a certain length
  - *semester* might be strings from the set {Fall, Winter, Spring, Summer}

# E-R notation of Composite Attributes

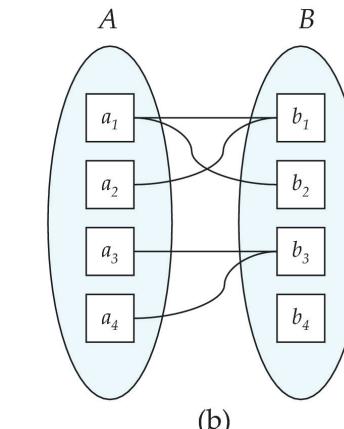
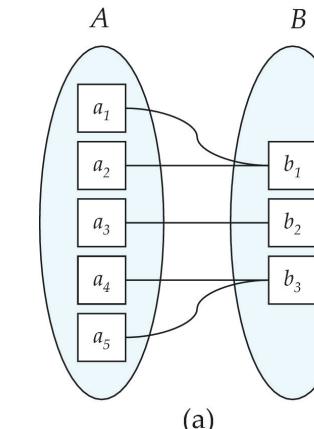
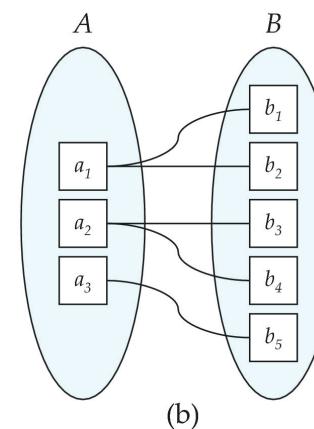
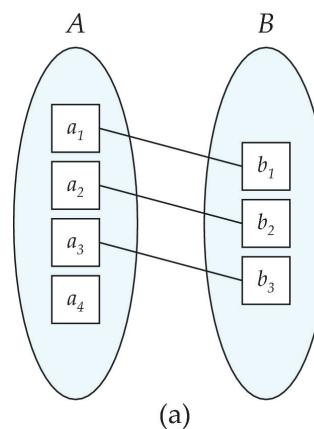
- E-R notations for
  - Composite attributes
  - Multivalued attributes
  - Derived attributes



|                         |
|-------------------------|
| <i>instructor</i>       |
| <i>ID</i>               |
| <i>name</i>             |
| <i>first_name</i>       |
| <i>middle_initial</i>   |
| <i>last_name</i>        |
| <i>address</i>          |
| <i>street</i>           |
| <i>street_number</i>    |
| <i>street_name</i>      |
| <i>apt_number</i>       |
| <i>city</i>             |
| <i>state</i>            |
| <i>zip</i>              |
| { <i>phone_number</i> } |
| <i>date_of_birth</i>    |
| <i>age()</i>            |

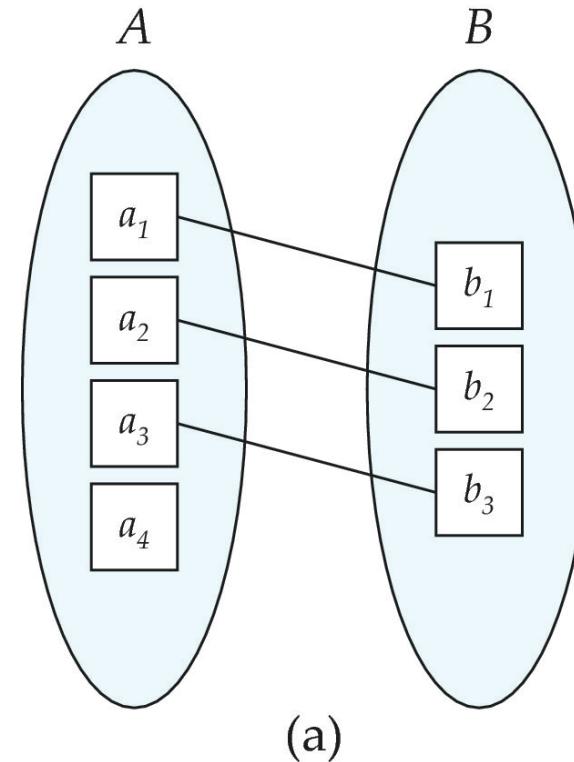
# Mapping Cardinality Constraints

- Mapping Cardinality (映射基数)
  - Express **the number of entities** to which **another entity** can be associated via a relationship set
    - Most useful in describing binary relationship sets
    - Can contribute to the description of non-binary relationship sets



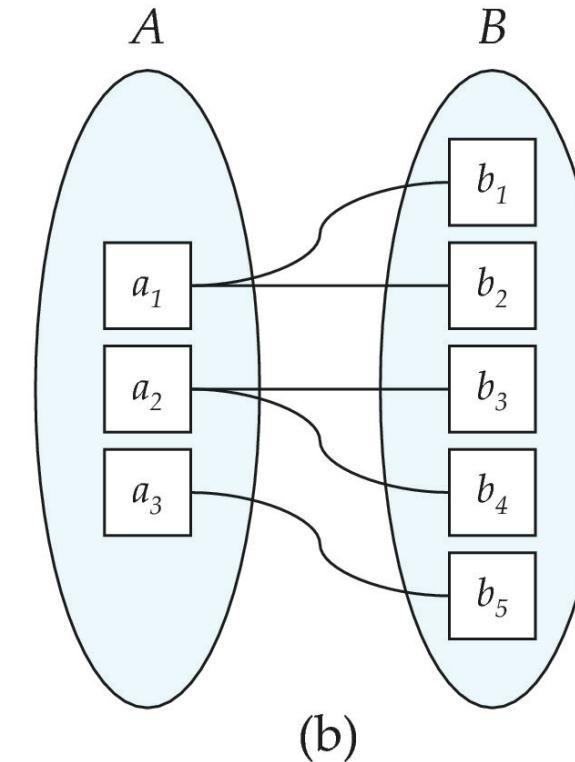
- For a binary relationship set, the mapping cardinality must be one of the following:
  - One to one, one to many, many to one, many to many

# Mapping Cardinalities



One to one

- Every entity in A is associated with at most one entity in B
- Same for entity set B

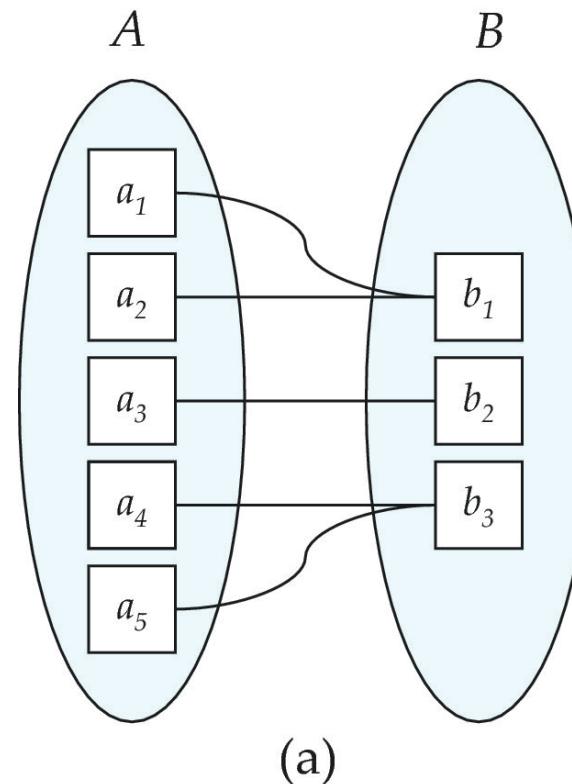


One to many

- Every entity in A is associated with 0, 1, or more entities in B
- Every entity in B is associated with at most one entity in A

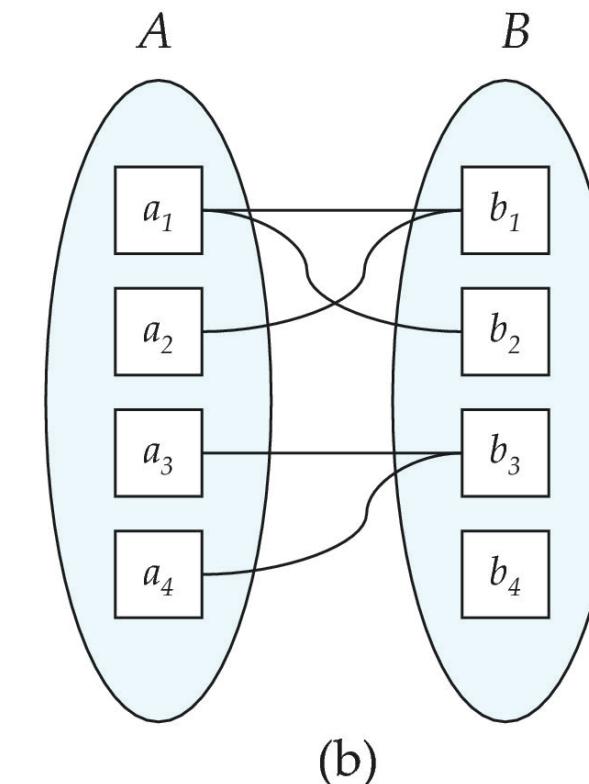
Note: Some entities in A and B may not be mapped to any entities in the other set

# Mapping Cardinalities



## Many to one

- Every entity in A is associated with at most one entity in B
- Every entity in B is associated with 0, 1 or more entities in A



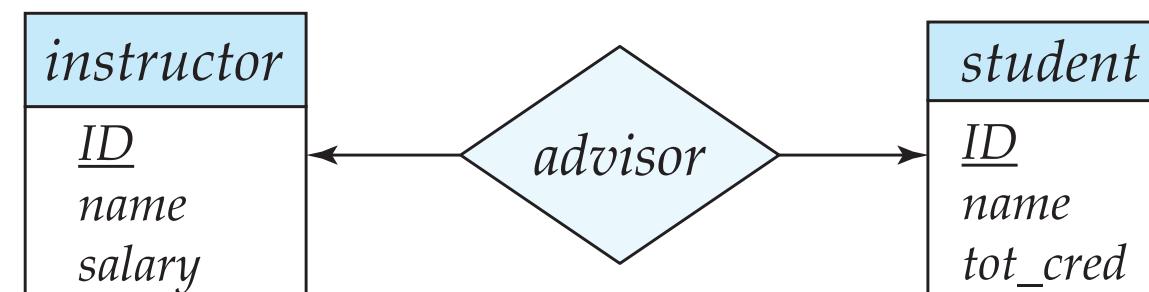
## Many to many

- Every entity in A is associated with 0, 1, or more entities in B
- Same for B

Note: Some entities in A and B may not be mapped to any entities in the other set

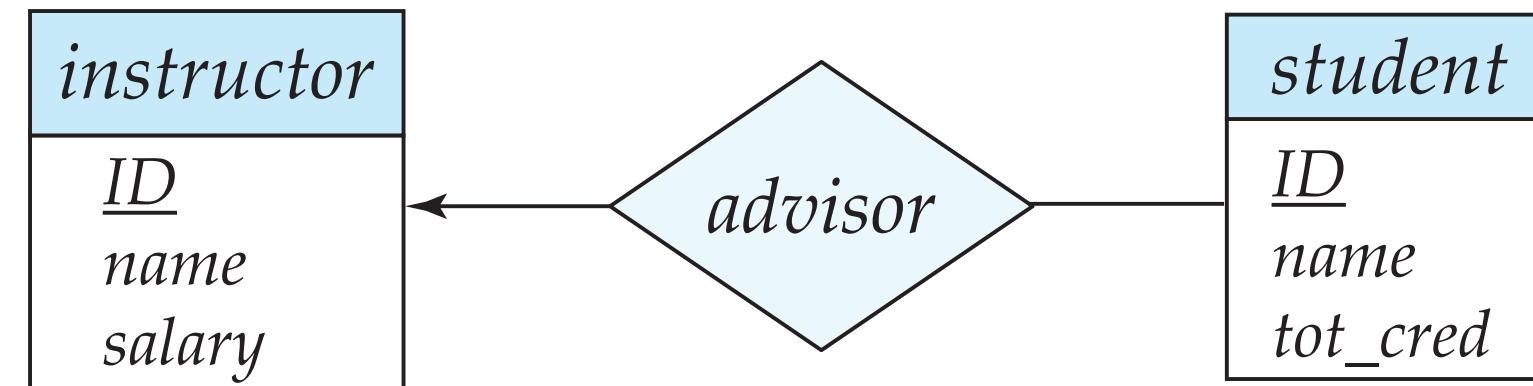
# Representing Cardinality Constraints in ER Diagram

- We express cardinality constraints by:
  - drawing either a directed line ( $\rightarrow$ ), signifying “one”
  - or an undirected line ( $-$ ), signifying “many”
- ... between the relationship set and the entity set
- One-to-one relationship between an instructor and a student :
  - A student is associated with at most one instructor via the relationship **advisor**



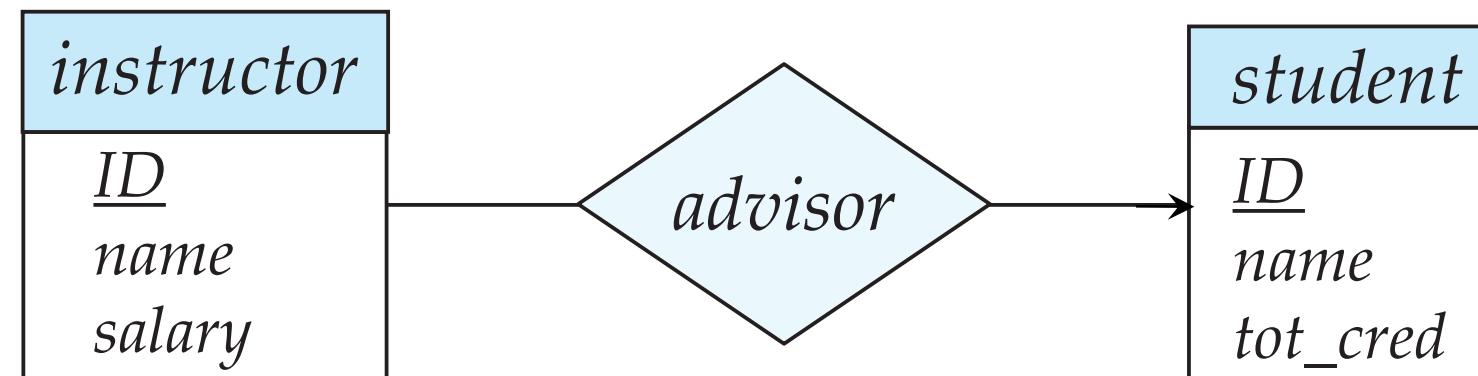
# Representing Cardinality Constraints in ER Diagram

- One-to-many relationship between an instructor and a student
  - an **instructor** is associated with **several (including 0) students** via **advisor**
  - a **student** is associated with **at most one instructor** via **advisor**



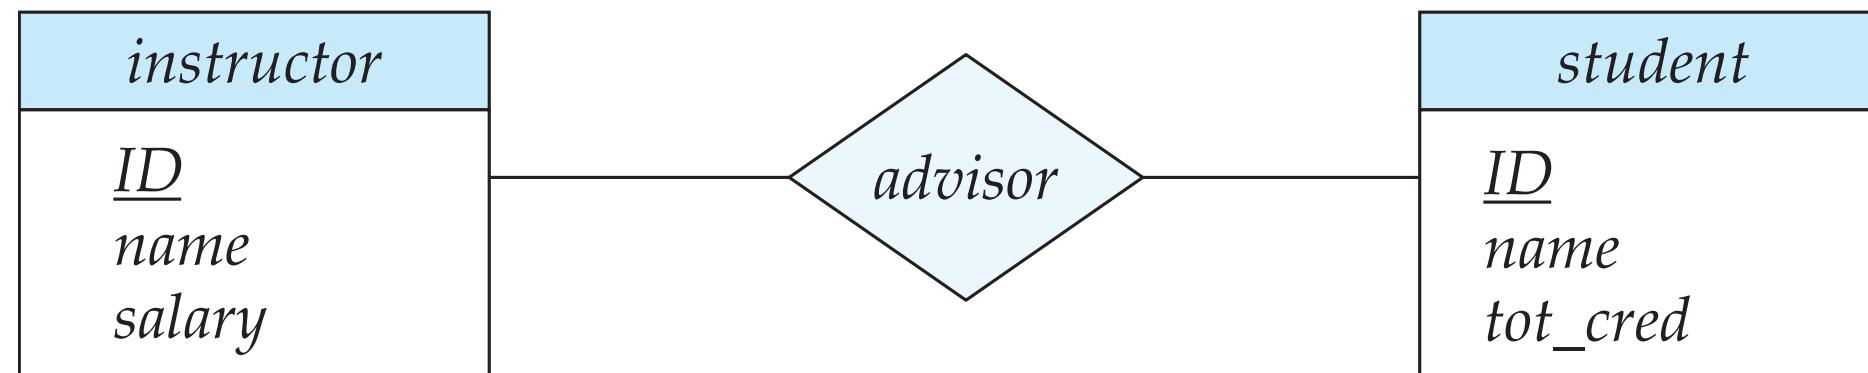
# Representing Cardinality Constraints in ER Diagram

- In a many-to-one relationship between an instructor and a student,
  - an **instructor** is associated with **at most one student** via **advisor**
  - and a **student** is associated with **several (including 0)** instructors via **advisor**



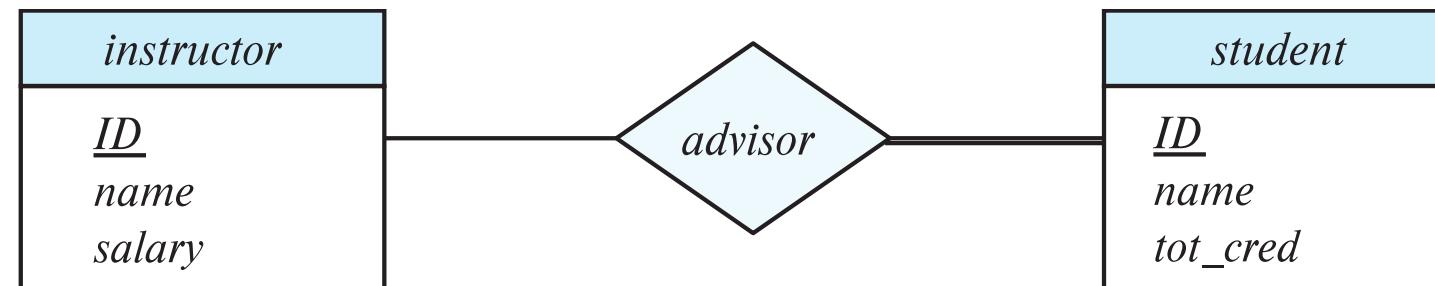
# Representing Cardinality Constraints in ER Diagram

- Many-to-many relationship:
  - An **instructor** is associated with **several (possibly 0) students** via **advisor**
  - A **student** is associated with **several (possibly 0) instructors** via **advisor**



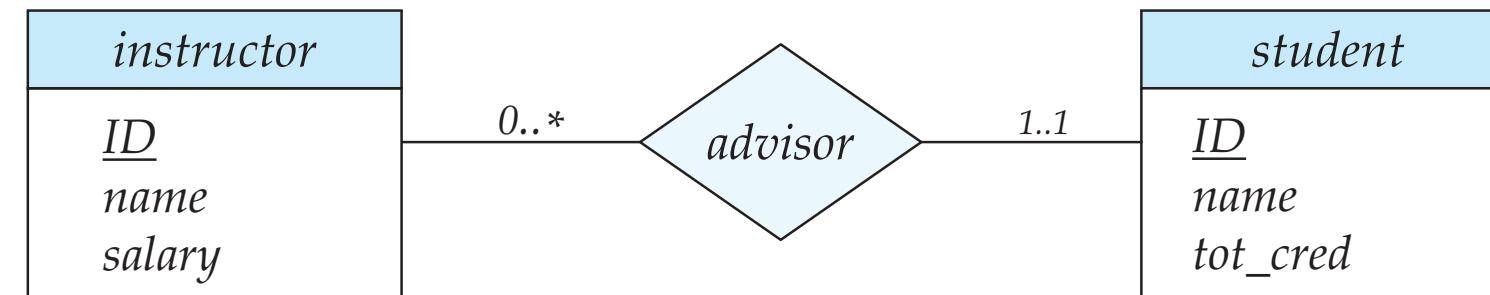
# Total and Partial Participation

- Total participation (indicated by *undirected double line, i.e., =*)
  - Every entity in the entity set participates in at least one relationship in the relationship set
  - Example: Participation of student in advisor relation is total
    - i.e., every student must have an associated instructor
- Partial participation (undirected line, i.e., –)
  - Some entities may not participate in any relationship in the relationship set
  - Example: participation of instructor in advisor is partial



# Notation for Expressing More Complex Constraints

- A line may have **an associated minimum and maximum cardinality**, shown in the form ***l..h***, where ***l*** is the minimum and ***h*** the maximum cardinality
  - A minimum value of 1 indicates total participation
  - A maximum value of 1 indicates that the entity participates in at most one relationship
  - A maximum value of \* indicates no limit



- Example
  - Instructor can advise 0 or more students
  - A student must have 1 advisor; cannot have multiple advisors

# Primary Key

- Primary keys provide a way to specify how entities and relations are distinguished

# Primary Key for Entity Sets

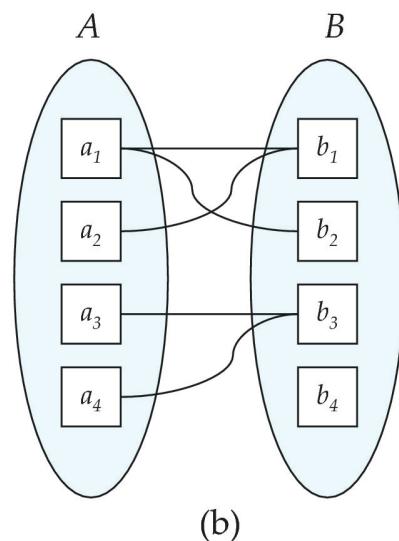
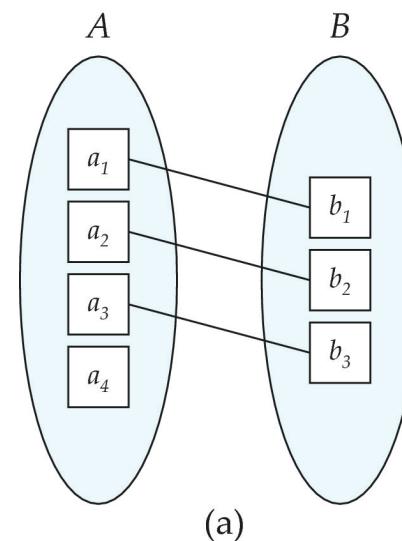
- By definition, individual entities are **distinct**
  - From database perspective, the differences among them must be expressed in terms of their attributes
- The values of the attribute values of an entity must be such that they can uniquely identify the entity
  - No two entities in an entity set are allowed to have exactly the same value for all attributes
- A **key** for an entity is a set of attributes that suffice to distinguish entities from each other

# Primary Key for Relationship Sets

- To **distinguish** among the various **relationships** of a relationship set, we **use** the **individual primary keys of the entities** in the relationship set.
  - Let R be a relationship set involving entity sets E1, E2, .. En
  - The **primary key for R** consists of the union of the primary keys of entity sets E1, E2, ..En
  - If the relationship set R has attributes  $a_1, a_2, \dots, a_m$  associated with it, the primary key of R also includes the attributes  $a_1, a_2, \dots, a_m$
- Example: relationship set “advisor”.
  - The primary key consists of **instructor.ID** and **student.ID**
- The choice of the primary key for a relationship set depends on the mapping cardinality of the relationship set

# Choice of Primary key for Binary Relationship

- Many-to-Many relationships
  - The **union of the primary keys** is a minimal superkey and is chosen as the primary key
- One-to-one relationships
  - The primary key of **either one of the participating entity sets** forms a minimal superkey, and either one can be chosen as the primary key.



- \* K is a **superkey** of R if values for K are sufficient to identify a unique tuple of each possible relation  $r(R)$
- Example: {ID} and {ID,name} are both superkeys of *instructor*.

# Choice of Primary key for Binary Relationship

- One-to-Many relationships
  - The **primary key of the “Many” side** is a minimal superkey and is used as the primary key
- Many-to-one relationships
  - The **primary key of the “Many” side** is a minimal superkey and is used as the primary key

