

# Computer **V**ision

CS308

Feng Zheng

SUSTech CS Vision Intelligence and Perception

Week 6



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



# Content

- Brief Review
- Fitting Techniques
  - Least Squares
  - Total Least Squares
- Random Sample Consensus (RANSAC)
- Hough Voting
- Image Alignment

# Brief Review



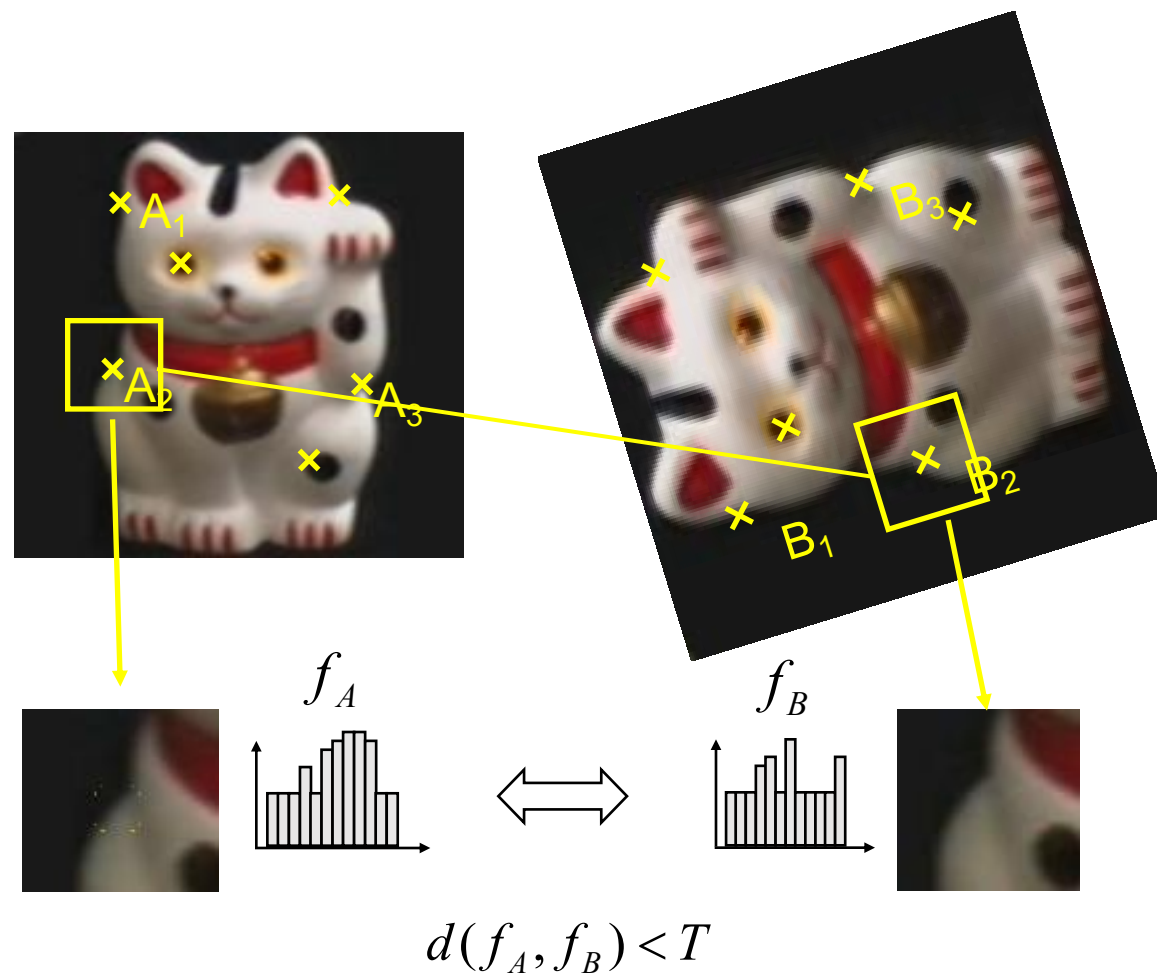
# Overview of Keypoint Matching

- Steps

- **Find** a set of distinctive keypoints
- Define a **region** around each keypoint
- Compute a local **descriptor** from the region
- **Match** local descriptors

- Goals

- Detect points that are **repeatable** and **distinctive**



# Fitting Techniques



# How Do We Build Panorama?

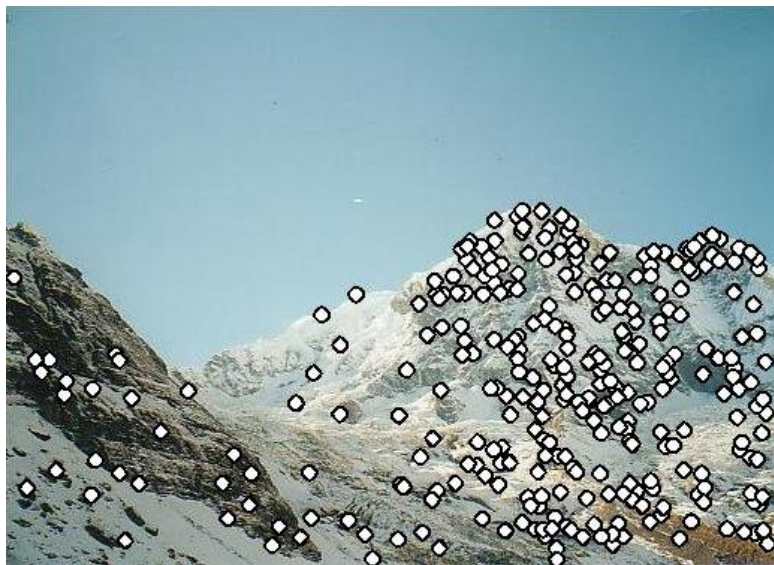
- We need to match (align) images





# Matching with Features

- Steps
  - Detect feature points in both images



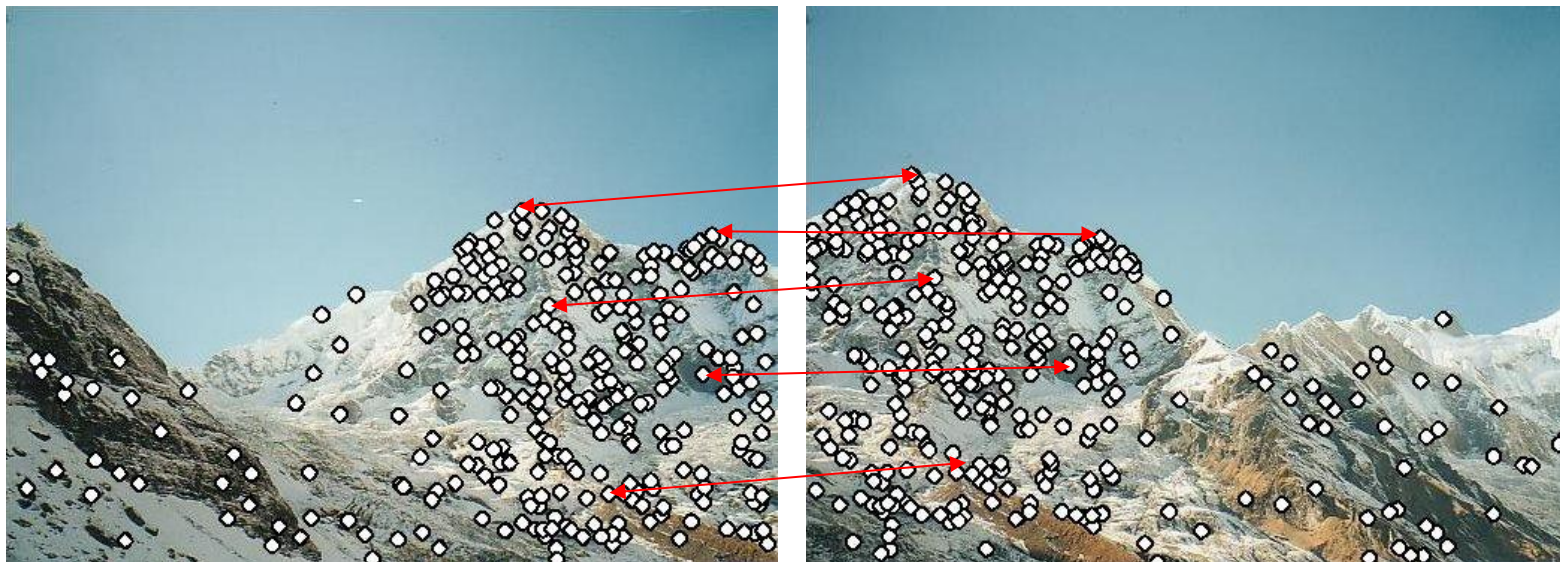




# Matching with Features

- Steps

- Detect feature points in both images
- Find corresponding pairs







# Matching with Features

- Steps


- Detect feature points in both images
  - Find corresponding pairs
  - Use these pairs to align images
- } Previous Lecture





# Fitting: Building a Model for a Set of Features


- Choose a **parametric model** to represent **a set of features**



Simple model: **lines**



Simple model: **circles**



Complicated model: **car**



# Fitting: Issues

- Case study: Line detection
  - **Noise** in the measured feature locations
  - **Extraneous** data: clutter (outliers), multiple lines
  - **Missing** data: occlusions





# Fitting: Issues


- If **we know which points belong to the line**, how do we find the "optimal" line parameters?
  - Least squares
- What if there are **outliers**?
  - Robust fitting, RANSAC
- What if there are **many lines**?
  - Voting methods: RANSAC, Hough transform
- What if we're **not even sure** it's a line?
  - Model selection



# Line Fitting: Ordinary Least Squares

- Data:  $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation:  $y_i = mx_i + b$
- Find  $(m, b)$  to minimize

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$




We know which points belong to the line



# Line Fitting: Ordinary Least Squares

- Data:  $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation:  $y_i = mx_i + b$
- Find  $(m, b)$  to minimize



$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$

$$\begin{aligned} E &= \sum_{i=1}^n \left( y_i - \begin{bmatrix} x_i & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right)^2 = \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right\|^2 = \|Y - XB\|^2 \\ &= (Y - XB)^T (Y - XB) = Y^T Y - 2(XB)^T Y + (XB)^T (XB) \end{aligned}$$





# Line Fitting: Ordinary Least Squares

- Normal equations: least squares solution to  $XB=Y$

$$\frac{dE}{dB} = 2X^T XB - 2X^T Y = 0$$

$$X^T XB = X^T Y$$

- Problem with "vertical" least squares
  - Not rotation-invariant
  - Fails completely for vertical lines

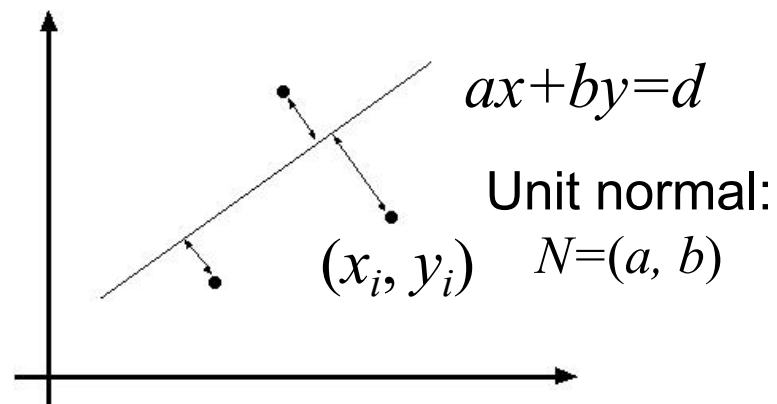
$$X^T X$$



# Total Least Squares

- Distance between point  $(x_i, y_i)$  and line  $ax+by=d$  ( $a^2+b^2=1$ ):  
 $|ax_i + by_i - d|$

$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$





# Total Least Squares

$$\frac{\partial E}{\partial d} = \sum_{i=1}^n -2(ax_i + by_i - d) = 0 \quad \Rightarrow \quad d = \frac{a}{n} \sum_{i=1}^n x_i + \frac{b}{n} \sum_{i=1}^n y_i = a\bar{x} + b\bar{y}$$

$$E = \sum_{i=1}^n (a(x_i - \bar{x}) + b(y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = (UN)^T (UN)$$

$\uparrow$   
 $U$

$\uparrow$   
 $N$

$$\frac{dE}{dN} = 2(U^T U)N = 0$$

- Solution to  $(U^T U)N = 0$ , subject to  $\|N\|^2 = 1$ : eigenvector of  $U^T U$  associated with the smallest eigenvalue (least squares solution to *homogeneous linear system*  $UN = 0$ )

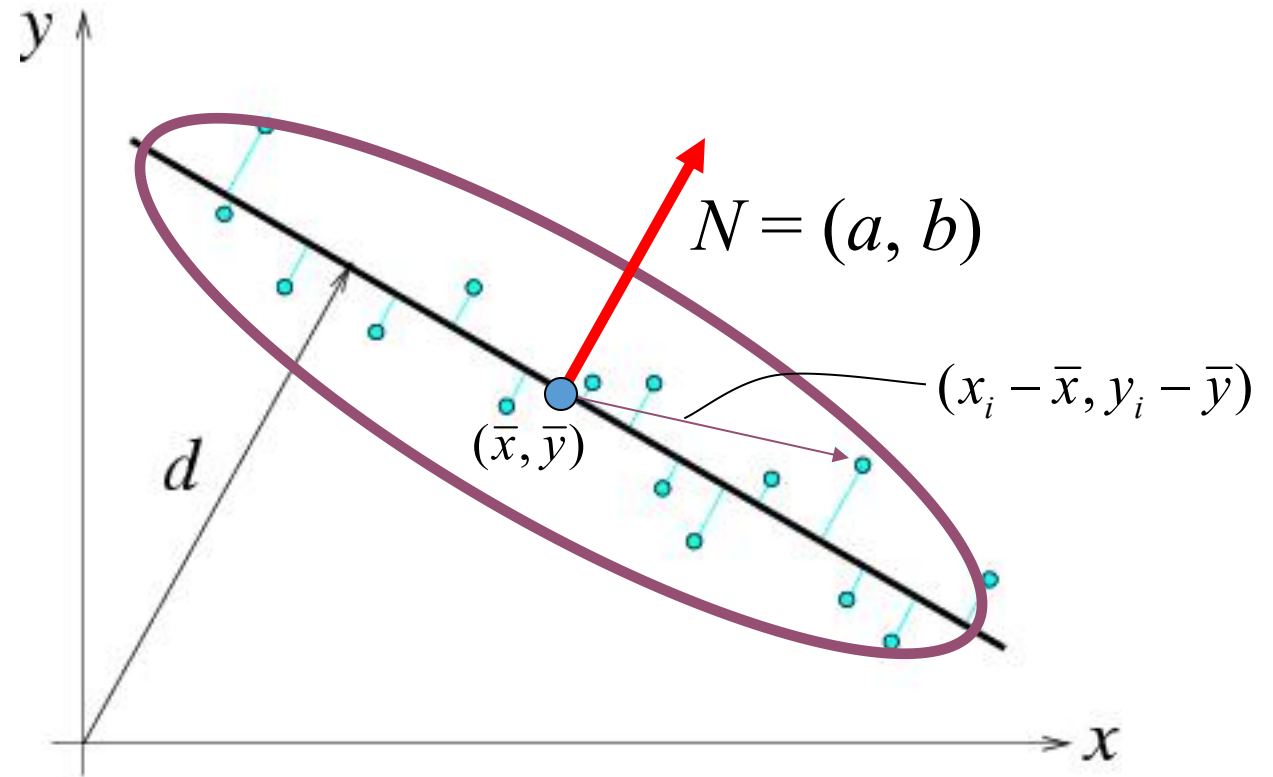


# Total Least Squares

- Second moment matrix

$$U = \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix}$$

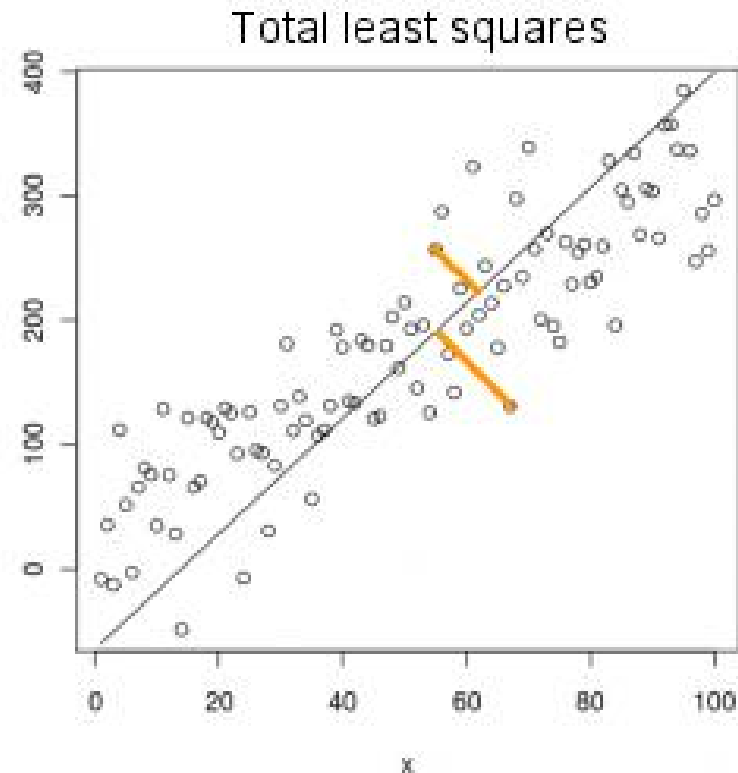

$$U^T U = \begin{bmatrix} \sum_{i=1}^n (x_i - \bar{x})^2 & \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) \\ \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) & \sum_{i=1}^n (y_i - \bar{y})^2 \end{bmatrix}$$





# OLS vs. TLS


- The difference between standard OLS regression and "orthogonal" TLS regression





# Total Least Squares

- Robustness to **noise**:  
least squares fit to the  
red points

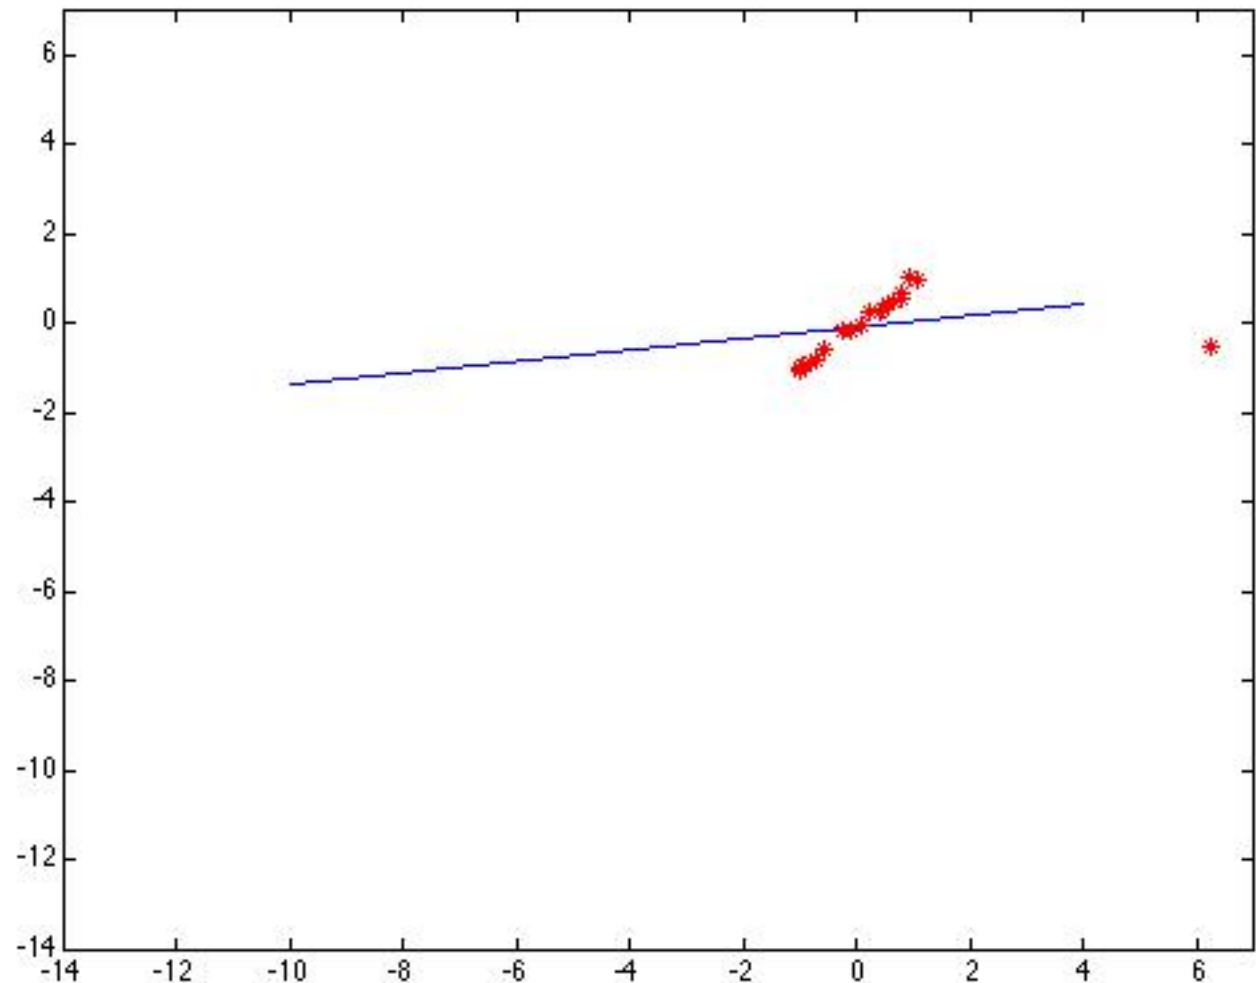






# Total Least Squares

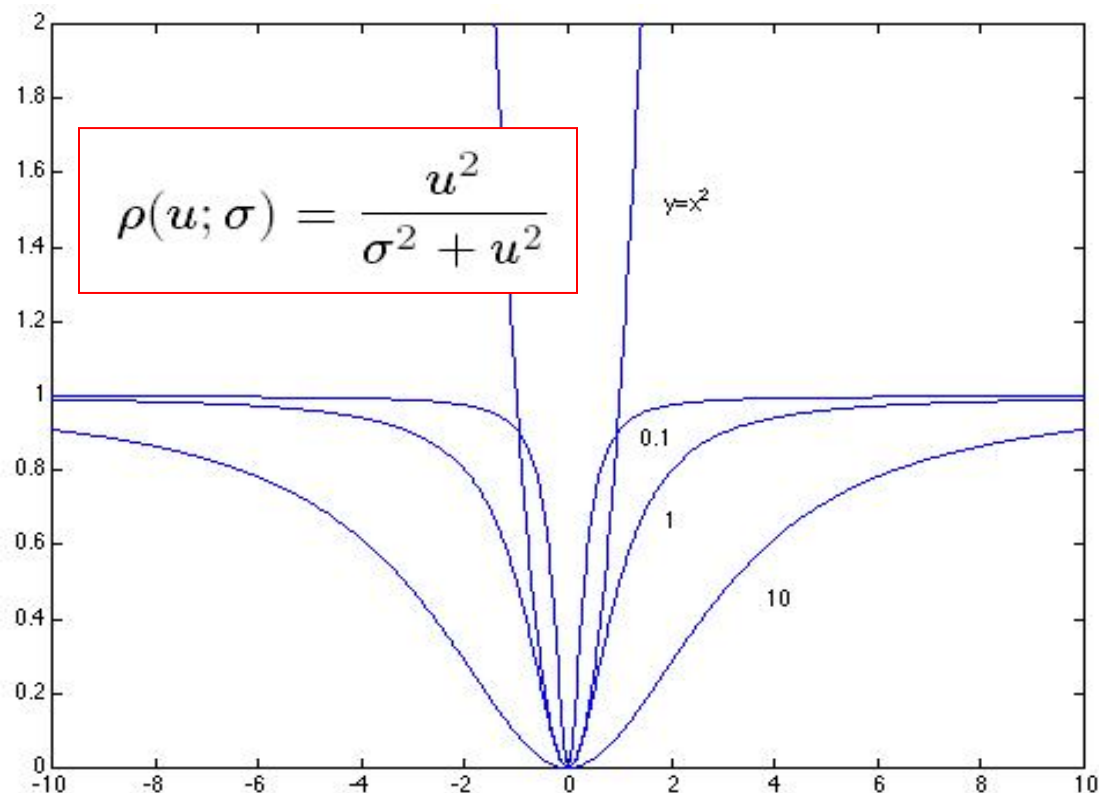
- Robustness to noise:  
Least squares fit with  
an **outlier**
- Problem: squared error  
**heavily** penalizes  
outliers





# Robust Estimators

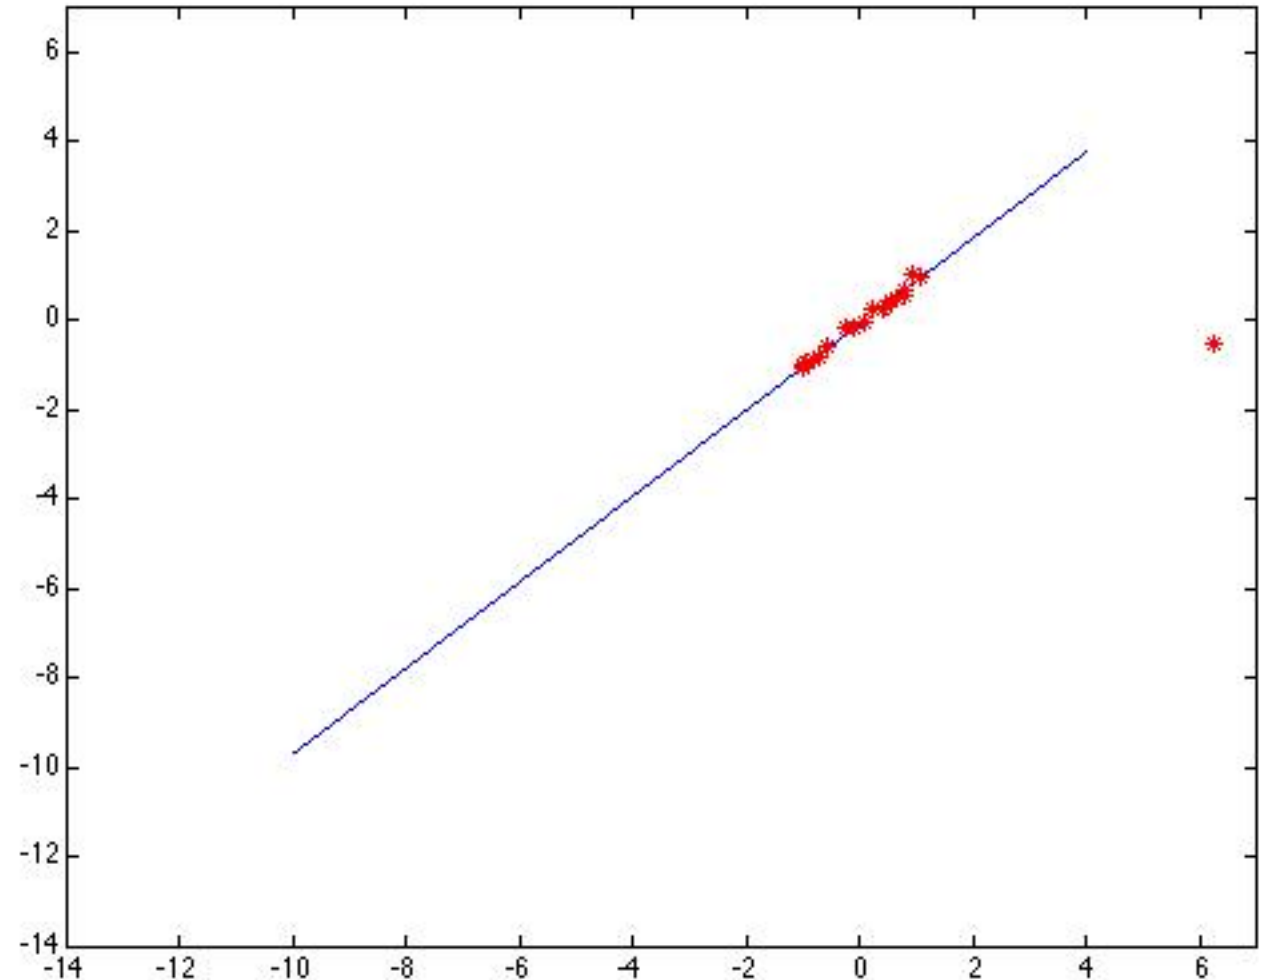
- General approach---minimize:  $\sum_i \rho(r_i(x_i, \theta); \sigma)$   
 $r_i(x_i, \theta)$  – **residual** of  $i$ th point  
w.r.t. model parameters  $\theta$   
 $\rho$  – robust function with  
scale parameter  $\sigma$
- The robust function  $\rho$   
behaves like squared  
distance for **small values** of  
the residual  $u$  but saturates  
for **larger values** of  $u$





# Choosing the Scale: Just Right

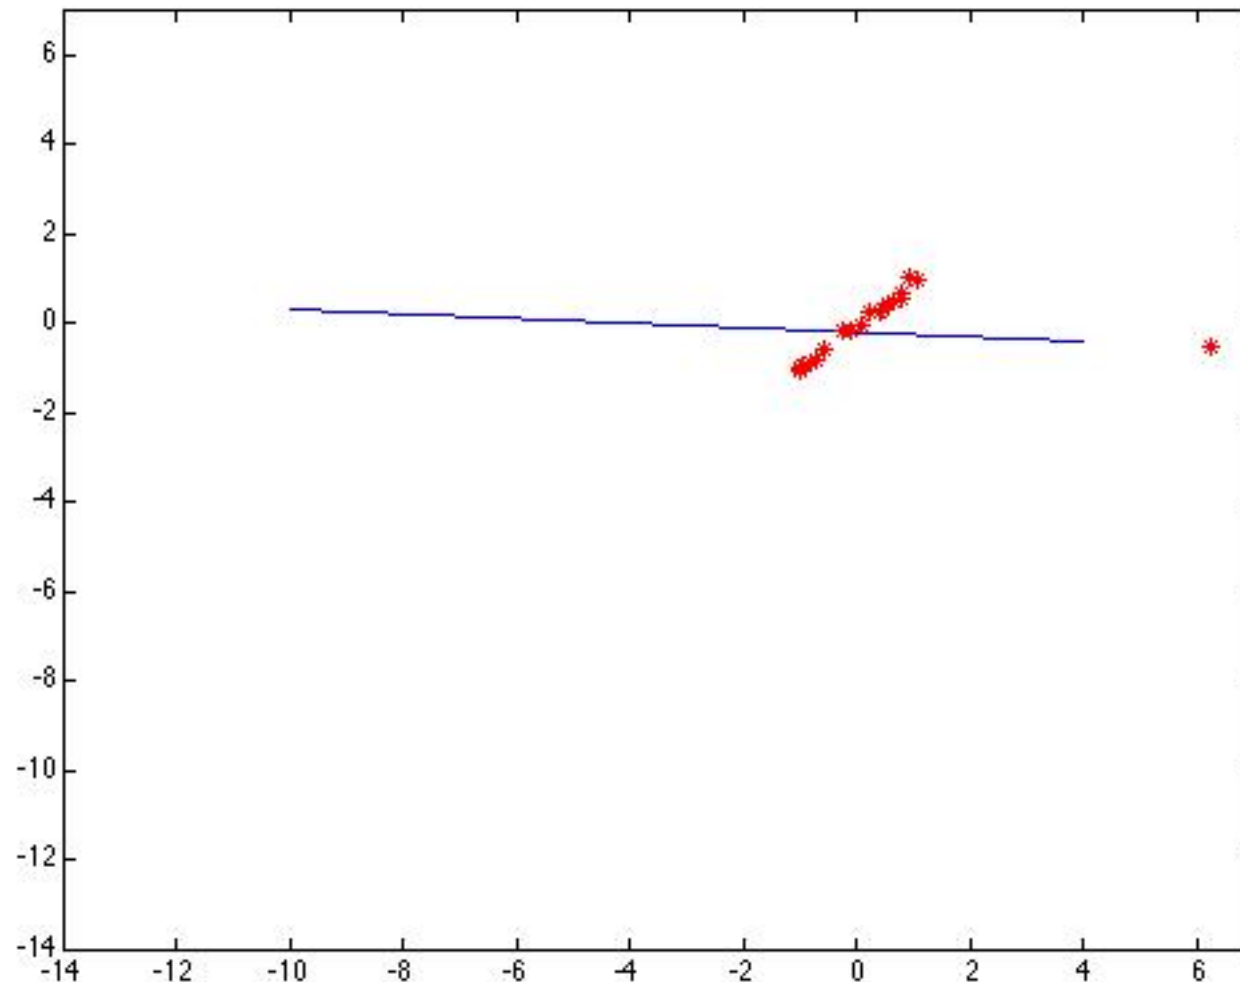
- The effect of the **outlier** is minimized, when choosing a **just right** scale





# Choosing the Scale: Too Small

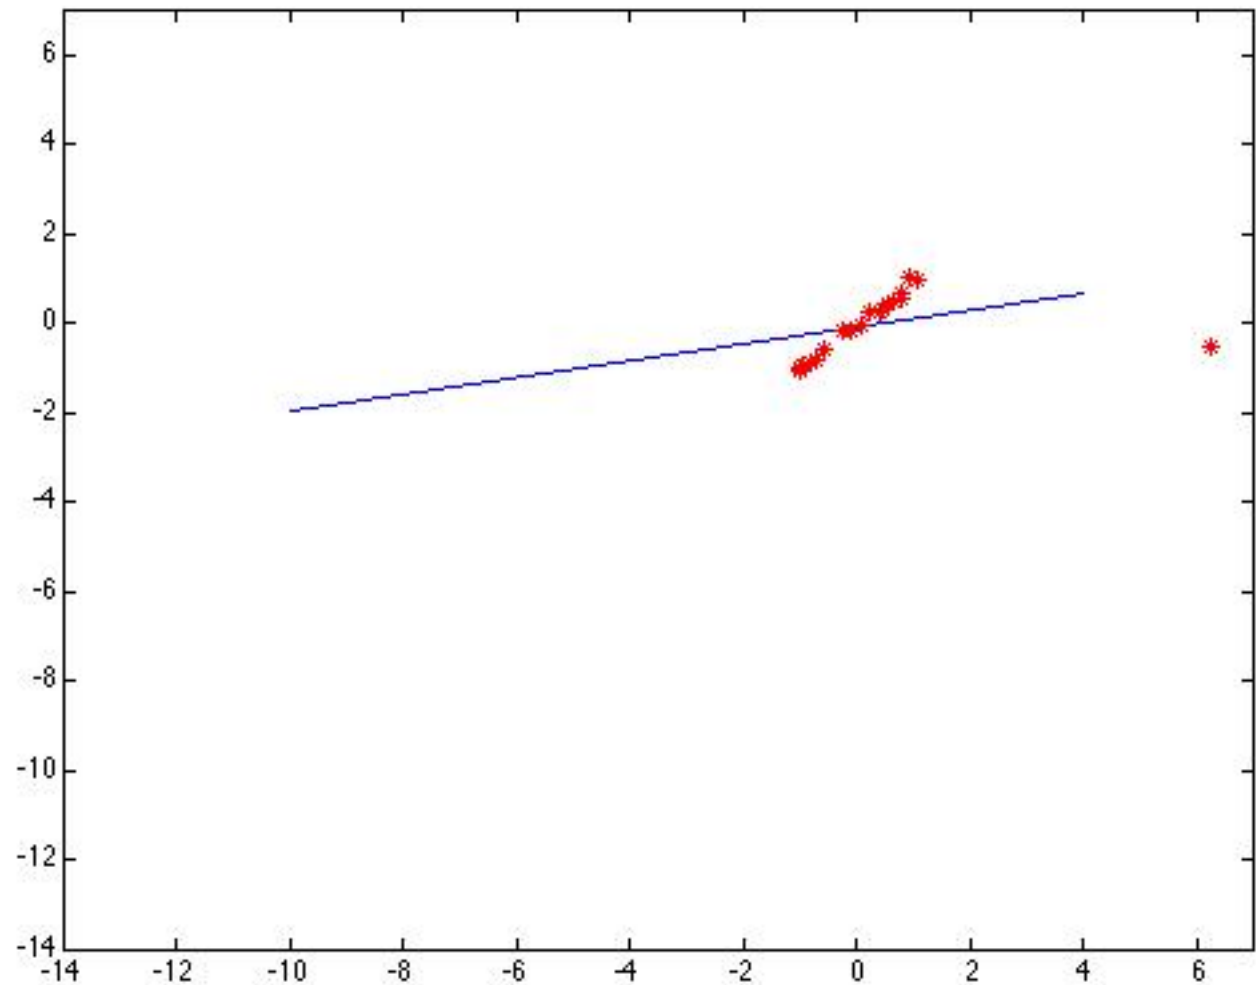
- The error value is almost the **same for every point** and the fit is very **poor**





# Choosing the Scale: Too Large

- Behaves much the same as least squares






# Curve Fitting

- **Find Polynomial:**  $y = f(x) = ax^3 + bx^2 + cx + d$ 
  - That best fits the given points  $(x_i, y_i)$

- **Minimize:**  $\frac{1}{N} \sum_i [y_i - (ax_i^3 + bx_i^2 + cx_i + d)]^2$

- **Using:**  $\frac{\partial E}{\partial a} = 0$  ,  $\frac{\partial E}{\partial b} = 0$  ,  $\frac{\partial E}{\partial c} = 0$  ,  $\frac{\partial E}{\partial d} = 0$

- **Note:**  $f(x)$  is **LINEAR** in the parameters  $(a, b, c, d)$





# Random Sample Consensus



# RANSAC

- Robust fitting (TLS) can deal with **a few** outliers - what if we have very **many**?
- Random sample consensus (RANSAC): Very general framework for model fitting in the **presence of outliers**
- Outline
  - Choose a **small subset** of points uniformly at random
  - Fit **a model** to that subset
  - Find all remaining points that are "**close**" to the **model** and reject the rest as outliers
  - Do this **many** times and choose the **best** model



# RANSAC for Line Fitting

- Algorithm
- Repeat  $N$  times:
  - Draw  $s$  points uniformly at random
  - Fit line to these  $s$  points (TLS)
  - Find inliers to this line among the remaining points (i.e., points whose distance from the line is less than  $t$ )
  - If there are  $d$  or more inliers, accept the line and refit using all inliers
- End
- Four parameters:  $s$ ,  $t$ ,  $d$  and  $N$



# Choosing the Parameters

- Initial number of points  $s$ 
  - Minimum number needed to fit the model  
✓ 2 points
- Distance threshold  $t$ 
  - (1) Choose  $t$  so probability for inlier is  $p$  (e.g. 0.95)
  - (2) Zero-mean Gaussian noise with standard deviation  $\sigma$ :  $t^2 = 3.84\sigma^2$



# Choosing the Parameters

- Number of times  $N$ 
  - Choose  $N$  so that, with probability  $p$ , at least one random sample is free from outliers (e.g.  $p=0.99$ )

Desired success rate after  $N$  times:  $p$

Outlier ratio (Unknown):  $e$

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$

$N$	proportion of outliers $e$						
s	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

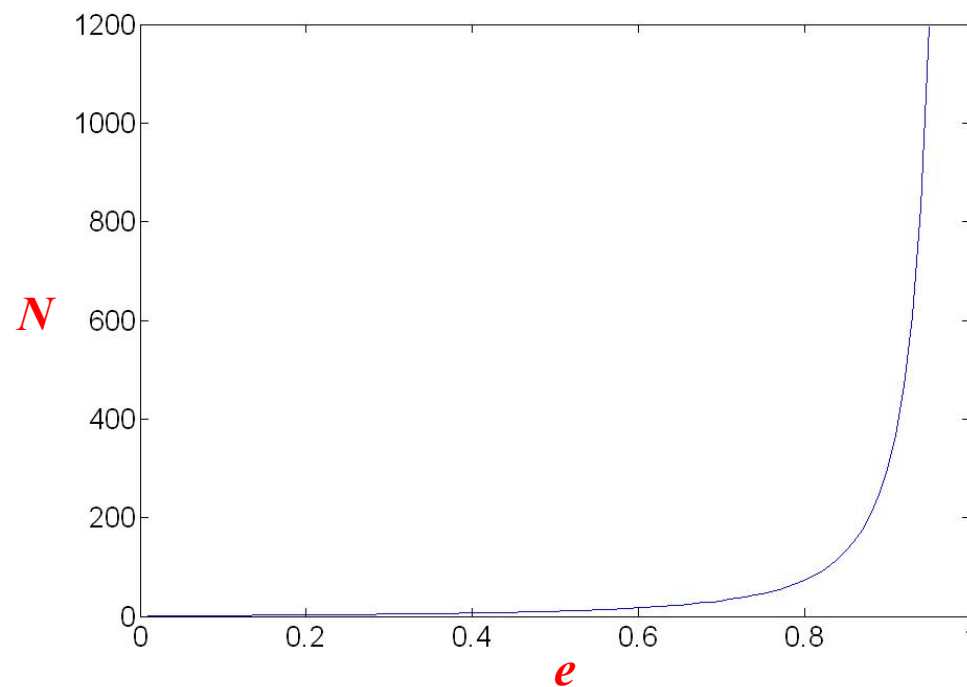


# Choosing the Parameters

- Consensus set size  $d$  (number of inliers)
  - Should match expected **inlier ratio**

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$







# Adaptively determining the number of samples

- Inlier ratio  $e$  is often unknown a priori, so pick worst case, e.g. 50%, and **adapt** if more inliers are found, e.g. 80% would yield  $e=0.2$
- Adaptive procedure:
  - $N=\infty, sample\_count=0$
  - While  $N > sample\_count$ 
    - ✓ Choose a sample (fitting) and count the number of inliers
    - ✓ Set  $e = 1 - (\text{number of inliers})/(\text{total number of points})$
    - ✓ Recompute  $N$  from  $e$ :
$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$
    - ✓ Increment the *sample\_count* by 1

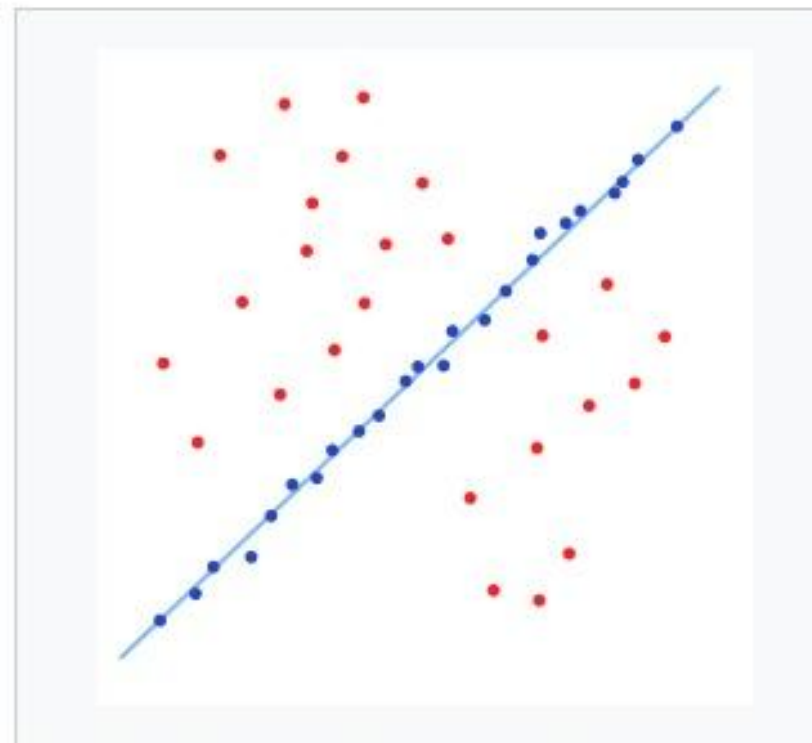


# RANSAC

- An example



A data set with many outliers for which a line has to be fitted.



Fitted line with RANSAC; outliers have no influence on the result.



# RANSAC pros and cons

- Pros

- Simple and general
- Applicable to many different problems
- Often works well in practice

- Cons

- **Lots** of parameters to tune
- **Can't always** get a good initialization of the model based on the minimum number of samples
- Sometimes **too many iterations** are required
- Can fail for extremely **low inlier ratios**
- We can often do better than **brute-force sampling**

Hough transform



# Voting Schemes

- Principal of voting
  - Let **each** feature (voter) vote for **all** the models that are compatible with it
  - Hopefully the **noise** features (voter) will **not vote consistently** for **any** single model (nominator)
  - **Missing data doesn't matter** as long as there are **enough** features remaining to agree on a good model



# Hough Transform

- An early type of voting scheme
- General outline:
  - Discretize parameter space into bins
  - For **each** feature point in the image, put a **vote** in every bin in the parameter space that **could have generated this point**
  - Find bins that have the **most** votes






Image space




Hough parameter space



# Parameter Space Representation


- A **line** in the image corresponds to a **point** in Hough space





# Parameter Space Representation

- What does a point  $(x_0, y_0)$  in the image space map to in the Hough space?







# Parameter Space Representation

- What does a point  $(x_0, y_0)$  in the image space map to in the Hough space?
  - Answer: the solutions of  $b = -x_0m + y_0$
  - This is a **line** in Hough space





Image space




Hough parameter space



# Parameter Space Representation



- Where is the line that contains both  $(x_0, y_0)$  and  $(x_1, y_1)$ ?





# Parameter Space Representation

- Where is the line that contains both  $(x_0, y_0)$  and  $(x_1, y_1)$ ?
  - It is the **intersection** of the lines  $b = -x_0m + y_0$  and  $b = -x_1m + y_1$






# Parameter Space Representation

- Problems with the  $(m, b)$  space:
  - Unbounded parameter domain
  - Vertical lines require infinite  $m$
- Alternative: polar representation

$$x \cos \theta + y \sin \theta = \rho$$




- Each point will add a sinusoid in the  $(\theta, \rho)$  parameter space



# Algorithm Outline

- Initialize accumulator  $H$  to all zeros
- For **each** edge point  $(x,y)$  in the image
  - For  $\theta = 0$  to  $180$ 
    - $\rho = x \cos \theta + y \sin \theta$
    - $H(\theta, \rho) = H(\theta, \rho) + 1$
  - end
- end
- Find the value(s) of  $(\theta, \rho)$  where  $H(\theta, \rho)$  is a **local maximum**
  - The detected line in the image is given by
$$\rho = x \cos \theta + y \sin \theta$$

H: accumulator array (votes)





# Basic Illustration

- A line





Image space



Votes

Horizontal axis is  $\theta$   
Vertical is  $\rho$ .



# Basic Illustration

- A line with noise





Image space



Votes



# Basic Illustration

- Scattered points





Image space



Votes





# Mechanics of the Hough transform

- Difficulties

- How big should the **cells** be? (too big, and we merge quite different lines; too small, and noise causes lines to be missed)

- How many lines?

- Count the **peaks** in the Hough array
- Treat **adjacent** peaks as a single peak

- Which points belong to each line?


- Search for points close to the line
- Solve again for line and iterate



# Real World Example




Original



Edge Detection



Parameter Space



Found Lines



# Finding Circles by Hough Transform



- Equation of Circle:

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- If radius is known:

- 2D Hough Space

- Accumulator Array:  $A(a, b)$





# Finding Circles by Hough Transform

- Equation of Circle:


$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- If radius is not known:

- 3D Hough space!

- Use Accumulator array:  $A(a, b, r)$

- What is the surface in the Hough space?





# Finding Circles by Hough Transform

- Hough transform for circles






Image space




Hough parameter space




# Finding Coins

Original



Edges (note noise)






# Finding Coins

- Note that because the quarters and penny are **different sizes**, a different Hough transform (with separate accumulators) was used for each circle size

Penny


Quarters





# Generalized Hough Transform

- We want to find a fixed shape (known) defined by its boundary points and a reference point








# Generalized Hough Transform

- We want to find **a fixed shape (known)** defined by its boundary points and a reference point
- For every boundary point  $p$ , we can compute the **displacement** vector  $r = a - p$  as a function of **gradient** orientation  $\theta$






# Generalized Hough Transform

- Construct a model for a shape:
  - Construct a table indexed by  $\theta$  storing displacement vectors  $r$  as function of **gradient direction**
- Detect using the model
  - For each **edge point**  $p$  with gradient orientation  $\theta$ :
    - ✓ Retrieve **all**  $r$  indexed with  $\theta$
    - ✓ For **each**  $r(\theta)$ , put a vote in the Hough space at  $p + r(\theta)$
  - **Peak** in this Hough space is **reference** point with most supporting edges
- Assumption: translation is the only transformation here, i.e., orientation and scale are fixed



# Example: a Known and Fixed Shape


- **Model** shape
  - Gradient orientation
  - No rotation





# Example: Building a Table


- Displacement vectors for model points





# Example: Detection


- Range of voting locations for **test** point





# Example: Detection


- Range of voting locations for test point





# Example: Detection

- **Votes** for points with  $\theta = \uparrow$




$r$		
$\theta$		



# Example: Building a Table

- Displacement vectors for model points








# Example: Detection

- Range of **voting** locations for test point






# Example: Detection

- Votes for points with

$$\theta = \swarrow$$




$r$		




# Application in Recognition

- Instead of indexing **displacements** by gradient orientation, index by "**visual codeword**"



training image

What is the codeword?



visual codeword with  
displacement vectors



# Application in Recognition

- Instead of indexing displacements by gradient orientation, index by "visual codeword"




test image

# Image Alignment



# Image Alignment



- Two broad approaches:
  - Direct (pixel-based) alignment
    - ✓ Search for alignment where most pixels agree
  - Feature-based alignment
    - ✓ Search for alignment where *extracted features* agree
    - ✓ Can be verified using pixel-based alignment





# Alignment as Fitting

- Alignment: fitting a model to a transformation between pairs of features (*matches*) in **two** images




Find transformation  $T$   
that minimizes

$$\sum_i \text{residual}(T(x_i), x'_i)$$



# Alignment as Fitting

- Previously: fitting a model to features in **one** image



Find model  $M$  that minimizes


$$\sum_i \text{residual}(x_i, M)$$






# 2D Transformation Models


- Similarity  
(translation, scale, rotation)



- Affine




- Projective  
(homography)





# Affine Transformations


- Simple fitting procedure (linear least squares)
- Approximates **viewpoint changes** for roughly planar objects and roughly orthographic cameras
- Can be used to initialize fitting for more complex models





# Affine Transformations

- Assume we know the **correspondences (???)**, how do we get the **transformation**?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$



$$\begin{bmatrix} x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & & & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$




# Affine Transformations

- Linear system with **six** unknowns
- Each match gives us **two linearly independent equations**: need at least **three** to solve for the transformation parameters

$$\begin{bmatrix} \dots & & & & & \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & & & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$




# Feature-Based Alignment Outline



- Extract features



# Feature-Based Alignment Outline




- Extract features
- Compute *putative matches*






# Feature-Based Alignment Outline



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$



# Feature-Based Alignment Outline



- Extract features
- Compute *putative* matches
- Loop:
  - *Hypothesize* transformation  $T$
  - *Verify* transformation (search for other matches consistent with  $T$ )





# Feature-Based Alignment Outline



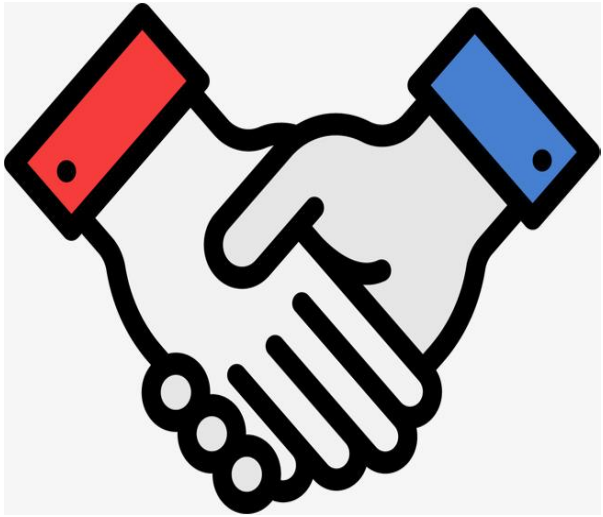
- Extract features
- Compute *putative matches*
- Loop:
  - Hypothesize transformation  $T$
  - Verify transformation (search for other matches consistent with  $T$ )

# Conclusions



# Conclusion

- Fitting techniques
  - Least Squares
  - Total Least Squares
- RANSAC
- Hough Voting
- Alignment as a fitting problem



Thanks



[zhengf@sustc.edu.cn](mailto:zhengf@sustc.edu.cn)