



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

STA-5007: Advanced Natural Language Processing

Lecture 2: Basics of Deep Learning



陈冠华 CHEN Guanhua

Department of Statistics and Data Science

Content

- Introduction
- History
- Model
- Optimization
- Training
- Coding

- To create a function to map an input X into an output Y , $Y = f(X)$
- Examples:

<u>Input X</u>	<u>Output Y</u>	<u>Task</u>
Text	Text in Other Language	Translation
Text	Response	Dialog
Text	Label	Text Classification
Text	Linguistic Structure	Language Analysis

- To create such a system, we can use
 - Manual creation of rules
 - Machine learning from paired data $\langle X, Y \rangle$

Machine Learning

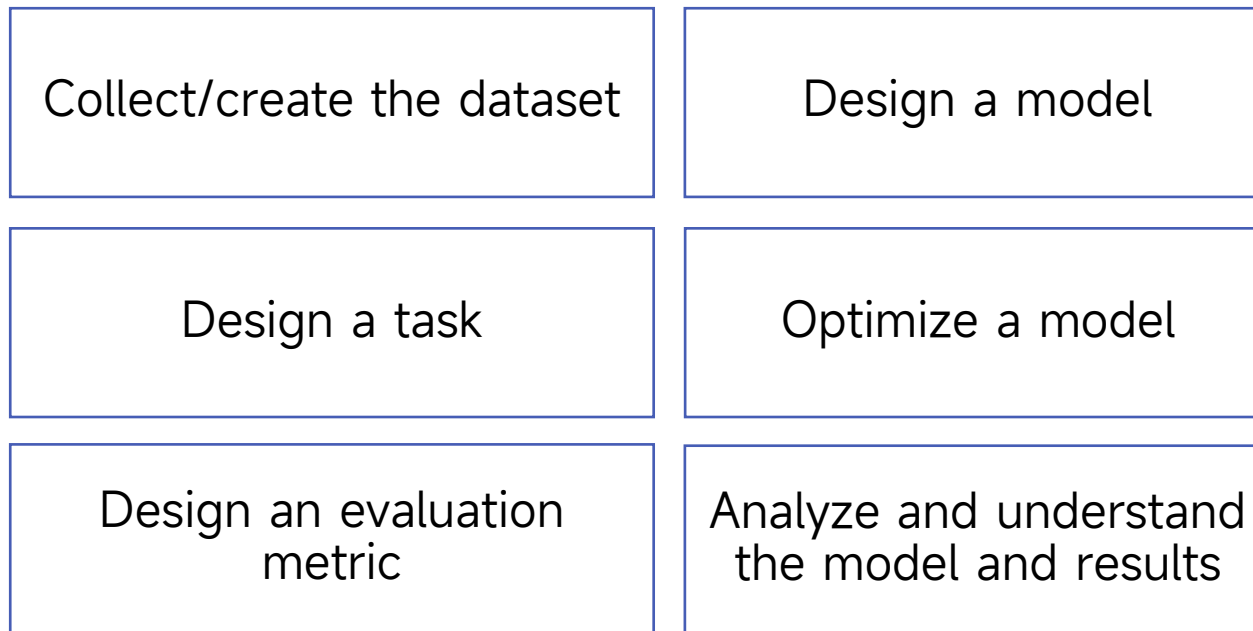
- Statistical approach
 - Generalized linear model, linear regression, logistic regression
 - Gaussian mixture models
 - Support vector machine (SVM)
 - Decision trees, random forests
- Deep learning approach
 - Modeling with different deep neural networks

Why didn't They Work Before?

- Datasets too small
 - For machine translation, not really better until you have 1M+ parallel sentences (and really need a lot more)
- Optimization not well understood
 - Good initialization
 - Momentum (Adagrad/Adam) work best out-of-the-box
- Other innovations
 - Word embedding
 - Dropout, layer normalization, residual connection
 - Large-scale computing system

Deep Learning

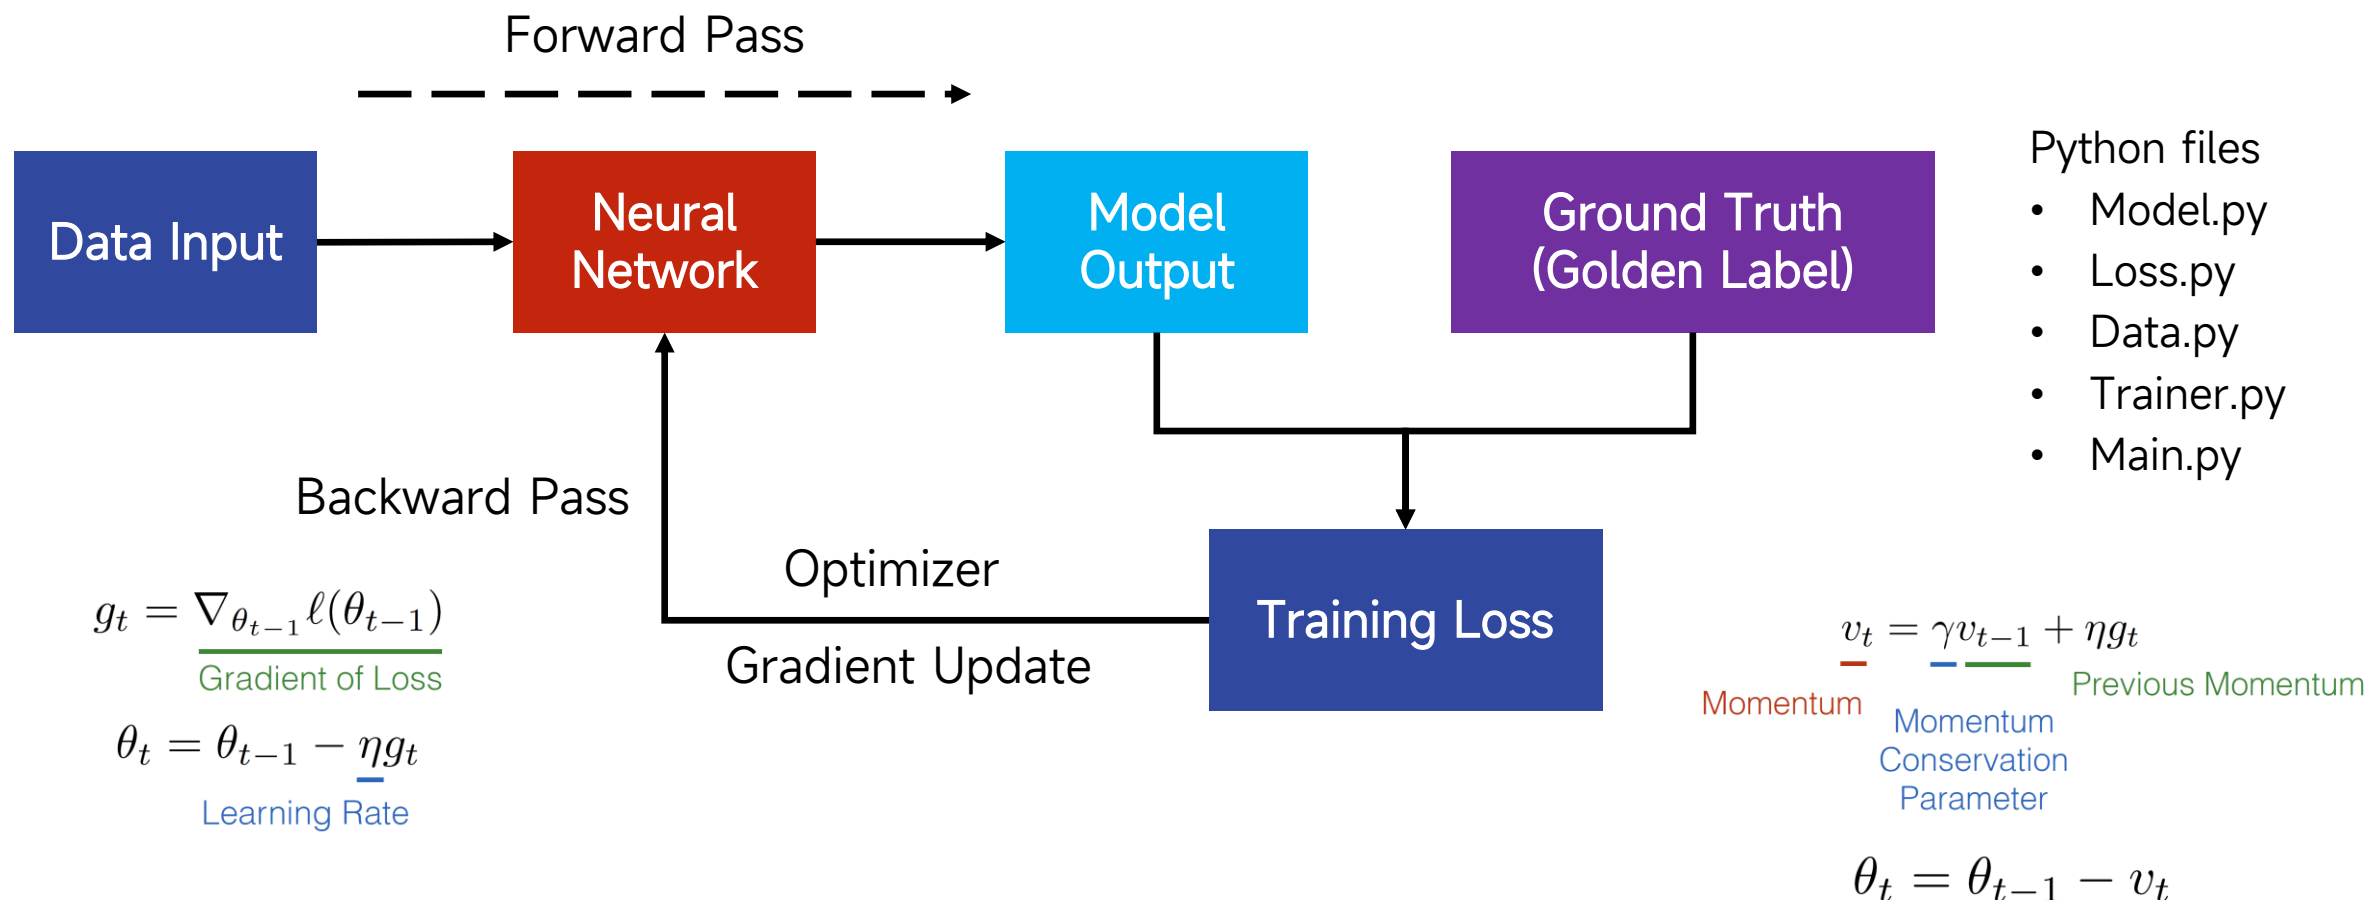
- Modeling with deep neural networks
- Optimized on big data
- Research Contributions



Deep Learning Algorithm Sketch

- Create a model and define a loss
- For each example
 - Forward process: calculate the result (prediction & loss) of that example
 - if training
 - Perform back propagation
 - Update parameters

Deep Learning Algorithm Sketch




Dataset Split

When creating a system, use three sets of data

- Training Set
 - Generally larger dataset, used during system design, creation, and learning of parameters.
- Development/validation Set
 - Smaller dataset for testing different design decisions ("hyper-parameters").
- Test Set
 - Dataset reflecting the final test scenario, do not use for making design decisions.

Data is Very Important



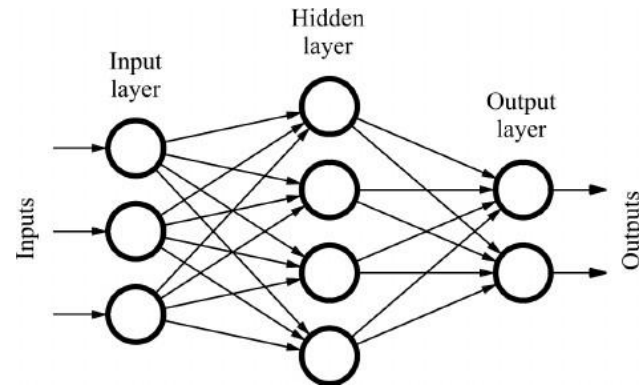
Deep Learning Algorithm Sketch

- Create a model and define a loss
- For each example
 - Forward process: calculate the result (prediction & loss) of that example
 - if training
 - Perform back propagation
 - Update parameters

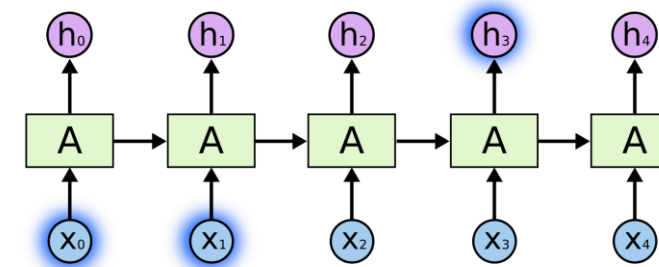
Different Model Structures



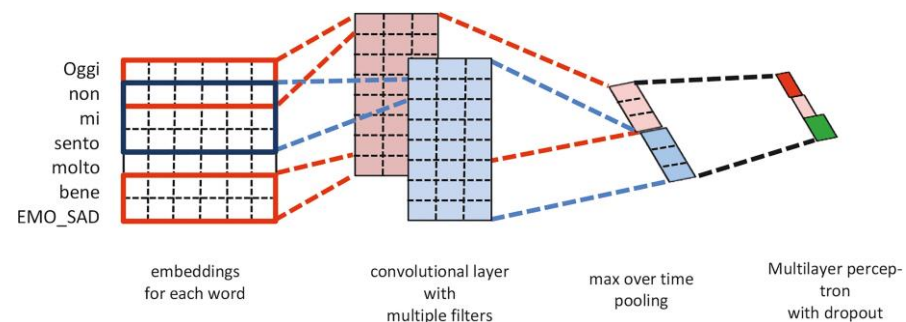
Feed-forward NNs



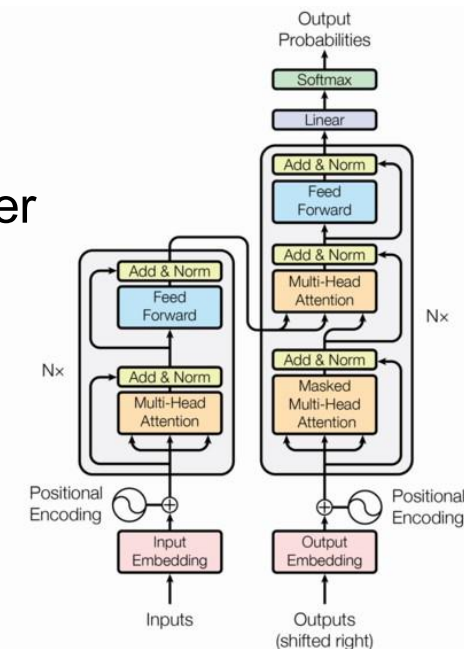
Recurrent NNs



Convolutional NNs

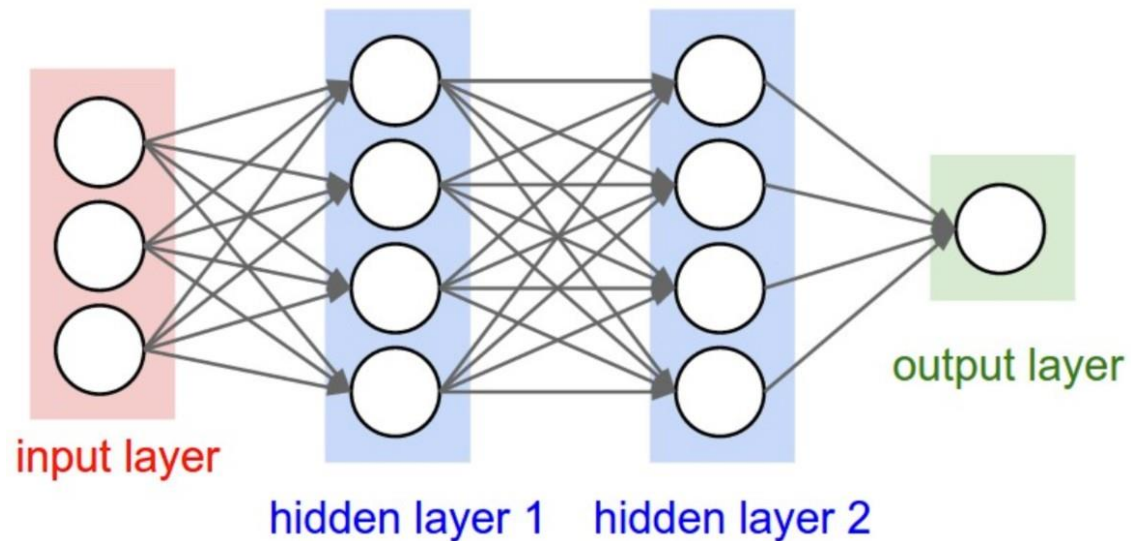


Transformer




Feed-forward Neural Network

- The units are connected with **no cycles**
- The outputs from units in each layer are passed to units in the next higher layer
- No outputs are passed back to lower layers
- Fully-connected (FC) layers



```
1 import torch.nn as nn
2 import torch.nn.functional as F
3
4 class Net(nn.Module):
5     def __init__(self):
6         super().__init__()
7         self.fc1 = nn.Linear(784, 128)
8         self.fc2 = nn.Linear(128, 64)
9         self.fc3 = nn.Linear(64, 10)
10
11     def forward(self, x):
12         x = F.relu(self.fc1(x))
13         x = F.relu(self.fc2(x))
14         x = self.fc3(x)
15         return x
16
```

Feed-forward Neural Network



- Input layer: $\mathbf{x} \in \mathbb{R}^d$

- Hidden layer 1:

$$\mathbf{h}_1 = f(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}) \in \mathbb{R}^{d_1}$$

$$\mathbf{W}^{(1)} \in \mathbb{R}^{d_1 \times d}, \mathbf{b}^{(1)} \in \mathbb{R}^{d_1}$$

- Hidden layer 2:


$$\mathbf{h}_2 = f(\mathbf{W}^{(2)}\mathbf{h}_1 + \mathbf{b}^{(2)}) \in \mathbb{R}^{d_2}$$

$$\mathbf{W}^{(2)} \in \mathbb{R}^{d_2 \times d_1}, \mathbf{b}^{(2)} \in \mathbb{R}^{d_2}$$

- Output layer:

$$\mathbf{y} = \mathbf{W}^{(o)}\mathbf{h}_2, \mathbf{W}^{(o)} \in \mathbb{R}^{C \times d_2}$$

Feed-forward Neural Network



$$h_1^{(1)} = f(w_{1,1}^{(1)}x_1 + w_{1,2}^{(1)}x_2 + w_{1,3}^{(1)}x_3)$$

$$h_2^{(3)} = f(w_{3,1}^{(2)}h_1^{(1)} + w_{3,2}^{(2)}h_1^{(2)} + w_{3,3}^{(2)}h_1^{(3)} + w_{3,4}^{(2)}h_1^{(4)})$$

Non-linearity (activation function) f : \tanh or $ReLU$

Feedforward Neural Network for Classification



- Use **softmax** to get the probability distribution

$$\mathbf{y} = \mathbf{W}^{(o)} \mathbf{h}_2, \mathbf{W}^{(o)} \in \mathbb{R}^{C \times d_2}$$

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{y}) \quad \text{softmax}(\mathbf{y})_k = \frac{\exp(y_k)}{\sum_{j=1}^C \exp(y_j)} \quad \mathbf{y} = [y_1, y_2, \dots, y_C]$$

Training loss:

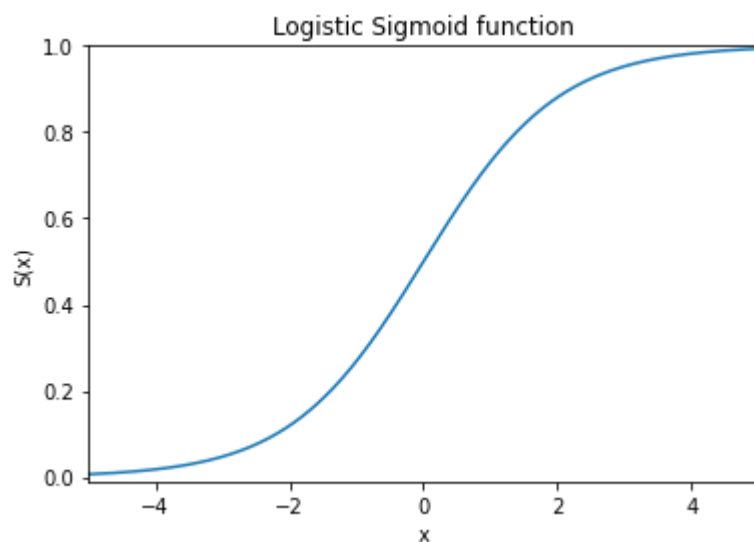
$$\min_{\mathbf{W}^{(1)}, \mathbf{W}^{(2)}, \mathbf{W}^{(o)}} - \sum_{(\mathbf{x}, y) \in D} \log \hat{\mathbf{y}}_y$$
$$\begin{aligned} \mathbf{h}^{(1)} &= \text{ReLU}(\mathbf{W}^{(1)} \mathbf{x}) \\ \mathbf{h}^{(2)} &= \text{ReLU}(\mathbf{W}^{(2)} \mathbf{h}^{(1)}) \\ \hat{\mathbf{y}} &= \text{softmax}(\mathbf{W}^{(o)} \mathbf{h}^{(2)}) \end{aligned}$$

Neural networks are difficult to optimize.
SGD can only converge to local minimum.
Initializations and optimizers matter a lot!

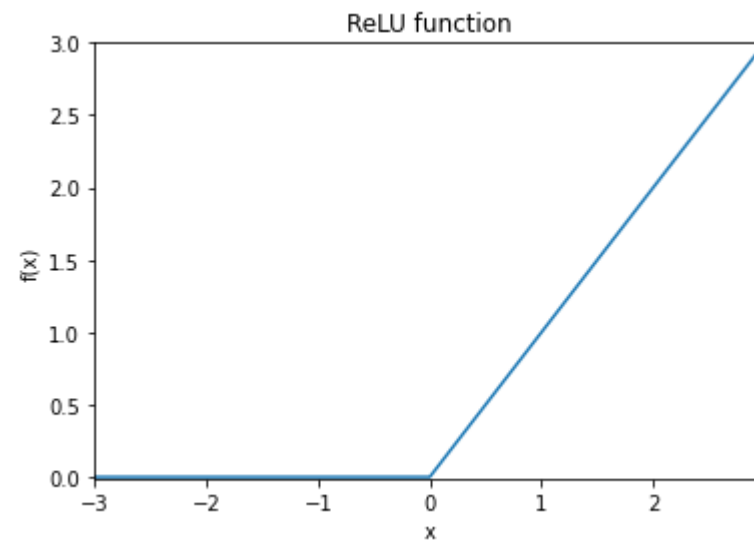
Activation Functions



- Add non-linearities into neural networks
- Allowing the neural networks to learn powerful operations



$$f(x) = \frac{1}{1 + e^{-x}}$$



$$f(x) = \max(x, 0)$$

Activation Functions



- GeLU (Gaussian Error Linear Unit)
 - Used in GPT-3, BERT, and many other models


Applies the Gaussian Error Linear Units function:

$$\text{GELU}(x) = x * \Phi(x)$$

where $\Phi(x)$ is the Cumulative Distribution Function for Gaussian Distribution.

When the approximate argument is 'tanh', Gelu is estimated with:

$$\text{GELU}(x) = 0.5 * x * (1 + \text{Tanh}(\sqrt{2/\pi} * (x + 0.044715 * x^3)))$$



- Given a labeled example (x, \hat{y}) , we use a neural network to estimate the conditional probability and predict the label as

$$\hat{y} = \operatorname{argmax} P_{\theta}(y|x)$$

- We compute how close our prediction w.r.t. the true label by a loss function
- Classification: Cross-Entropy

$$\mathcal{L}(x, y^*) = -\log P_{\theta}(y = y^*|x)$$

- Regression: L1 loss, L2 loss (a.k.a. Mean Square Error)

$$\mathcal{L}(x, y^*) = \|f(x) - y^*\|_1$$

$$\mathcal{L}(x, y^*) = \|f(x) - y^*\|_2$$

Deep Learning Algorithm Sketch

- Create a model and define a loss (i.e., construct a computation graph)
- For each example
 - **Forward process**: calculate the result (prediction & loss) of that example
 - if training
 - **Perform back propagation**
 - Update parameters

Backpropagation

Forward propagation:
from input to output layer

Given: x_1, x_2, x_3
and the class
label y
(a single training
example)

Goal:

$$\frac{\partial L}{\partial W^{(1)}}, \frac{\partial L}{\partial W^{(2)}}, \frac{\partial L}{\partial W^{(o)}}$$

Forward step 1:
Compute $h_1^{(1)}, h_2^{(1)}$

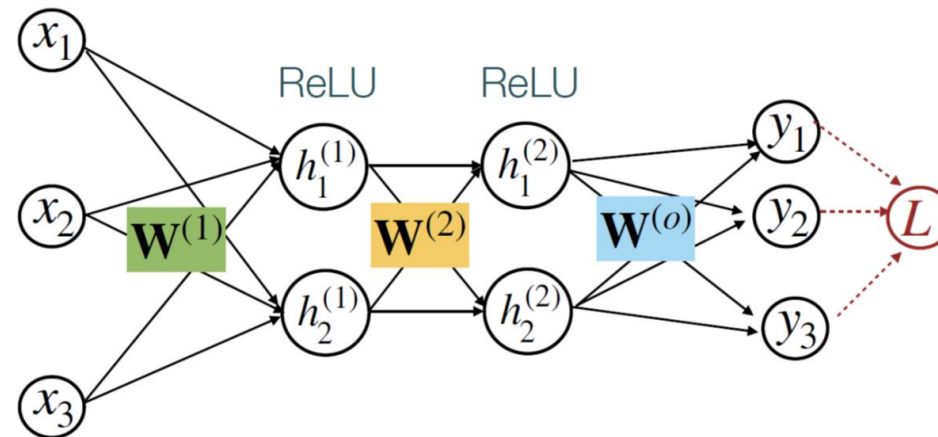
Forward step 2:
Compute $h_1^{(2)}, h_2^{(2)}$

Forward step 3:

Compute y_1, y_2, y_3 and
 $[\hat{y}_1, \hat{y}_2, \hat{y}_3] = \text{softmax}[y_1, y_2, y_3]$

Forward step 4:

Compute loss
 $L = -\log \hat{y}_y$



Back propagation:
from output to input layer

Back step 4:
Compute
 $\frac{\partial L}{\partial W^{(1)}}$

Back step 3:
Compute
 $\frac{\partial L}{\partial h_1^{(1)}}, \frac{\partial L}{\partial h_2^{(1)}}, \frac{\partial L}{\partial W^{(2)}}$

Back step 2:
Compute
 $\frac{\partial L}{\partial h_1^{(2)}}, \frac{\partial L}{\partial h_2^{(2)}}, \frac{\partial L}{\partial W^{(o)}}$

Back step 1:

Compute
 $\frac{\partial L}{\partial y_1}, \frac{\partial L}{\partial y_2}, \frac{\partial L}{\partial y_3}$

Back-propagation in PyTorch



```
1 import torch.nn as nn
2 import torch.nn.functional as F
3
4 class Net(nn.Module):
5     def __init__(self):
6         super().__init__()
7         self.fc1 = nn.Linear(784, 128)
8         self.fc2 = nn.Linear(128, 64)
9         self.fc3 = nn.Linear(64, 10)
10
11     def forward(self, x):
12         x = F.relu(self.fc1(x))
13         x = F.relu(self.fc2(x))
14         x = self.fc3(x)
15         return x
```

```
1 import torch.optim as optim
2
3 net = Net()
4 criterion = nn.CrossEntropyLoss()
5 optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

```
1 outputs = net(inputs)
2 loss = criterion(outputs, labels)
3 loss.backward()
4 optimizer.step()
```

PyTorch did back-propagation for you in this one line of code!

[A toy pytorch example to train an NN model](#)

Deep Learning Algorithm Sketch

- Create a model and define a loss (i.e., construct a computation graph)
- For each example
 - Forward process: calculate the result (prediction & loss) of that example
 - if training
 - Perform back propagation
 - Update parameters

Optimizer Update

- Most deep learning toolkits implement the parameter updates by calling `optimizer.step()` function

```
[59] # Before gradient update
for name, param in model.named_parameters():
    print(name, param)

A Parameter containing:
tensor([[ -0.2267,  0.6521, -0.8193,  0.7723, -0.6456],
        [ 1.2410, -2.4380, -0.5612, -0.1144, -0.2687],
        [ 0.4792,  0.4543, -1.3530,  1.2934, -0.9943],
        [ 0.7565,  0.9449,  0.2796,  0.4703,  0.2926],
        [ 0.4143,  0.5891,  0.4370,  0.6060,  0.0161]])
b Parameter containing:
tensor([ 0.3592,  0.3455, -0.2517, -0.5678, -0.6016], :
c Parameter containing:
tensor([-1.5490], requires_grad=True)
```

Before optimizer update

```
[61] # Gradient update:
optimizer.step()
# After gradient update
for name, param in model.named_parameters():
    print(name, param)

A Parameter containing:
tensor([[ -0.2167,  0.6621, -0.8093,  0.7623, -0.6556],
        [ 1.2510, -2.4280, -0.5512, -0.1244, -0.2787],
        [ 0.4892,  0.4643, -1.3430,  1.2834, -1.0043],
        [ 0.7465,  0.9349,  0.2696,  0.4803,  0.3026],
        [ 0.4043,  0.5791,  0.4270,  0.6160,  0.0261]])
b Parameter containing:
tensor([ 0.3492,  0.3355, -0.2617, -0.5578, -0.5916], :
c Parameter containing:
tensor([-1.5390], requires_grad=True)
```

After optimizer update

Can be updated with standard SGD or Adam optimizer

Standard SGD



- Standard stochastic gradient descent


$$g_t = \nabla_{\theta_{t-1}} \ell(\theta_{t-1})$$

Gradient of Loss

$$\theta_t = \theta_{t-1} - \eta g_t$$

Learning Rate

```
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```



Adam Optimizer

- Most standard optimization option in NLP and beyond
- Considers rolling average of gradient g_t , and momentum m_t

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t \quad \longrightarrow \quad \text{Momentum}$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t \odot g_t \quad \longrightarrow \quad \text{Rolling Average of Gradient}$$

$$\hat{m}_t = \frac{m_t}{1 - (\beta_1)^t} \quad \hat{v}_t = \frac{v_t}{1 - (\beta_2)^t} \quad \longrightarrow \quad \text{Correction of bias}$$

[Further reading: how to use the optimizer in Pytorch](#)

$$\theta_t = \theta_{t-1} - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t \quad \longrightarrow \quad \text{Final update the parameter}$$

```
optimizer = torch.optim.Adam(model.parameters(), lr=0.0005, betas=(0.99, 0.999))
```

Adam Optimizer



- [Gradient descent | Khan Academy](#)
- [Intuition of Adam Optimizer](#)
- [Blog: An updated overview of recent gradient descent algorithms](#)
- [\(paper\) Convex Optimization: Algorithms and Complexity](#)
- [Course: Optimization for Machine Learning](#)

Learning Rate





Figure 4: Effect of learning rate (adapted figure from: <https://www.jeremyjordan.me/nn-learning-rate/>)

[learning rate schedule](#) [[another link](#)], warmup

Tensors

- An n-dimensional array



- Widely used in neural networks
- Parameters in NNs consist of different shape of tensors, which store both their values and gradients (e.g., \mathbf{x} , $\mathbf{x}.\text{grad}$)

Tensor Operations



```
import numpy as np
x = torch.Tensor([[2, 3], [1, 2]])
x = torch.Tensor(np.array([[ -1, 1], [2, 4]]))
x = torch.zeros([2, 3], dtype=torch.int32)
```

create tensors from list, numpy.array

```
import torch
import torch.nn as nn
```

```
x = torch.randn((4, 2))
W = torch.randn((3, 4))
print(x)
print(W)
```

```
tensor([[ -1.5372,  0.0845],
        [ 0.5752,  0.7634],
        [ 0.4265, -0.1287],
        [-1.8629, -0.8520]])
tensor([[ -1.1272, -1.1810,  0.0867,  0.0676],
        [-0.1070,  3.2586,  0.7446, -1.2094],
        [-1.7670,  0.2900, -1.0881, -1.5555]])
```

```
[10] torch.matmul(W, x) # results in a [3,2] matrix
```

```
tensor([[ 0.9646, -1.0656],
        [ 4.6092,  3.4132],
        [ 5.3167,  1.5373]])
```

matrix multiply

```
x = torch.randn((4,2))
z = torch.randn((4,2))
print(x)
print(z)
```

```
tensor([[ 0.0762,  1.5145],
        [-0.4747, -0.9141],
        [ 0.7106,  0.4888],
        [ 0.6959, -0.5305]])
tensor([[ -0.1766,  0.6187],
        [ 0.9254, -0.5931],
        [-0.9162,  0.3209],
        [ 0.0216, -0.7116]])
```

```
[13] x * z # results in a [4,2] matrix
```

```
tensor([[ -0.0135,  0.9371],
        [-0.4392,  0.5421],
        [-0.6510,  0.1569],
        [ 0.0151,  0.3775]])
```


Element-wise matrix multiply

Efficiency Tricks: Mini-batching




- On modern hardware 10 operations of size 1 is much slower than 1 operation of size 10
- Mini-batching combines together smaller operations into one big one
- About padding


Operations w/o Minibatching



Operations with Minibatching



Deep Learning Algorithm Sketch



Different Learnings





南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

- Supervised/unsupervised learning
- Self-supervised learning
- Transfer learning
- Few-shot/zero-shot learning

[Stanford STATS214 / CS229M: Machine Learning Theory](#)

Further Reading

- [\(book\) Information Theory From Coding to Learning](#)
- [Neural Networks: Zero to Hero \(Andrej Karpathy\)](#)
- [Course: Introduction to Deep Learning](#)
- [Course: MIT Introduction to Deep Learning](#)
- [Github free resources for students](#)





南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Thank you